

# PolicyBench

Benchmarking no-tool tax-and-benefit estimation in frontier language models

Max Ghenis\*

## Table of contents

|                                                          |    |
|----------------------------------------------------------|----|
| Abstract                                                 | 1  |
| 1 Introduction                                           | 2  |
| 2 Related work                                           | 2  |
| 3 Benchmark design                                       | 4  |
| 3.1 Headline metric: exact match . . . . .               | 4  |
| 3.2 Near-miss companion: within 1% . . . . .             | 5  |
| 3.3 Secondary metric: bounded continuous score . . . . . | 5  |
| 3.4 Frozen snapshot and open-set status . . . . .        | 7  |
| 4 Data and scenario construction                         | 15 |
| 4.1 United States . . . . .                              | 15 |
| 4.2 Disability in the household facts . . . . .          | 15 |
| 4.3 Reference-output credibility . . . . .               | 17 |
| 5 Results                                                | 25 |
| 5.1 United States leaderboard . . . . .                  | 25 |
| 5.2 Uncertainty . . . . .                                | 26 |
| 5.3 Cost and latency . . . . .                           | 26 |
| 5.4 Simple baselines . . . . .                           | 28 |
| 5.5 Sensitivity to benchmark view . . . . .              | 29 |
| 5.6 Hardest benchmark targets . . . . .                  | 30 |
| 5.7 Zero and positive cases . . . . .                    | 31 |
| 5.8 Open-weight time lag . . . . .                       | 33 |

---

\*PolicyEngine and Axiom Foundation.

|   |                                     |    |
|---|-------------------------------------|----|
| 6 | Failure modes                       | 33 |
| 7 | Limitations                         | 35 |
| 8 | Conclusion                          | 36 |
|   | Appendix A: Structured-output audit | 37 |
|   | Competing interests                 | 38 |

Source: [Article Notebook](#)

## Abstract

PolicyBench evaluates whether frontier language models can estimate household tax and benefit outputs from household facts without tools. This release covers the United States, drawing sampled households from the certified PolicyEngine populace microdataset, scored through `policyengine.py` 4.16.1. It evaluates 42 frontier models on 100 households across 18 output groups, and reports three complementary scores: an exact-match rate as the headline deployability bar (a prediction counts only if it matches the PolicyEngine reference to the dollar for amounts or to the eligibility flag for booleans), a within-1% hit rate as a near-miss-tolerant companion, and a continuous bounded score that awards partial credit for close answers. The headline is exact match on a household-impact-weighted leaderboard: 86% of US reference outputs are exact zeros, and the weighting concentrates scores in the high-dollar output groups a hedge-to-zero strategy cannot answer, so the top model leads an always-zero baseline of 69.5% weighted exact by 24.8 points (Section 3.1 details the construction). The manuscript snapshot is a 100-household public preview, not a protected held-out leaderboard: the current public scenario explorer exposes prompts and reference outputs, so open-set leakage is a central limitation of any public ranking. In the frozen snapshot used for this manuscript (2026-09-22), the top-scoring US model on exact match is GPT-6 Sol at 94.3%, while GPT-5.4 nano is lowest at 64.8%. Within model families, newer is not always stronger: Claude Opus 4.8 scores 76.0% exact, below the 80.7% of the earlier Claude Opus 4.7. Multi-step tax quantities and positive-dollar benefit cases are materially harder than zero cases. The frozen audit annotations cover 7,545 rows whose legacy threshold score is below 1. That universe contains 7,541 of the snapshot's 7,541 exact-match misses and 4 exact hits; 1,843 additional rows with bounded score below 100 were not selected or annotated. Every reference-suspect flag was adjudicated against primary sources and engine traces. The latest audit found defects in `policyengine-us` 1.755.4: the 15 scored references resting on defects fixed upstream after the freeze are regenerated with the fixes, and every output whose reference a defect not yet fixed upstream moves is excluded from scoring for every model (28 are recorded as engine-defect exclusions), as are outputs whose reference depends on an input the prompt never lists. References that rested on projected rather than published 2026 amounts were regenerated with the amounts published before the

references were frozen. The 0.8% of cells with a missing or unparsable answer remain scored as misses rather than dropped. The live benchmark is available at <https://policybench.org>.

## 1 Introduction

Household tax-and-benefit estimation sits between arithmetic and policy discussion. Each case has numeric labels, but generating those labels requires filing status, household composition, income concepts, program thresholds, and jurisdiction-specific rules. PolicyBench measures whether models can map household records to those outputs.

Most current language-model evaluations do not test this mapping directly. General math benchmarks emphasize symbolic manipulation and exact final answers (Cobbe et al. 2021; Hendrycks et al. 2021), while generic question-answering benchmarks emphasize recall or instruction following. PolicyBench instead asks whether models can transform household facts into policy outputs without access to tools or simulators.

This matters for public-facing calculators, analyst workflows, and tax-preparation or screening systems. The benchmark is intended to separate verbal policy fluency from household-level quantitative prediction.

## 2 Related work

The most relevant prior work is tax and statutory reasoning. SARA frames statutory interpretation, including tax-relevant rule application, as a language-understanding problem (Holzenberger et al. 2021). LegalBench broadens this view and includes tax-oriented numeric tasks such as `sara_numeric` (Guha et al. 2023). RuleArena evaluates rule-guided reasoning in regulation-style settings (Zhou et al. 2025), and TaxCalcBench evaluates frontier models on tax calculation from structured return-like inputs (Bock et al. 2025). Shanahan et al. (2025) report a similar split on the Volunteer Income Tax Assistance (VITA) test: models do better on tax knowledge questions than on open-ended calculation tasks.

PolicyBench also sits within a broader literature on quantitative reasoning benchmarks. Canonical math evaluations such as GSM8K and MATH normalized exact final-answer scoring for numeric tasks (Cobbe et al. 2021; Hendrycks et al. 2021). Recent work has questioned pure exact-match scoring in numeric QA and temporal reasoning, arguing for error-aware metrics instead (Abbood et al. 2025). PolicyBench treats this as a measurement design choice rather than a forced pick, leading with a household-impact-weighted exact-match rate and reporting near-miss and bounded companions; the metric definitions, and the weighting that keeps the zero-inflated panel from rewarding hedge-to-zero answers, are specified in the benchmark design section.

Finance-domain benchmarks provide another relevant precedent. FinQA and TAT-QA both show that realistic quantitative reasoning often requires combining structured numbers with domain knowledge rather than solving stylized school-math problems (Chen et al. 2021; Zhu et al. 2021). PolicyBench differs in that its reference outputs are generated by executable tax-benefit microsimulation rather than annotated document questions, but the underlying motivation is similar: domain-specific quantitative reasoning deserves dedicated evaluation.

The nearest public-benefits benchmark is the concurrent Public Benefits Bench, built by Vals AI with the Center for Civic Futures and Code for America (Vals AI 2026). It evaluates models on 459 expert-validated SNAP guidance scenarios (230 in the published test set) spanning all 50 US states plus Guam and the Virgin Islands, and grades free-text answers against expert-written rubric criteria using a judge model that agreed with SNAP policy experts on 80.6% of grading decisions. Where PolicyBench isolates parametric knowledge by withholding tools, Public Benefits Bench varies tool access directly: across models, multi-turn conversation raises pass rates by roughly 7.6 percentage points and web search by about 6.9. The two benchmarks probe complementary failure surfaces — conversational guidance quality for a single program with tools available, versus no-tool numeric estimation across tax and transfer outputs scored deterministically against microsimulation references — and reach consistent conclusions. The best model passes 71.7% of rubric criteria under Public Benefits Bench’s strongest condition, and its authors conclude that “no general-purpose AI model performs well enough to be trusted with SNAP benefits guidance,” matching the deployability gaps PolicyBench reports for unaided calculation.

Finally, PolicyBench depends on two infrastructure literatures that are not themselves large language model (LLM) benchmarks. One is structured-output reliability, since the benchmark relies on multi-output JSON responses and parse coverage rather than free-form prose (Shorten et al. 2024). The other is tax-benefit microsimulation, where systems such as EUROMOD provide the methodological precedent for evaluating policy rules over household microdata (Sutherland and Figari 2013). There is also operational work on applying and evaluating AI systems in public-benefits settings, including caseworker-assist and Supplemental Nutrition Assistance Program (SNAP)-focused evaluations (Nava Labs 2026, 2025; ZenML LLMops Database 2025). PolicyBench combines these strands into a public cross-model benchmark over household-level tax and benefit outputs.

### 3 Benchmark design

The canonical PolicyBench request asks each model to predict every benchmark output for a household in a whole-scenario response under its structured answer contract. 10 of the 42 models received the documented chunked accommodation and instead answered the same prompt template over one- or three-output subsets, reducing each response’s completion demand while preserving the household-level task.

### 3.1 Headline metric: exact match

The headline metric is the exact-match rate: the share of requested outputs the model gets right to the dollar (or, for eligibility flags, to the boolean). This is the deployability bar. A tax filer, a benefit estimator, or a caseworker cannot ship “close” — a prediction off by \$50 is no more usable than one off by \$500, because neither matches the bottom line a downstream system or a household needs. For currency amounts, “exact” means within 1 currency unit of the reference value after numeric parsing. For binary outputs, “exact” means the parsed value is exactly the same 0/1 flag as the reference.

Exact match can be a deceptive headline in a zero-inflated panel — but only when the rate is computed unweighted. In the frozen manuscript snapshot, 86% of reference outputs are exact zeros: most sampled households are not eligible for any given program, owe no capital gains tax, do not phase into a refundable credit, and so on. A model that hedges to zero on every requested output earns full exact credit on every zero case, so a naive unweighted exact rate is dominated by those zeros — a hedge-to-zero model scores roughly 86% unweighted, leaving little room above the floor to separate models that actually understand the rules.

The public leaderboard corrects for this by household-impact weighting. Each output group’s weight is its population-average share of the household’s dollar stake, computed as  $|ref| / \max(|household\_net\_income|, \Sigma |ref|)$  per source household and averaged with calibrated household weights over the full populace population (Section 5.5 gives the full construction). Zero-reference rows are not removed — every row carries its output group’s weight — but the weighting concentrates mass in high-dollar output groups, which are rarely zero, so the share of weight on zero-reference rows falls from their 86% row share to 69.5%. An always-zero baseline therefore scores 69.5% weighted exact rather than the 86% it earns unweighted, and the top model clears that baseline by 24.8 points — a margin comparable to within-1%. Across the 42-model roster, the weighted exact rate spreads about as widely as within-1% and ranks models in nearly the same order, so exact match discriminates between rule comprehension and zero-hedging while preserving the to-the-dollar deployability meaning.

The exact rate is aggregated to a household score and then to a country score, with the same population-derived output weights that the within-1% companion and the bounded score use. Equal household weight preserves comparability across scenarios and slices. The leaderboard reports US results on this fixed test set.

### 3.2 Near-miss companion: within 1%

The within-1% hit rate is the near-miss-tolerant companion to the headline: the share of requested outputs the model gets within one percent of the PolicyEngine reference value. For currency amounts, “within 1%” means  $|pred - ref| \leq 0.01 \times |ref|$  when the reference is nonzero and  $|pred| = 1$  (the same one-currency-unit absolute tolerance used by exact match) when

the reference is zero. Binary outputs are requested as integer 0/1 eligibility flags and scored identically to the exact-match metric.

Within-1% relaxes the exact bar by tolerating rounding and small rate or parameter drift on positive-reference cases, so it credits a model that is consistently close but rarely to-the-dollar. It uses the same per-output household-impact weighting and the same household-equal aggregation as the headline exact rate, and it is reported alongside exact in every leaderboard table. Because both metrics share the weighting, the two rank models almost identically; within-1% sits a few points above exact and is the right view when near-miss accuracy, rather than strict deployability, is the question.

### 3.3 Secondary metric: bounded continuous score

PolicyBench also reports a bounded 0-100 score that awards smooth partial credit for close-but-not-exact amount answers. This is informative for tracking conceptual progress year over year and captures the qualitative gap between a model that is consistently close and one that is consistently far from the reference value. For amount variables, the row-level bounded score is  $\max(0, 1 - |\text{pred} - \text{ref}| / |\text{ref}|)$  when the reference value is nonzero and  $1\{\text{pred} = 0\}$  when the reference value is zero. For binary outputs, it uses exact 0/1 matching. Note that the zero-reference branch is stricter than the headline’s one-currency-unit tolerance — a \$0.50 prediction on a zero-reference amount counts as exact but earns no bounded credit — so the bounded score is not a uniform relaxation of exact match. The exact, within-1%, and within-10% hit rates are reported separately as threshold diagnostics; they are not averaged to create the bounded score.

PolicyBench also tracks mean absolute error and related error metrics. These remain secondary diagnostics under the headline exact-match rate, the within-1% near-miss companion, and the bounded score. The three-metric design follows recent numeric-evaluation literature that questions pure exact-match scoring (Abbood et al. 2025) while keeping exactness as the production bar.

All requested US outputs are annual amounts or annual eligibility indicators for tax year 2026, scored under the tolerances defined above. Missing, unparseable, or non-0/1 answers receive zero score for that requested output.

Aggregation proceeds in three steps. First, each household-output prediction receives a 0-100 score. Second, each household receives one model score: requested output rows are weighted by output-group weights constructed from the full weighting population, then renormalized within the household over outputs that are actually requested for that household. Person-level coverage flags are scored at the person row; their output-group weight is split across the relevant people in the household and is based on PolicyEngine value proxies where available. Third, the country score averages those household scores with equal household weight. The leaderboard reports US results on this fixed test set. Equal-output-group scores and other alternative views are reported as sensitivity checks.

The benchmark requires each model response to include numeric answers and one explanation per requested output. Both the exact-match rate and the bounded score use only the numeric answers. Explanations are retained for scenario exploration and qualitative error analysis; they should not be interpreted as faithful traces of model reasoning.

Because explanations are required, the canonical task measures policy estimation under a public-facing structured-response contract, not isolated arithmetic accuracy. Prompt fairness is part of the benchmark contract. The current release uses one prompt template per country, with no model-specific tuning. Models receive the same household facts and requested outputs. They receive no web or external tool access. The prompt sets unlisted numeric inputs to 0, unlisted boolean or status facts to false, and household characteristics as constant over the tax-benefit year. Provider-specific differences are limited to structured-output transport, request shape, and the reasoning setup in Table 2. The harness sends no reasoning-control or sampling parameters. It leaves reasoning effort, temperature, and related decoding controls at provider defaults. The frozen GPT-5.5 wave uses provider-default reasoning effort.

The forced `tool_choice` suppresses Claude’s extended thinking in this snapshot ([sensitivity/claude-thinking-2026-08.md](#)). A labeled sensitivity with `tool_choice`: "auto" is published beside the board. Claude Fable 5.1, released September 1, 2026, closes the interaction from the API side: it rejects forced tool use with an error, so its row answers through the JSON transport and reasons at the provider default; a sensitivity run with the tool declared under `tool_choice`: "auto" scores 0.9 points above that row. Claude Opus 5.5, released September 21, 2026, rejects forced tool use the same way, so its row also answers through the JSON transport and reasons at the provider default; it has no sensitivity run. The next board version moves every model to `tool_choice`: "auto" so each provider’s default reasoning posture can engage.

Two dimensions on which requests were not identical across models are request shape and answer transport (Table 2). 10 of the 42 models could not reliably complete the whole-scenario request because their serving stacks rejected the structured-output call, exhausted completion budgets mid-response, or timed out. They instead answered the same prompt template over subsets of the requested outputs, one to three outputs per request. Subsetting gives each response fewer requested outputs and a fresh completion budget over the same household facts, so the accommodation, if it moves scores at all, should favor the chunked models; their scores are not strictly comparable to whole-scenario scores. This accommodation is closed going forward: a model added after this snapshot either answers the canonical whole-scenario request or is listed as not scorable rather than accommodated.

One row, GLM-5.3-Flash (preview), answered as a cloaked OpenRouter listing called "Ox Alpha" and is labeled a preview. It was publicly callable under the canonical whole-scenario request and forced-tool transport, and its date is its OpenRouter listing date. After the row’s run, Z.ai identified the listing as GLM-5.3-Flash (OpenRouter 2026); the row carries the disclosed name and keeps the preview qualifier, because whether the checkpoint that answered is the one later released under that name has not been stated, and it is distinct from the GLM-5.3 row, which is Z.ai’s larger model called under its own id. Its artifact key stays `ox-alpha`, the name it was queried under.

### 3.4 Frozen snapshot and open-set status

Leaderboard positions are version-sensitive, so manuscript claims refer to the frozen source-run exports in Table 1 rather than to the live site. The committed source-run exports are the manuscript artifacts. The public site exposes the current prompts, predictions, explanations, and reference outputs for transparency. This makes the public leaderboard open-set: models, model providers, or benchmark users could learn from the released cases before later runs. Protected leaderboard claims would require a separate held-out or rotating set.

Source: [Article Notebook](#)

Table 1: Frozen manuscript snapshot.

| Item                                 | Value                                                                                                |
|--------------------------------------|------------------------------------------------------------------------------------------------------|
| Snapshot date                        | 2026-09-22                                                                                           |
| Model response date                  | 2026-06-12 to 2026-09-22                                                                             |
| Policy period                        | US tax year 2026                                                                                     |
| Frozen export                        | paper/snapshot/20260501/runs/<br>us_full_run_20260612_poli-<br>cyengine_4_16_1_populace/data.json.gz |
| Frozen export SHA-256 prefix         | US 3c5f5bbf3612                                                                                      |
| Snapshot manifest                    | paper/snapshot/20260501/manifest.json                                                                |
| Audit annotations                    | annotations/us_full_run_20260612_poli-<br>cyengine_4_16_1_populace                                   |
| US scenarios SHA-256 prefix          | 71b16212f0c0                                                                                         |
| US reference outputs SHA-256 prefix  | 800a19bc1c22                                                                                         |
| Benchmark spec SHA-256 prefix        | ce233f8cbb05                                                                                         |
| Frozen prompt payload SHA-256 prefix | US 1c3f94362d0b                                                                                      |
| Scoring code SHA-256 prefix          | baaf2506800b                                                                                         |
| US run label                         | us_full_run_20260612_poli-<br>cyengine_4_16_1_populace                                               |
| Household sample seed                | 42                                                                                                   |
| PolicyEngine.py                      | 4.16.1                                                                                               |
| PolicyEngine-US                      | policyengine-us 1.755.4                                                                              |
| US dataset                           | populace_us_2024<br>(populace-us-2024-5da5a95-20260611)                                              |
| Households                           | 100 US                                                                                               |
| Models                               | 42                                                                                                   |
| Output groups                        | 18 US                                                                                                |
| Condition                            | No tools, no web access; structured responses<br>covering every requested output                     |
| Response contract                    | Numeric answer and non-empty explanation<br>for every requested output                               |

Table 2: Model run configuration in the frozen manuscript snapshot. Every model answers through the transport its model card records, with no external tools, provider-default temperature and sampling, no reasoning-control parameters, and the same household facts and requested outputs. Transport is the answer contract: a forced answer-schema tool call (tool\_choice set to the answer tool) or a JSON object, the latter where the provider rejects a forced tool or where the model card selects JSON. Supervised-run fingerprints pin answer contract, request shape, and completion ceiling for 13 rows; tool choice for 12 rows; reasoning setup and timeouts for three rows. Reasoning setup and timeouts for the other ten fingerprinted rows, and all fields for the other 29 rows, are the harness registry as frozen in the snapshot’s serving-configuration file. Rows with an output count answered the template over subsets of the requested outputs, an accommodation that predates the canonical whole-scenario rule described in the text. Released is the first public availability (paid tiers count; trusted-tester previews do not), compiled from vendor announcements and contemporaneous press; grok-build-0.1’s date rests on secondary trackers. Leaderboard artifact keys match the final segment of the provider id (qwen-3.7-max adds a hyphen). A dagger marks open-weight models; Kimi K3’s weights, announced at its API launch (Willison 2026), shipped on Hugging Face on July 27, 2026 under a custom license (Tech Times 2026). GLM-5.3-Flash (preview) was queried as the OpenRouter stealth listing Ox Alpha (artifact key ox-alpha); Z.ai identified the listing as GLM-5.3-Flash after its run (OpenRouter 2026), and its Released date is its August 20, 2026 OpenRouter listing date.

| Model            | Provider ID               | Released   | Parsed          | Transport   | Request shape        | Reasoning setup                                                |
|------------------|---------------------------|------------|-----------------|-------------|----------------------|----------------------------------------------------------------|
| Claude Fable 5   | claude-fable-5            | 2026-06-09 | 1,932/<br>1,932 | forced tool | 1 output/<br>request | thinking<br>does not<br>engage<br>under<br>forced tool<br>call |
| Claude Fable 5.1 | claude-fable-5-1          | 2026-09-01 | 1,932/<br>1,932 | JSON        | whole<br>scenario    | provider<br>default;<br>16,384-<br>token<br>shared<br>budget   |
| Claude Haiku 4.5 | claude-haiku-4-5-20251001 | 2025-10-15 | 1,932/<br>1,932 | forced tool | 1 output/<br>request | provider<br>default                                            |

| Model                  | Provider ID                                         | Released    | Parsed          | Transport   | Request shape        | Reasoning setup                                    |
|------------------------|-----------------------------------------------------|-------------|-----------------|-------------|----------------------|----------------------------------------------------|
| Claude Opus 4.7        | claude-opus-4-7                                     | 2026-04-16  | 1,932/<br>1,932 | forced tool | 1 output/<br>request | provider default                                   |
| Claude Opus 4.8        | claude-opus-4-8                                     | 2026-05-28  | 1,932/<br>1,932 | forced tool | 1 output/<br>request | provider default                                   |
| Claude Opus 5          | claude-opus-5                                       | 2026-07-24  | 1,932/<br>1,932 | forced tool | whole scenario       | thinking<br>does not engage under forced tool call |
| Claude Opus 5.5        | claude-opus-5-5                                     | 2026-09-21  | 1,932/<br>1,932 | JSON        | whole scenario       | provider default;<br>16,384-token shared budget    |
| Claude Sonnet 4.6      | claude-sonnet-4-6                                   | 2026-02-17  | 1,932/<br>1,932 | forced tool | 1 output/<br>request | provider default                                   |
| Claude Sonnet 5        | claude-sonnet-5                                     | 2026-06-30  | 1,932/<br>1,932 | forced tool | 1 output/<br>request | thinking<br>does not engage under forced tool call |
| DeepSeek V4 Flash 0731 | open-router/<br>deepseek/<br>deepseek-v4-flash-0731 | 2026-07-31† | 1,932/<br>1,932 | forced tool | whole scenario       | provider default;<br>98,304-token shared budget    |
| DeepSeek V4 Pro        | deepseek/<br>deepseek-v4-pro                        | 2026-04-24† | 1,932/<br>1,932 | JSON        | whole scenario       | provider default;<br>16,384-token shared budget    |

| Model                         | Provider ID                                       | Released    | Parsed          | Transport   | Request shape  | Reasoning setup                                 |
|-------------------------------|---------------------------------------------------|-------------|-----------------|-------------|----------------|-------------------------------------------------|
| DeepSeek V4 Pro 0813          | open-router/<br>deepseek/<br>deepseek-v4-pro-0813 | 2026-08-13† | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>98,304-token shared budget |
| Gemini 3 Flash Preview        | gemini/<br>gemini-3-flash-preview                 | 2025-12-17  | 1,932/<br>1,932 | JSON        | whole scenario | provider default                                |
| Gemini 3.1 Flash Lite Preview | gemini/<br>gemini-3.1-flash-lite-preview          | 2026-03-03  | 1,932/<br>1,932 | JSON        | whole scenario | provider default                                |
| Gemini 3.1 Pro Preview        | gemini/<br>gemini-3.1-pro-preview                 | 2026-02-19  | 1,932/<br>1,932 | JSON        | whole scenario | provider default                                |
| Gemini 3.5 Flash              | gemini/<br>gemini-3.5-flash                       | 2026-05-19  | 1,932/<br>1,932 | JSON        | whole scenario | provider default                                |
| Gemini 3.5 Flash-Lite         | gemini/<br>gemini-3.5-flash-lite                  | 2026-07-21  | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |
| Gemini 3.6 Flash              | gemini/<br>gemini-3.6-flash                       | 2026-07-21  | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |
| Gemini 3.7 Flash              | gemini/<br>gemini-3.7-flash                       | 2026-08-13  | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |

| Model            | Provider ID                      | Released    | Parsed          | Transport   | Request shape     | Reasoning setup                                 |
|------------------|----------------------------------|-------------|-----------------|-------------|-------------------|-------------------------------------------------|
| Gemini 3.8 Flash | gemini/<br>gemini-3.8-flash      | 2026-09-02  | 1,932/<br>1,932 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| GLM-5.2          | open-router/<br>z-ai/<br>glm-5.2 | 2026-06-13† | 1,799/<br>1,932 | JSON        | 3 outputs/request | provider default;<br>16,384-token shared budget |
| GLM-5.3          | open-router/<br>z-ai/<br>glm-5.3 | 2026-08-14† | 1,861/<br>1,932 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| GPT-5.4 mini     | gpt-5.4-mini                     | 2026-03-17  | 1,932/<br>1,932 | forced tool | whole scenario    | provider default                                |
| GPT-5.4 nano     | gpt-5.4-nano                     | 2026-03-17  | 1,932/<br>1,932 | forced tool | whole scenario    | provider default                                |
| GPT-5.5          | gpt-5.5                          | 2026-04-23  | 1,932/<br>1,932 | forced tool | 3 outputs/request | provider default;<br>16,384-token shared budget |
| GPT-5.6 Luna     | gpt-5.6-luna                     | 2026-07-09  | 1,932/<br>1,932 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| GPT-5.6 Sol      | gpt-5.6-sol                      | 2026-07-09  | 1,932/<br>1,932 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |

| Model          | Provider ID        | Released   | Parsed          | Transport   | Request shape  | Reasoning setup                                 |
|----------------|--------------------|------------|-----------------|-------------|----------------|-------------------------------------------------|
| GPT-5.6 Terra  | gpt-5.6-terra      | 2026-07-09 | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |
| GPT-6 Astra    | gpt-6-astra        | 2026-09-04 | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |
| GPT-6 Luna     | gpt-6-luna         | 2026-09-22 | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |
| GPT-6 Sol      | gpt-6-sol          | 2026-09-22 | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |
| Grok 4.3       | xai/<br>grok-4.3   | 2026-04-17 | 1,932/<br>1,932 | forced tool | whole scenario | provider default                                |
| Grok 4.5       | xai/<br>grok-4.5   | 2026-07-08 | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |
| Grok 4.6       | xai/<br>grok-4.6   | 2026-08-12 | 1,932/<br>1,932 | forced tool | whole scenario | provider default;<br>16,384-token shared budget |
| Grok Build 0.1 | xai/grok-build-0.1 | 2026-05-29 | 1,932/<br>1,932 | forced tool | whole scenario | provider default                                |

| Model                      | Provider ID                                   | Released    | Parsed          | Transport   | Request shape     | Reasoning setup                                    |
|----------------------------|-----------------------------------------------|-------------|-----------------|-------------|-------------------|----------------------------------------------------|
| Inkling                    | open-router/<br>thinking-machines/<br>inkling | 2026-07-15† | 1,932/<br>1,932 | forced tool | whole scenario    | provider default;<br>49,152-token<br>shared budget |
| Kimi K2.6                  | open-router/<br>moon-shotai/<br>kimi-k2.6     | 2026-04-20† | 1,538/<br>1,932 | JSON        | 3 outputs/request | provider default;<br>16,384-token<br>shared budget |
| Kimi K3                    | open-router/<br>moon-shotai/<br>kimi-k3       | 2026-07-16† | 1,873/<br>1,932 | JSON        | whole scenario    | provider default;<br>49,152-token<br>shared budget |
| MiniMax M3                 | open-router/<br>minimax/<br>minimax-m3        | 2026-06-01† | 1,932/<br>1,932 | forced tool | whole scenario    | provider default;<br>16,384-token<br>shared budget |
| GLM-5.3-Flash<br>(preview) | open-router/<br>stealth/<br>ox-alpha          | 2026-08-20  | 1,932/<br>1,932 | forced tool | whole scenario    | provider default;<br>16,384-token<br>shared budget |
| Qwen3.7-max                | open-router/<br>qwen/<br>qwen3.7-max          | 2026-05-19  | 1,932/<br>1,932 | JSON        | 3 outputs/request | provider default;<br>16,384-token<br>shared budget |

| Model       | Provider ID                          | Released   | Parsed          | Transport | Request shape  | Reasoning setup                                 |
|-------------|--------------------------------------|------------|-----------------|-----------|----------------|-------------------------------------------------|
| Qwen3.8-Max | open-router/<br>qwen/<br>qwen3.8-max | 2026-08-03 | 1,932/<br>1,932 | JSON      | whole scenario | provider default;<br>98,304-token shared budget |

Source: [Article Notebook](#)

## 4 Data and scenario construction

### 4.1 United States

The US benchmark is built from the certified PolicyEngine populace dataset (populace\_us\_2024) using PolicyEngine US. The dataset contains 160,858 people in 75,112 households, with household, tax-unit, SPM-unit, family, marital-unit, and person records for PolicyEngine calculations. The sampled households are filtered to keep a single-tax-unit, single-family, single-Supplemental Poverty Measure (SPM)-unit structure with at least one adult and a supported filing status. Of the 75,112 households, 63,128 (84.0%) pass the filter and form the eligible draw (this filter-pass rate is coincidentally close to the 86% zero-output share discussed under the benchmark design; the two figures are unrelated). The 16.0% excluded by the filter include multi-tax-unit households (e.g., adult roommates), multi-family households, multi-SPM-unit households, and households whose head reports a filing status outside the supported set. These excluded compositions are exactly the kind of cases where federal/state credit allocations and benefit-unit rules become hardest, so the eligible draw is a tractable subset rather than the full distribution of US households. Prompts include nonzero promptable raw inputs across relevant entities rather than a hand-curated summary, so the models see many of the same facts the simulator receives. Filing status is not stated in the prompt; the reference computation infers it from tax-unit role flags. Models therefore see the same household facts that drive the reference filing-status assignment, but they do not receive that assignment as a label.

### 4.2 Disability in the household facts

The household facts carry one disability fact, while the programs apply several. The certified microdata marks a person as disabled when they report any of the six disability-difficulty items in the Current Population Survey, and the prompt passes that flag through as is disabled: 33

of the 177 people in the benchmark households carry it. The law defines disability program by program, and policyengine-us follows it with a separate input for each determination. SSI requires its own disability criterion and no substantial gainful activity (42 U.S.C. 1382c(a)(3)). SNAP counts a member as disabled through receipt-based routes, among them receiving SSI or Social Security disability benefits and a qualifying veterans' status (7 U.S.C. 2012(j)); policyengine-us approximates the SSI route with SSI disability status, which does not require receiving SSI. Medicare's disability route before age 65 requires 24 months of entitlement to Social Security disability benefits (42 U.S.C. 426(b)), the route policyengine-us models. The tax code uses permanent and total disability for dependents (26 U.S.C. 152), retirement on disability for the credit for the elderly and disabled (26 U.S.C. 22), and incapacity for self-care for the child and dependent care credit (26 U.S.C. 21). The general flag enters none of these determinations; in the reference engine it reaches mainly state provisions, and in SNAP only the work-requirement and student rules. No benchmark person carries any of the program-specific inputs. The certified build never sets SSI's disability criterion, and the months of Social Security disability receipt are integer-typed, which the prompt builder does not expose. The references of a person the prompt calls disabled therefore take the non-disabled path through each of these determinations unless another listed fact establishes it: listed Social Security disability income, for example, makes a person disabled for SNAP. A careful reader of is disabled could take either path, so wherever the other reading moves a reference, the output is excluded from scoring for every model (Section 4.3): SSI and SNAP outputs under SSI's disability criterion, and the SNAP output whose engine defect turns on whether its head is a disabled member. Separately, the prompt lists Social Security disability income but not how long it has been received, so the Medicare eligibility of the five household heads with that income, which turns on 24 months of it, is excluded whether or not the prompt calls them disabled. Where the other reading moves nothing, the output stays scored. The next scenario refresh defines the prompt's disability fact and states the program-specific determinations as separate facts (PolicyEngine/policybench#165, PolicyEngine/microcosm#876).

The current US release requests 18 scored output groups spanning federal income tax, refundable credits, payroll and self-employment tax, state and local income tax, Supplemental Nutrition Assistance Program (SNAP), Supplemental Security Income (SSI), Temporary Assistance for Needy Families (TANF), school-meal eligibility, and person-level coverage eligibility for the Special Supplemental Nutrition Program for Women, Infants, and Children (WIC), Medicaid, the Children's Health Insurance Program (CHIP), Medicare, Head Start, and Early Head Start; several of these categories span two scored groups each (federal and state refundable credits, payroll and self-employment tax, state and local income tax, free and reduced-price school meals), which is how the named categories reconcile to the 18 output groups. The source run also requested the Affordable Care Act (ACA) Premium Tax Credit (PTC), but explanation audits showed that the prompt could be misleading when households lacked plan-specific Marketplace information. We therefore preserve those raw responses but exclude PTC from the canonical scored leaderboard until the prompt contract is revised.

The output scope is intentionally narrower than the full PolicyEngine model. Table 3 summarizes the inclusion rule. The benchmark asks for WIC eligibility rather than a WIC dollar amount;

WIC dollar values are used only as impact-weight proxies for coverage flags, not as requested model outputs.

Source: [Article Notebook](#)

Table 3: Output-selection rationale.

| Scope decision          | Rationale                                                                                                                                                                                                                     |
|-------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Included                | Direct tax, credit, benefit, health-support, and coverage outputs that a household-facing model could plausibly be asked to estimate from household facts.                                                                    |
| Excluded                | Intermediate tax bases, payroll subcomponents, and outputs that mainly require unavailable history, restricted local market data, restricted program-administration data, or take-up assignment rather than rule calculation. |
| Binary coverage outputs | Requested as 0/1 eligibility flags and scored as classification tasks; their dollar values are used only as impact-weight proxies, not as requested model outputs.                                                            |
| WIC                     | The benchmark asks for person-level WIC eligibility. It does not ask models to estimate a WIC dollar amount.                                                                                                                  |

Source: [Article Notebook](#)

### 4.3 Reference-output credibility

PolicyBench treats PolicyEngine outputs as benchmark reference outputs, not as administrative records. The reference source is nevertheless stronger than an ad hoc answer key: PolicyEngine is open source, used for household calculators and reform analysis, and externally checked in specific domains. No. 10 Downing Street’s data science team adapted PolicyEngine’s open-source microsimulation model for experimental policy simulation, with validation against external projections before use (Woodruff 2026; Ghenis 2026). In the US, PolicyEngine (2024) reports matching the National Bureau of Economic Research (NBER) TAXSIM-35 model (Feenberg and Coutts 1993) to the cent on the vast majority of cases for the 2021 tax year across hundreds of thousands of tax units per state, with state-specific differences documented in the integration tests. We do not restate that comparison as a single percentage because the published source uses qualitative phrasing rather than a headline accuracy number. PolicyEngine has also signed

a memorandum of understanding (Ghenis and Makarchuk 2025) with the Federal Reserve Bank of Atlanta for future validation work against its Policy Rules Database (2026). The Atlanta Fed sources are a caveat rather than evidence of completed validation for this benchmark: they document planned collaboration and the comparison source, not finished checks of the frozen PolicyBench outputs. Taken together, these sources support using PolicyEngine as a transparent reference implementation with partial external validation, but they do not validate every benchmark output.

This does not make PolicyEngine infallible, so we treated the reference outputs as claims to be attacked rather than assumptions. Across development and four audit waves, the model-assisted review covered 7,545 rows whose legacy threshold score is below 1. That selected universe contains 7,541 exact-match misses and 4 exact hits; 1,843 additional rows with bounded score below 100 were not annotated. Reviewers adjudicated each reference-suspect flag by hand against primary sources, including statutes, regulations, and published program parameters, and against the engine’s computation traces. Before the September 22 wave, reviewers had raised 35 reference-suspect flags. Thirty-one dissolved under adjudication, with the reviewer’s hypothesis failing against a primary source (stale thresholds, missed program mergers, missed federal-conformity elections, misread eligibility-income definitions) or against the engine trace. Four were genuine: two Head Start eligibility flags and one New Jersey state income tax cell reflected policyengine-us defects (an age-window bound and a filing-threshold floor the benchmarked path bypassed), and one SNAP cell rested on a scenario whose encoded inputs did not match the disability status its prompt implied. The engine defects were fixed upstream (policyengine-us #8845 and #8846, plus #8844, found in the same review but affecting no benchmark cell), and this snapshot’s reference outputs are generated with the fixed engine, policyengine-us 1.755.4. The mis-encoded scenario is disclosed under Limitations pending the next scenario refresh. Table 4 summarizes the reviewed discrepancy classes and outcomes.

The audit wave run with the September 22 additions used a stronger judge, Claude Opus 5.5, and it found what the earlier waves had not. The judge flagged 39 references. Adjudication against primary sources and sandbox runs of the reference engine affirmed 5 of them, traced 8 to inputs the prompt never lists, replaced 6 with a regenerated reference, and confirmed 20 as defects in policyengine-us 1.755.4. The judge’s stated hypothesis was not always the defect: a Wisconsin flag blamed a missing retirement exclusion that the engine does apply, on a path the benchmarked output skips, and a Montana flag cited a child tax credit that was never enacted, while the reference was wrong for another reason. The judge also sees only the cases a new model joins, so its flags sample each defect rather than bound it. Each root cause’s rule was therefore implemented as a sandbox fix on the same engine version and every reference recomputed under it. Of those root causes, 8 had been fixed in policyengine-us after the references were frozen (#8839, #9162, #9301, #9313, #9318, #9363 and #9586): capital gain distributions, New York’s renter cap, the CalEITC’s lookup at adjusted gross income, and the SNAP rounding and SNAP child support option described below. The scored references they move are regenerated on the same engine version with the fix applied, 15 in all; an output one of them moves that an unfixed cause also moves stays excluded. The capital gain fix as first built also extended Wisconsin’s 30 percent capital gain exclusion (Wis. Stat. 71.05(6)(b)9) to

the distributions. The upstream fix does not, so that part is recorded as its own defect, not fixed upstream, measured on top of the upstream fix and Wisconsin’s published amounts, and the two outputs it moves are excluded (one of them already excluded for another Wisconsin defect). Every output a defect not yet fixed upstream moves by more than a dollar is excluded from scoring for every model: 28 outputs in 20 households are recorded as engine-defect exclusions across 11 root causes, and 8 of them had never been flagged (an output such a defect moves that was already excluded for an unlisted input keeps that record). They stay excluded until references built on a fixed engine replace them. A second model, GPT-6 Astra, then re-derived the sandbox fixes the first pass built from primary sources and reran their sweeps. It narrowed several fixes to the benchmark year and filters they reach, reclassified two root causes as unlisted inputs (the type of a household’s survivor benefits, and whether listed mortgage interest is on the home the household occupies), and found that a California credit move came from a separate defect: the engine never looks the CalEITC up a second time at adjusted gross income, as the credit worksheet requires.

The same wave tested a failure mode no judge flagged. `policyengine-us` fills a parameter’s values past its last encoded year by projecting it with a price index, so a 2026 reference can rest on a forecast rather than on the amount a government published. Every reference was recomputed with each projected 2026 value held at its last encoded value, and each projected amount behind a moved output was then sourced. Where the government published the 2026 amount before the freeze, the reference takes it: Minnesota’s, Michigan’s and Missouri’s brackets, deductions and exemptions, and Maryland’s withholding allowance, which reaches federal tax through the state and local tax deduction. Wisconsin had published its 2026 brackets and deductions too, but every Wisconsin output they move is excluded for an engine defect, so they enter only those outputs’ corrected values. Where no 2026 amount published before the freeze was found, the reference takes the last amount published: California’s indexed amounts, whose 2026 factor rests on June 2026 prices; the IRS optional sales tax tables, whose 2026 edition had not been found by September 22; Idaho’s zero-rate threshold; Maryland’s return deduction; and SNAP’s FY2027 figures, published on 2026-08-21. Illinois’s exemption and Minnesota’s child credit projections matched the published amounts. These corrections change parameter values and no rule, so the affected references were regenerated with the same engine version rather than excluded: 23 references, most of them SNAP. With the upstream fixes, 26 references in 24 households are regenerated, 13 of them SNAP. The reference sidecar records each regenerated value with the fix that produced it, and the fixes, the recomputation of every reference under each, and the verification reports are committed with the snapshot (`reference_audit/2026-09-22`).

The SNAP convention as first built also corrected the engine’s SNAP arithmetic, and an independent review of the snapshot separated those corrections from the parameter holds. They are engine defects, each fixed in `policyengine-us` after the freeze (`#9162` and `#9318`), so the SNAP references apply the fixes: the engine kept cents in the allotment, where 7 CFR 273.10(e)(2)(ii)(A) requires a whole-dollar allotment; it returned a minimum benefit of \$23.84 a month, where 7 CFR 273.10(e)(2)(ii)(C) rounds 8 percent of the one-person maximum to the nearest dollar (\$24); it floored net income, which 7 CFR 273.10(e)(1)(ii) requires a state to

round to the nearest dollar at each step or to calculate by its TANF procedure, which may keep cents; and it compared income with unrounded income standards. Given the engine’s household size, eligibility and floored net income, the Axiom rules engine’s encoding of 7 U.S.C. 2017(a) reproduces the whole-dollar allotment of all 19 households the first SNAP convention module changed, in each month from January to September (171 of 171 monthly comparisons), and the federal minimum benefit; it did not test New Jersey’s state minimum, and it covers the allotment rounding, not the net income rounding the published references also apply. Recomputing SNAP with net income’s cents kept, rather than rounded to the nearest dollar as the upstream fix does, moves no scored reference by more than the dollar tolerance. One SNAP defect is not fixed upstream: the engine still grants the heat-and-eat standard utility allowance to households without an elderly or disabled member, which P.L. 119-21 section 10103 ended in July 2025, and the scored output it moves is excluded. That household’s prompt calls its head disabled but states none of the routes by which 7 U.S.C. 2012(j) defines a disabled member, such as receiving SSI or Social Security disability benefits; read the other way, the allowance would stand, and so would the reference the publication conventions and upstream SNAP fixes give (\$3,576, which the output would carry if scored), so the output is excluded on either reading. One more SNAP defect was found after the September 22 release and fixed upstream on 2026-09-24: the engine’s values for the SNAP child support option carried the opposite meaning from the one its formulas read. A state may exclude legally obligated child support paid from gross income (7 CFR 273.9(c)(17)); a state that does not must deduct it from net income (273.9(d)(5)). For 2026 the engine excluded it in 37 jurisdictions that USDA’s SNAP State Options Report lists as deducting it, Michigan among them, and deducted it in 9 of the 14 it lists as excluding it. Recomputed under the fix (policyengine-us #9586), one reference moves: a one-person Michigan household that pays child support stayed within the 200 percent gross limit of Michigan’s broad-based categorical eligibility only because the engine excluded the payments, so the engine granted it the \$24 monthly minimum; counting the payments, as Michigan does, leaves it ineligible, and the regenerated reference is \$0. A revision of the release excluded that output while the fix was open; the current release regenerates it with the fix.

The audit tooling itself required hardening, which we report because it changes how much the annotations can be trusted. The derivation narratives displayed alongside reference values are generated from engine computation traces. After earlier model-written narratives were found to misattribute eligibility pathways, mechanical validators now enforce that each narrative states the engine’s actual mechanism — for example, that a child ineligible for CHIP because they qualify for Medicaid is never described as failing a CHIP income test. The failure-audit judge reads these grounded narratives together with engine facts and must return a per-model diagnosis naming the specific rule or computation the model missed. It is barred from re-adjudicating the reference in prose: hedged or reference-adjudicating verdicts are mechanically rejected and re-judged, and genuine reference doubt is routed to a structured flag that triggers the manual adjudication described above.

After freezing the snapshot and completing response-contract repairs, we annotated the 7,545 rows whose legacy threshold score is below 1 and their scenario-output cases. Table 5 reports the final classifications for that legacy-threshold audit universe. It contains 7,541 of the snapshot’s

7,541 exact-match misses and 4 exact hits. The 1,843 other rows with bounded score below 100 fall outside the selection and have no audit annotation. Each annotated row on a scored output resolves to one of two final classes: `llm_error` (a substantive miscalculation, with a per-model diagnosis naming the rule or computation the model missed) or `parse_contract_failure` (a missing or unparsable answer, classified deterministically). Two further classes describe outputs the snapshot does not score, and each such output is removed from scoring for every model, so no model gains or loses from it. `prompt_ambiguity` marks 24 outputs whose reference depends on an input the prompt never lists, so a careful reader could take the stated facts either way; for each, the reference was recomputed under the other reading with the engine version that produced it, and it moved. `reference_engine_defect` marks the 28 outputs where that engine misapplies the law on stated facts. Together they remove 52 outputs in 36 households, and of the 1,793 annotated rows on those outputs, 1743 carry one of these classes as description rather than as a scored class and the rest are answers that never parsed. A scored reference also follows from law published before the references were frozen on 2026-07-03. USDA published the FY2027 SNAP cost-of-living figures on 2026-08-21, after the freeze and after most rows were answered, so grading any model on them would score it against law it could not have seen; SNAP's October to December months therefore hold the FY2026 figures, and the 23 references the publication conventions regenerate, described above, apply the same rule to every projected amount. Every model is therefore scored on 1,932 of its 1,984 requested outputs. The exclusion record (`reference_exclusions.json`) travels with the frozen run and the manifest pins it. No scored row remains classified as a prompt ambiguity, unresolved reference issue, or needs-review item, and zero annotated rows carry a standing reference-suspect flag. The audit pipeline is model-assisted, with developer adjudication of each reference-suspect or prompt-ambiguity flag rather than independent external validation; this snapshot records 63 cases of developer adjudication, kept beside the annotations with the judge's verdict and the reasoning. Three judge models produced the verdicts, and all are board rows: GPT-5.6 Sol, called through the Codex CLI, judged the cases whose verdicts predate the September 2026 additions; Claude Opus 5, called through the Claude Code CLI, judged the cases the September 5 additions joined (183 cases); and Claude Opus 5.5 judged the cases the September 22 additions joined or the regenerated references changed (171 of 668 cases), its own row's cases included. Each verdict is recorded in a per-case provenance sidecar in the audit tree. Judge verdicts classify misses after scoring and change no score, and every reference flag a judge raised is settled by the developer adjudication above. The unlisted inputs behind the unlisted-input exclusions are the SSI disability criterion (`meets_ssi_disability_criteria`), false for every person in the certified June 2026 microdata build; months of SSDI receipt (`months_receiving_social_security_disability`), never carried by the data and integer-typed, which the prompt builder does not expose; weekly hours worked (`weekly_hours_worked_before_lsr`), which the engine defaulted to 40 where the prompt treats unlisted numbers as zero; whether survivor benefits other than Social Security are a pension or annuity, which the prompt does not say and which decides both their taxability and whether they count toward Wisconsin's homestead credit income; who paid for the coverage behind listed employment disability benefits, which decides whether they are taxable; whether listed home mortgage interest is on the home the household occupies, which decides whether SNAP counts it as a shelter cost; whether an adult tax dependent is the claiming filers' child,

which decides whose income counts in that dependent’s Medicaid household; and whether listed interest comes from a Massachusetts bank deposit, which moves only an output already recorded as an engine-defect exclusion. One definition was also open: the output definition of federal income tax before refundable credits does not say whether it includes the net investment income tax, which the reference adds and Form 1040 reports after its line for tax after nonrefundable credits, so the outputs that tax moves are excluded as well. The prompt lists a general is disabled fact and never a program-specific determination (Section 4.2). The exclusion set was derived mechanically rather than from which models missed: every reference was recomputed under each alternative reading, and under each engine defect’s sandbox fix, and every output that moved under an alternative reading or a defect not fixed upstream was excluded, which is why it includes three SNAP outputs that move through the elderly-or-disabled SNAP rules (in two of those households SSI itself stays at zero; in the third an aged member already receives SSI and the disabled member would add a second benefit).

Source: [Article Notebook](#)

Table 4: Discrepancy review across the audit waves.

| Discrepancy class                  | Review outcome                                                                                                                                                                                                                                                  |
|------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| US federal credits                 | Reviewed cases where explanations omitted EITC or refundable CTC; discrepancies reflected model credit-treatment errors, not PolicyEngine errors.                                                                                                               |
| US health coverage (Medicaid/CHIP) | Reviewed threshold-cluster suspects; all dissolved against 42 CFR 435 pathway rules (aged/disabled income tests, tax-dependent household composition, CHIP’s exclusion of Medicaid-eligible children) confirmed in engine traces.                               |
| US SNAP, SSI, and school meals     | Reviewed false-zero suspects; references confirmed against categorical-eligibility, aged/blind/disabled, and universal-free-meal rules. One SNAP cell traced to a scenario encoding that omits the disability its prompt implies (disclosed under Limitations). |
| US state and local income tax      | Reviewed state-rule and bracket suspects; most applied wrong schedules or missed conformity elections and forgiveness programs. One genuine defect: New Jersey’s filing-threshold floor was bypassed on the benchmarked path (fixed, policyengine-us #8845).    |

Table 4: Discrepancy review across the audit waves.

| Discrepancy class                           | Review outcome                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|---------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| US Head Start eligibility                   | Two flags reflected a genuine age-window defect (fixed, policyengine-us #8846); references regenerated with the fix.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| US payroll and overtime                     | Reviewed during development; upstream data fixes and prompt clarifications were applied before the frozen snapshot.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| US federal income tax (September 2026 wave) | Engine defects confirmed: the traditional IRA deduction ignores the compensation limit and the active-participant phase-out (IRC 219(b), 219(g)); estate income never enters gross income; elective deferrals count as EITC and refundable-CTC earned income. Affected outputs excluded. Capital gain distributions reported without Schedule D never reached AGI, a defect fixed upstream after the freeze; the references it alone moves are regenerated with the fix. The taxability of employment disability benefits and whether the output includes the net investment income tax excluded as unlisted inputs. The educator-expense flag dissolved against the \$350 2026 cap. |

Table 4: Discrepancy review across the audit waves.

| Discrepancy class                                     | Review outcome                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|-------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| US state income tax and credits (September 2026 wave) | Engine defects confirmed in New Jersey worker contributions, Wisconsin’s elected retirement exclusion, Wisconsin’s capital gain exclusion (which leaves out capital gain distributions), Idaho’s health-premium subtraction, California itemized-deduction conformity, and Massachusetts dividends taxed without the short-term capital loss offset; affected outputs excluded. The CalEITC’s missing second lookup at adjusted gross income and New York’s renter cap were fixed upstream after the freeze; the affected references are regenerated with the fixes. Wisconsin homestead income turned on the unlisted type of survivor benefits (excluded as an unlisted input). The New Jersey EITC, New York real property tax credit and Idaho child tax credit flags dissolved against the statute (Idaho’s credit expired after 2025). |

Table 4: Discrepancy review across the audit waves.

| Discrepancy class                                  | Review outcome                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
|----------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| US SNAP, SSI and CHIP (September 2026 wave)        | Whether listed mortgage interest is on the occupied home, unlisted hours worked, survivor-benefit taxability, and whether an adult tax dependent is the filers' child (three Medicaid outputs) excluded as unlisted inputs. SNAP October-December references regenerated with the FY2026 figures, the last published before the freeze, and with the engine's SNAP rounding defects fixed as policyengine-us fixed them after the freeze (cents in the allotment, an unrounded minimum benefit, floored net income, unrounded income standards). One Michigan SNAP reference regenerated for all twelve months with the fix for the SNAP child support option, whose engine values carried the opposite meaning (Michigan deducts child support from net income, and the engine excluded it from gross income; fixed upstream in policyengine-us #9586 on 2026-09-24). The heat-and-eat utility allowance granted without an elderly or disabled member, which P.L. 119-21 ended, excluded as an engine defect. The CHIP flag dissolved: employer coverage bars CHIP. |
| US projected 2026 parameters (September 2026 wave) | No judge flag; a recomputation with every projected 2026 value held found references resting on price-index projections. Wisconsin, Minnesota, Michigan and Missouri amounts and Maryland's withholding allowance had been published before the freeze; California's indexed amounts, the IRS sales tax tables, Idaho's threshold and Maryland's return deduction had not. Scored references regenerated with the published or last-published amounts; Wisconsin's reach only the corrected values of Wisconsin outputs excluded for engine defects.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |

Source: [Article Notebook](#)

Table 5: Annotation outcomes for the legacy-threshold audit universe in the frozen snapshot.

| Country | Legacy-threshold rows audited | LLM response errors | Parse-contract failures | Unresolved prompt/reference rows | Missing row annotations | Missing case notes |
|---------|-------------------------------|---------------------|-------------------------|----------------------------------|-------------------------|--------------------|
| US      | 7545                          | 6888                | 657                     | 0                                | 0                       | 0                  |
| Total   | 7545                          | 6888                | 657                     | 0                                | 0                       | 0                  |

Source: [Article Notebook](#)

## 5 Results

### 5.1 United States leaderboard

The US leaderboard for the frozen manuscript snapshot is shown in Table 6. The top three models in that snapshot are GPT-6 Sol (94.3 exact, 95% CI 92.2–96.1; 95.3 within-1%), Claude Opus 5.5 (93.0 exact, 95% CI 90.6–95.1; 94.4 within-1%), and GPT-5.6 Sol (92.7 exact, 95% CI 90.3–94.9; 93.8 within-1%).

Source: [Article Notebook](#)

Table 6: Top US benchmark models in the frozen manuscript snapshot, ranked by the headline exact-match rate with the within-1% near-miss companion alongside. Exact 95% CIs and rank ranges are household-resampling bootstrap intervals (see Uncertainty).

|   | Model           | Exact % | Exact 95% CI | Exact rank | Within 1% % | Within 10% | Bounded score |
|---|-----------------|---------|--------------|------------|-------------|------------|---------------|
| 0 | GPT-6 Sol       | 94.3    | 92.2–96.1    | 1–2        | 95.3        | 96.9       | 97.9          |
| 1 | Claude Opus 5.5 | 93.0    | 90.6–95.1    | 1–4        | 94.4        | 96.7       | 97.4          |
| 2 | GPT-5.6 Sol     | 92.7    | 90.3–94.9    | 1–5        | 93.8        | 95.4       | 96.6          |
| 3 | GPT-6 Luna      | 91.4    | 88.9–93.8    | 2–7        | 93.4        | 96.3       | 97.2          |
| 4 | GPT-6 Astra     | 91.2    | 88.8–93.4    | 2–7        | 92.4        | 94.4       | 97.0          |

Source: [Article Notebook](#)

## 5.2 Uncertainty

Each leaderboard score is a mean over a sample of 100 households, so the rankings carry sampling uncertainty. We quantify it with a household-level bootstrap: the 100 households are resampled with replacement 10,000 times, and each model’s population-weighted exact-match rate is recomputed on every resample. The resampling unit is the household, carried with its full output vector, so within-household correlations across outputs are preserved. The same resampled set is used for every model in each replicate (a paired bootstrap), so the rank ranges in Table 6 reflect genuine ordering uncertainty rather than independent per-model noise.

Sampling uncertainty is substantial at this sample size. In the frozen snapshot, 6 of the other 41 models have an exact-match 95% interval that overlaps the leader’s. Small differences in the leaderboard ordering should therefore be read as ties rather than as a strict ranking, and the rank-range column makes the set of plausible positions explicit. These intervals are a manuscript artifact: the public leaderboard and app report point estimates only, because the benchmark sample, not run-to-run model variation, is the dominant source of uncertainty here.

## 5.3 Cost and latency

Accuracy is one axis; the price and speed of a no-tools answer are another. Table 7 reports each model’s estimated cost to compute one household’s full output vector and its median per-household latency, next to the headline exact rate. Per-household cost spans \\\\$0.002 (GPT-5.4 nano) to \\\\$0.541 (Claude Fable 5), and median per-household latency 3 s (Gemini 3.5 Flash-Lite) to 1023 s (Kimi K3). Neither dimension cleanly tracks accuracy: the top-scoring model, GPT-6 Sol at 94.3% exact, sits mid-pack at \\\\$0.027 per household and 23 s.

Source: [Article Notebook](#)

Table 7: Estimated cost and median latency per household for the frozen US snapshot, alongside the headline exact-match rate. Each row uses its recorded per-call cost: reconstructed from token counts at the list price configured at request time, or the provider-reported charge where no reconstruction was available. List-price overrides apply at request time, not retroactively to recorded costs. Models without per-call costs use the frozen release-metadata cost. Latency is the median across households of each household’s summed request time. A dash marks usage or wall-clock latency that was not measured per call.

|   | Model           | Exact (%) | Cost / household | Latency (median) |
|---|-----------------|-----------|------------------|------------------|
| 0 | GPT-6 Sol       | 94.3      | \\\$0.027        | 23 s             |
| 1 | Claude Opus 5.5 | 93.0      | \\\$0.067        | 24 s             |
| 2 | GPT-5.6 Sol     | 92.7      | \\\$0.090        | 36 s             |
| 3 | GPT-6 Luna      | 91.4      | \\\$0.002        | 29 s             |

Table 7: Estimated cost and median latency per household for the frozen US snapshot, alongside the headline exact-match rate. Each row uses its recorded per-call cost: reconstructed from token counts at the list price configured at request time, or the provider-reported charge where no reconstruction was available. List-price overrides apply at request time, not retroactively to recorded costs. Models without per-call costs use the frozen release-metadata cost. Latency is the median across households of each household’s summed request time. A dash marks usage or wall-clock latency that was not measured per call.

|    | Model                         | Exact (%) | Cost / household | Latency (median) |
|----|-------------------------------|-----------|------------------|------------------|
| 4  | GPT-6 Astra                   | 91.2      | \\$0.126         | 32 s             |
| 5  | Claude Fable 5.1              | 90.2      | \\$0.257         | 49 s             |
| 6  | Kimi K3                       | 89.9      | \\$0.470         | 1023 s           |
| 7  | GPT-5.6 Luna                  | 88.5      | \\$0.019         | 14 s             |
| 8  | Inkling                       | 87.9      | \\$0.091         | 250 s            |
| 9  | GLM-5.3-Flash (preview)       | 87.9      | —                | 242 s            |
| 10 | GPT-5.5                       | 87.5      | \\$0.262         | —                |
| 11 | GPT-5.6 Terra                 | 87.0      | \\$0.039         | 19 s             |
| 12 | Grok 4.6                      | 86.1      | \\$0.087         | 219 s            |
| 13 | Gemini 3.8 Flash              | 85.4      | \\$0.068         | 56 s             |
| 14 | Grok 4.5                      | 84.4      | \\$0.073         | 112 s            |
| 15 | Claude Fable 5                | 83.5      | \\$0.541         | 98 s             |
| 16 | Claude Opus 5                 | 83.4      | \\$0.067         | 16 s             |
| 17 | DeepSeek V4 Flash 0731        | 83.1      | \\$0.003         | 195 s            |
| 18 | Gemini 3.7 Flash              | 82.9      | \\$0.014         | 9 s              |
| 19 | DeepSeek V4 Pro 0813          | 82.8      | \\$0.199         | 808 s            |
| 20 | Gemini 3.6 Flash              | 82.5      | \\$0.055         | 26 s             |
| 21 | Gemini 3.5 Flash-Lite         | 81.7      | \\$0.003         | 3 s              |
| 22 | Gemini 3.1 Pro Preview        | 81.4      | \\$0.085         | 45 s             |
| 23 | GLM-5.3                       | 80.8      | \\$0.169         | 134 s            |
| 24 | Claude Opus 4.7               | 80.7      | \\$0.293         | 53 s             |
| 25 | Gemini 3 Flash Preview        | 80.7      | \\$0.048         | 58 s             |
| 26 | Grok 4.3                      | 80.4      | \\$0.015         | 22 s             |
| 27 | Claude Sonnet 4.6             | 80.1      | \\$0.186         | 114 s            |
| 28 | Gemini 3.5 Flash              | 79.8      | \\$0.056         | 24 s             |
| 29 | Grok Build 0.1                | 79.8      | \\$0.045         | 135 s            |
| 30 | Gemini 3.1 Flash Lite Preview | 79.6      | \\$0.002         | 3 s              |
| 31 | DeepSeek V4 Pro               | 79.6      | \\$0.011         | 170 s            |
| 32 | Qwen3.7-max                   | 77.2      | \\$0.139         | —                |
| 33 | GLM-5.2                       | 76.5      | \\$0.082         | —                |
| 34 | Claude Opus 4.8               | 76.0      | \\$0.255         | 55 s             |
| 35 | MiniMax M3                    | 75.3      | \\$0.007         | —                |

Table 7: Estimated cost and median latency per household for the frozen US snapshot, alongside the headline exact-match rate. Each row uses its recorded per-call cost: reconstructed from token counts at the list price configured at request time, or the provider-reported charge where no reconstruction was available. List-price overrides apply at request time, not retroactively to recorded costs. Models without per-call costs use the frozen release-metadata cost. Latency is the median across households of each household’s summed request time. A dash marks usage or wall-clock latency that was not measured per call.

|    | Model            | Exact (%) | Cost / household | Latency (median) |
|----|------------------|-----------|------------------|------------------|
| 36 | Qwen3.8-Max      | 74.7      | \\$0.224         | 861 s            |
| 37 | Claude Haiku 4.5 | 74.5      | \\$0.055         | 48 s             |
| 38 | GPT-5.4 mini     | 73.2      | \\$0.006         | 5 s              |
| 39 | Claude Sonnet 5  | 72.4      | \\$0.172         | 64 s             |
| 40 | Kimi K2.6        | 67.0      | \\$0.391         | —                |
| 41 | GPT-5.4 nano     | 64.8      | \\$0.002         | 6 s              |

Source: [Article Notebook](#)

Cost preserves each call’s recorded total: a request-time reconstruction from token counts where one was available, and the provider-reported charge otherwise. Configured list-price overrides do not reprice recorded calls. Rows without per-call costs use the frozen release-metadata cost. Latency is wall-clock request time as observed by the harness, summed within each household and taken as the median across the 100 households. The per-call timer spans any provider-side retry or rate-limit backoff, so the median is reported in preference to the mean, which a few throttled calls inflate. Latency is not comparable across serving regimes: Claude-family models issue chunked serial calls that lengthen wall-clock time, several models ran through batch or asynchronous dispatch whose logged times reflect submission rather than generation (shown as a dash), and Claude Fable 5’s usage and latency are recorded at the run level rather than per call. Both cost and latency are properties of this run’s providers, prompts, and serving configurations at evaluation time rather than fixed model constants.

## 5.4 Simple baselines

Simple baselines help interpret score levels in a zero-heavy benchmark. Table 8 reports an always-zero response and a median-reference-by-output response on the same frozen household sample, scored with the household-impact weighting. These are not model competitors; they show how much of each metric can be earned without household-specific policy calculation. The always-zero baseline scores 69.5% on the headline exact-match rate — below the 86% of outputs that are exact zeros because the weighting concentrates mass in high-dollar output

groups that are rarely zero (Section 3.1). The top model clears the baseline by 24.8 points, and the within-1% companion shows the same separation with near-miss tolerance.

Source: [Article Notebook](#)

Table 8: Simple baseline scores on the frozen manuscript snapshot, household-impact-weighted. The always-zero exact rate sits below the unweighted zero share because the weighting concentrates on high-dollar output groups that are rarely zero.

|   | Baseline                   | Exact % | Within 1% % | Bounded score |
|---|----------------------------|---------|-------------|---------------|
| 0 | Always zero                | 69.5    | 69.5        | 69.5          |
| 1 | Median reference by output | 63.5    | 63.5        | 66.5          |

Source: [Article Notebook](#)

## 5.5 Sensitivity to benchmark view

The primary leaderboard gives each household equal weight. Table 9 reports several alternative views, including the older equal-output-group score. The top model is stable across several views, but lower ranks move across sensitivity views. Scores are therefore best read with the weighting choice visible rather than collapsed into a single ordering.

Source: [Article Notebook](#)

Table 9: Leaderboard sensitivity to alternative scoring views.

|   | View                          | Rank 1                 | Rank 2                 | Rank 3                  |
|---|-------------------------------|------------------------|------------------------|-------------------------|
| 0 | Exact-match headline          | GPT-6 Sol (94.3)       | Claude Opus 5.5 (93.0) | GPT-5.6 Sol (92.7)      |
| 1 | Within-1% companion           | GPT-6 Sol (95.3)       | Claude Opus 5.5 (94.4) | GPT-5.6 Sol (93.8)      |
| 2 | Bounded score                 | GPT-6 Sol (97.9)       | Claude Opus 5.5 (97.4) | GPT-6 Luna (97.2)       |
| 3 | Equal-output-group score      | Claude Opus 5.5 (98.1) | GPT-6 Astra (98.1)     | GPT-6 Sol (98.1)        |
| 4 | Amount outputs only           | Claude Opus 5.5 (97.7) | GPT-6 Astra (97.6)     | Claude Fable 5.1 (97.4) |
| 5 | Positive reference cases only | GPT-6 Astra (89.5)     | GPT-5.5 (89.1)         | Gemini 3.5 Flash (86.6) |
| 6 | Zero reference cases only     | Claude Opus 5.5 (99.5) | GPT-6 Sol (99.5)       | Gemini 3.8 Flash (99.2) |

Source: [Article Notebook](#)

The bounded score is a global-output-weighted average of row scores. Amount outputs use  $\max(0, 1 - |\text{pred} - \text{ref}| / |\text{ref}|)$  when the reference is nonzero and  $1\{\text{pred} = 0\}$  when the reference is zero. Binary outputs use exact 0/1 matching. The exact-match rate uses the same weighting and aggregation but replaces the amount row score with the headline indicator:

$1\{|\text{pred} - \text{ref}| \geq 1\}$  (one currency unit). Each output group’s default household-impact weight is computed from a full weighting population, not from the 100 benchmark households: the certified PolicyEngine US populace dataset. For each source household, the share is  $|\text{ref}_{ij}| / \max(|\text{household\_net\_income}_i|, \sum_k |\text{ref}_{ik}|)$ ; shares are averaged with calibrated household weights and renormalized so output weights sum to one. The  $\max(\dots)$  denominator anchors per-household shares to net income when net income is the dominant flow and falls back to the gross tax-benefit flow only when programs cancel each other out, so a \$1 benefit to a high-earner contributes essentially zero weight and per-household shares never exceed one. Booleans carry weight through PolicyEngine’s paired per-capita value (for example `medicaid_value`), so eligibility calls are graded as accuracy but weighted by the dollar stake. Table 10 compares the bounded-score ranking against two opt-in alternatives shown for transparency: equal weighting (every output the same) and budget-weighted (each output’s share of total absolute reference dollars in the same full weighting population).

Source: [Article Notebook](#)

Table 10: Top country ranks under the bounded secondary metric across three weightings (household, equal, budget).

|   | Weighting                    | Rank 1                 | Rank 2                 | Rank 3             |
|---|------------------------------|------------------------|------------------------|--------------------|
| 0 | Household-weighted (default) | GPT-6 Sol (97.9)       | Claude Opus 5.5 (97.4) | GPT-6 Luna (97.2)  |
| 1 | Equal weights                | Claude Opus 5.5 (98.2) | GPT-6 Sol (98.1)       | GPT-6 Astra (97.9) |
| 2 | Budget-weighted              | GPT-6 Sol (97.5)       | GPT-6 Luna (96.5)      | GPT-6 Astra (96.5) |

Source: [Article Notebook](#)

## 5.6 Hardest benchmark targets

The lowest-scoring US variables are multi-step tax quantities and sparse positive-dollar benefits. As shown in Table 11, federal income tax before refundable credits, state income tax before refundable credits, state refundable credits, SNAP, and federal refundable credits are the hardest US outputs by bounded score in the frozen manuscript snapshot.

Source: [Article Notebook](#)

Table 11: Hardest US variables by bounded score (the secondary tracking metric, used here because it orders the hardest variables differently from the hit rates and separates them more finely; on their positive-reference cases the hit rates collapse toward single digits — see the zero-versus-positive table).

|    | Variable                                     | Score | Exact | Within 1% | Within 10% |
|----|----------------------------------------------|-------|-------|-----------|------------|
| 0  | federal_income_tax_before_refundable_credits | 81.3  | 58.5  | 60.6      | 69.2       |
| 15 | state_income_tax_before_refundable_credits   | 83.5  | 63.7  | 66.1      | 75.2       |
| 16 | state_refundable_credits                     | 85.1  | 81.7  | 81.8      | 82.7       |
| 13 | snap                                         | 90.3  | 85.0  | 86.1      | 88.8       |
| 1  | federal_refundable_credits                   | 91.3  | 86.1  | 87.2      | 89.6       |

Source: [Article Notebook](#)

## 5.7 Zero and positive cases

Overall hit rates can overstate performance on sparse programs because correct zeros are common. Table 12 therefore reports within-10% performance for all cases, positive-reference cases, and zero-reference cases. The gap is large for several outputs: models often identify that a program or tax does not apply, but miss the amount when the reference value is positive. Figure 1 plots the same split by output group.

Source: [Article Notebook](#)

Table 12: US within-10% performance by zero versus positive reference cases.

|   | Variable                                     | All cases | Positive cases | Zero cases |
|---|----------------------------------------------|-----------|----------------|------------|
| 0 | federal_income_tax_before_refundable_credits | 70.2      | 40.3           | 92.9       |
| 1 | state_income_tax_before_refundable_credits   | 76.0      | 46.3           | 91.7       |
| 2 | state_refundable_credits                     | 83.8      | 16.3           | 98.9       |
| 3 | snap                                         | 89.2      | 37.6           | 96.8       |
| 4 | federal_refundable_credits                   | 91.0      | 55.7           | 95.5       |
| 5 | payroll_tax                                  | 91.5      | 87.5           | 98.3       |
| 6 | person_medicaid_eligible                     | 92.1      | 81.8           | 95.2       |
| 7 | person_chip_eligible                         | 96.4      | —              | 96.4       |

Source: [Article Notebook](#)



Figure 1: Output-level performance on zero-reference and positive-reference cases.

## 5.8 Open-weight time lag

Source: [Article Notebook](#)

Source: [Article Notebook](#)

The leaderboard also measures time. Following the lag construction the UK AI Security Institute applied to cyber capabilities (UK AI Security Institute 2026), the open-weight lag is the time between the closed frontier first reaching a score and the leading open-weight model matching it. Two cuts of the frozen roster give the trajectory. Among models released before July 2026, the leading open-weight model is DeepSeek V4 Pro at 79.6% exact, whose weights shipped on April 24, 2026 (DeepSeek 2026); the earliest-released closed model above that score is Gemini 3 Flash Preview, released December 17, 2025 (TechCrunch 2025) — a lag of 128 days, about 4.2 months. On the full roster, Kimi K3 at 89.9% exact trails 6 closed models; the earliest released of them, GPT-5.6 Sol, generally available July 9, 2026 (TechCrunch 2026), puts the lag at 7 days by API availability, or 18 days dated to Kimi K3’s July 27 weight release on Hugging Face (Tech Times 2026).

Both figures are lower bounds in one direction: a closed model outside the evaluated roster that crossed either score earlier would lengthen the corresponding lag. They also carry the uncertainty of any single-benchmark, single-domain estimate — the AISI analysis reports substantially longer lags on cyber tasks — and a lag measured at a moving frontier compresses mechanically when a strong open-weight release lands days after a frontier release.

## 6 Failure modes

The benchmark surfaces a few recurring failure patterns.

First, models miss positive tax and benefit quantities more often than zero cases. The lowest-scoring outputs by bounded score (Table 11) require the model to choose the right income concepts, exclusions, program thresholds, and sequencing before applying any final subtraction. Positive benefit cases remain difficult, so the result should not be read as a general claim that benefits are easy.

Second, joint accuracy across interacting components can be lower than marginal accuracy on either component. Table 13 shows within-10% accuracy for `federal_refundable_credits`, `state_refundable_credits`, and the conjunction of both within the same household. The joint hit rate can be no higher than either marginal and is strictly lower than both for every model except Claude Fable 5.1, Claude Opus 5.5, GPT-6 Astra, GPT-6 Sol, and GPT-5.6 Sol. Leaderboard scores that average across outputs therefore overstate how often a model gets a single household’s federal/state credit allocation jointly correct.

Source: [Article Notebook](#)

Table 13: US within-10% accuracy on federal vs state refundable credits and the household-level joint.

|    | Model                         | Federal within 10% | State within 10% | Joint within 10% |
|----|-------------------------------|--------------------|------------------|------------------|
| 1  | Claude Fable 5.1              | 100.0              | 93.8             | 93.8             |
| 6  | Claude Opus 5.5               | 100.0              | 91.8             | 91.8             |
| 28 | GPT-6 Astra                   | 100.0              | 89.7             | 89.7             |
| 30 | GPT-6 Sol                     | 99.0               | 88.7             | 88.7             |
| 26 | GPT-5.6 Sol                   | 99.0               | 87.6             | 87.6             |
| 29 | GPT-6 Luna                    | 99.0               | 88.7             | 87.6             |
| 19 | Gemini 3.8 Flash              | 95.9               | 87.6             | 85.6             |
| 24 | GPT-5.5                       | 95.9               | 87.6             | 85.6             |
| 33 | Grok 4.6                      | 96.9               | 85.6             | 84.5             |
| 17 | Gemini 3.6 Flash              | 95.9               | 84.5             | 83.5             |
| 11 | DeepSeek V4 Pro 0813          | 94.8               | 84.5             | 82.5             |
| 18 | Gemini 3.7 Flash              | 91.8               | 85.6             | 81.4             |
| 35 | Inkling                       | 94.8               | 83.5             | 81.4             |
| 12 | Gemini 3 Flash Preview        | 91.8               | 85.6             | 80.4             |
| 27 | GPT-5.6 Terra                 | 95.9               | 82.5             | 80.4             |
| 37 | Kimi K3                       | 94.8               | 82.5             | 80.4             |
| 0  | Claude Fable 5                | 92.8               | 84.5             | 79.4             |
| 10 | DeepSeek V4 Pro               | 91.8               | 83.5             | 79.4             |
| 14 | Gemini 3.1 Pro Preview        | 91.8               | 83.5             | 79.4             |
| 25 | GPT-5.6 Luna                  | 94.8               | 81.4             | 79.4             |
| 32 | Grok 4.5                      | 93.8               | 81.4             | 79.4             |
| 39 | GLM-5.3-Flash (preview)       | 91.8               | 83.5             | 78.4             |
| 15 | Gemini 3.5 Flash              | 91.8               | 81.4             | 77.3             |
| 5  | Claude Opus 5                 | 88.7               | 83.5             | 76.3             |
| 7  | Claude Sonnet 4.6             | 93.8               | 81.4             | 76.3             |
| 9  | DeepSeek V4 Flash 0731        | 89.7               | 82.5             | 75.3             |
| 13 | Gemini 3.1 Flash Lite Preview | 86.6               | 81.4             | 75.3             |
| 16 | Gemini 3.5 Flash-Lite         | 87.6               | 81.4             | 75.3             |
| 23 | GPT-5.4 nano                  | 87.6               | 81.4             | 75.3             |
| 34 | Grok Build 0.1                | 90.7               | 81.4             | 75.3             |
| 4  | Claude Opus 4.8               | 85.6               | 81.4             | 73.2             |
| 31 | Grok 4.3                      | 85.6               | 81.4             | 73.2             |
| 21 | GLM-5.3                       | 87.6               | 79.4             | 72.2             |
| 22 | GPT-5.4 mini                  | 82.5               | 81.4             | 72.2             |
| 8  | Claude Sonnet 5               | 81.4               | 81.4             | 71.1             |
| 38 | MiniMax M3                    | 81.4               | 81.4             | 70.1             |
| 3  | Claude Opus 4.7               | 86.6               | 79.4             | 69.1             |
| 20 | GLM-5.2                       | 81.4               | 74.2             | 68.0             |

Table 13: US within-10% accuracy on federal vs state refundable credits and the household-level joint.

|    | Model            | Federal within 10% | State within 10% | Joint within 10% |
|----|------------------|--------------------|------------------|------------------|
| 2  | Claude Haiku 4.5 | 76.3               | 81.4             | 67.0             |
| 40 | Qwen3.7-max      | 79.4               | 81.4             | 66.0             |
| 41 | Qwen3.8-Max      | 68.0               | 81.4             | 58.8             |
| 36 | Kimi K2.6        | 46.4               | 50.5             | 30.9             |

Source: [Article Notebook](#)

Third, structured-output reliability is part of the benchmark contract. Missing or unparseable numeric values are not dropped. Appendix A documents parser recovery and the supervised runner’s bounded retry rounds. The 657 cells (0.8%) that still returned no parsable value after retries remain in the canonical files and score zero as parse-contract failures. Appendix A gives the per-model breakdown.

## 7 Limitations

PolicyBench is not a substitute for a production tax-and-benefit calculator. Several caveats matter:

- output-contract reliability required a repair workflow; raw provider responses are retained within the canonical prediction files wherever the transport exposed them (no raw payload is retained for Claude Fable 5 (1,984 rows), GLM-5.3 (78 rows), and Kimi K3 (64 rows)), and cells that never parsed score zero
- request shape is not identical across models: 10 of 42 models answered the prompt template over subsets of the requested outputs rather than the whole scenario (Table 2), an accommodation that, if it moves scores at all, should favor those models; later additions must answer the whole-scenario request or are listed as not scorable
- zero-heavy outputs require separate positive-case interpretation
- the choice of headline metric (exact match), near-miss companion (within-1%), and secondary tracking metric (bounded score) is one benchmark view; all three are reported, but downstream evaluations should also consider error-magnitude metrics (mean absolute error, mean absolute percentage error) where appropriate
- the public scenario explorer exposes the current test set and reference outputs, so open-set leakage is a prominent limitation rather than a minor implementation detail
- the 100-household manuscript snapshot should be treated as a preview until larger frozen runs are published

- the annotations cover the 7,545-row legacy-threshold audit universe, including 7,541 exact-match misses and 4 exact hits, but not 1,843 additional rows with bounded score below 100; the review is developer-adjudicated and not an independent validation of each reference value
- one frozen scenario encodes household inputs that do not match its prompt narrative (a South Carolina household whose prompt implies a qualifying disability the inputs omit), and two input-semantics conventions (disability, which the prompt states as one general fact where each program applies its own determination, and adult tax-dependent relationships) are under-specified in ways a scheduled scenario refresh will address (Section 4.2); outputs whose reference moves under the other reading of either convention are excluded, while the mis-encoded scenario’s cells score models against its encoded inputs
- the reference engine carries known defects: 28 outputs, recorded as engine-defect exclusions, rest on 11 root causes in policyengine-us 1.755.4 that were not fixed upstream and are excluded until a reference version built on a fixed engine restores them; the 15 references regenerated for defects fixed upstream apply sandbox versions of those fixes to the frozen engine version, not the upstream releases themselves, and the audit that found them reviews only the cases a new model joins, so defects that no reviewed case exposes can remain
- the references were frozen on 2026-07-03 for tax year 2026, so amounts published later reach the benchmark only through a stated convention: where policyengine-us projected a 2026 amount, the reference takes the amount published before the freeze or, failing that, the last one published (SNAP’s FY2026 figures rather than USDA’s FY2027 figures of 2026-08-21; California’s 2025 indexed amounts; the IRS’s 2025 sales tax tables), so a model that knew later law is scored against the earlier law, and the projection audit sourced only the amounts whose projection moved an output
- benchmark success should not be interpreted as policy-advice readiness

The current paper is therefore an evaluation of model performance under a specific structured-output benchmark, not a general certification of tax or benefit competence.

## 8 Conclusion

PolicyBench shows a consistent pattern across the US benchmark. Models often identify non-applicability, but positive tax and benefit amounts remain difficult, especially for multi-step income tax and positive benefit cases. In the frozen manuscript snapshot, GPT-6 Sol is the top-scoring US model.

These results support a narrow conclusion: unaided frontier models still struggle to reproduce selected household-level microsimulation outputs under a structured public benchmark. They do not show that models cannot assist policy analysis, and they do not validate PolicyEngine

outputs as administrative truth. They suggest that future evaluations should separate no-tool estimation, tool-using system design, and reference-output validation more explicitly.

Next steps are to expand country coverage, increase frozen sample sizes, and add protected or rotating evaluation sets so that public rankings are less exposed to open-set leakage. The benchmark should also continue reporting sensitivity views, because scores are useful summaries only when their weighting choices are visible.

# Appendix A: Structured-output audit

The benchmark requires one numeric value and one explanation for every requested output. Parse failures are benchmark failures, not missing data to drop from the denominator. Before publishing the manuscript snapshot, we audited missing numeric values and explanations, extended the parser only where explicit variable-keyed value and explanation fields were recoverable, and retried failing scenarios under the supervised runner’s bounded rounds. Table 14 summarizes that sequence.

Source: [Article Notebook](#)

Table 14: Structured-output parser audit and repair sequence.

| Step                         | Finding                                                                                                                                                                                                                                       |
|------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Initial parse-contract audit | The source run was audited for rows missing a parsed numeric value or non-empty explanation; those rows were not dropped from the denominator.                                                                                                |
| Parser repair                | The parser was extended to recover explicit value and non-empty explanation blocks from nested, escaped, or partially truncated provider JSON without scraping prose numbers.                                                                 |
| Retry policy                 | Later waves ran under a resumable supervised runner that retries failed or timed-out scenarios in bounded rounds; every model’s canonical file covers all 100 households. No separate whole-response retry artifacts exist for this snapshot. |

Table 14: Structured-output parser audit and repair sequence.

| Step                 | Finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
|----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Final parse coverage | 657 of the 81,144 canonical model-output rows (0.8%) have no parsable numeric value after retries. The failures comprise Kimi K2.6 (394), GLM-5.2 (133), GLM-5.3 (71), and Kimi K3 (59). They score zero as parse-contract failures. A further 61 rows (Grok 4.3 (56), Kimi K2.6 (4), and Claude Haiku 4.5 (1)) carry a parsed numeric value but no explanation; they are scored on the numeric value and counted here as explanation-contract violations, so 718 rows in all fall short of the numeric-plus-explanation contract. Every other row has both. |
| Preservation rule    | The snapshot retains raw provider responses in the compressed source-run prediction files wherever the transport exposed them; no raw payload is retained for Claude Fable 5 (1,984 rows), GLM-5.3 (78 rows), and Kimi K3 (64 rows) (batch-adapter rows carry no per-row payload, and those parse failures captured none).                                                                                                                                                                                                                                   |

Source: [Article Notebook](#)

Model waves after June ran under a resumable supervised runner that retries failed or timed-out scenarios in bounded rounds. No separate retry artifacts exist; repairs are folded into each model’s canonical prediction file. After retries, 657 cells (0.8% of canonical rows) still returned no parsable numeric value. They comprise Kimi K2.6 (394), GLM-5.2 (133), GLM-5.3 (71), and Kimi K3 (59). These cells remain in the canonical files, with their raw responses preserved wherever the transport exposed one (no raw payload is retained for Claude Fable 5 (1,984 rows), GLM-5.3 (78 rows), and Kimi K3 (64 rows)). They score zero as parse-contract failures, and the denominator is never reduced.

## Competing interests

The author leads PolicyEngine, which develops the microsimulation software used to produce the benchmark reference outputs, and is a founder of the Axiom Foundation, a nonprofit developing

deterministic rules-as-code infrastructure. Both organizations’ missions are advanced by the finding that unaided language models do not reliably reproduce rule calculations.

Source: [Article Notebook](#)

Abbood, Auss, Zaiqiao Meng, and Nigel Collier. 2025. “Time to Revisit Exact Match.” Findings of the Association for Computational Linguistics: EMNLP 2025 (Suzhou, China), 11903–26. <https://doi.org/10.18653/v1/2025.findings-emnlp.637>.

Bock, Michael R., Kara Molisee, Zachary Ozer, and Sumit Shah. 2025. “TaxCalcBench: Evaluating Frontier Models on the Tax Calculation Task.” arXiv Preprint arXiv:2507.16126, ahead of print. <https://doi.org/10.48550/arXiv.2507.16126>.

Chen, Zhiyu, Wenhui Chen, Charese Smiley, et al. 2021. “FinQA: A Dataset of Numerical Reasoning over Financial Data.” arXiv Preprint arXiv:2109.00122. <https://arxiv.org/abs/2109.00122>.

Cobbe, Karl, Vineet Kosaraju, Mohammad Bavarian, et al. 2021. “Training Verifiers to Solve Math Word Problems.” arXiv Preprint arXiv:2110.14168. <https://arxiv.org/abs/2110.14168>.

DeepSeek. 2026. DeepSeek-V4 Preview Release. DeepSeek API news, April 24, 2026. <https://api-docs.deepseek.com/news/news260424/>.

Federal Reserve Bank of Atlanta. 2026. Policy Rules Database. Federal Reserve Bank of Atlanta. <https://www.atlantafed.org/what-we-study/workforce-development/advancing-careers-for-low-income-families/policy-rules-database>.

Feenberg, Daniel R., and Elisabeth Coutts. 1993. “An Introduction to the TAXSIM Model.” *Journal of Policy Analysis and Management* 12 (1): 189–94.

Ghenis, Max. 2026. PolicyEngine Powers Rapid Policy Analysis at No 10 Downing Street. PolicyEngine. <https://www.policyengine.org/us/research/policyengine-10-downing-street>.

Ghenis, Max, and Pavel Makarchuk. 2025. PolicyEngine and Atlanta Fed Sign Memorandum of Understanding for Policy Rules Database Validation. PolicyEngine. <https://www.policyengine.org/us/research/policyengine-atlanta-fed-mou-prd>.

Guha, Neel, Julian Nyarko, Daniel E. Ho, et al. 2023. “LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models.” arXiv Preprint arXiv:2308.11462. <https://arxiv.org/abs/2308.11462>.

Hendrycks, Dan, Collin Burns, Saurav Kadavath, et al. 2021. “Measuring Mathematical Problem Solving with the MATH Dataset.” *NeurIPS Datasets and Benchmarks Track*.

<https://arxiv.org/abs/2103.03874>.

- Holzenberger, Nils, Benjamin Van Durme, Sarah Lawsky, and Kyle Richardson. 2021. “Factoring Statutory Reasoning as Language Understanding Challenges.” arXiv Preprint arXiv:2105.07903. <https://arxiv.org/abs/2105.07903>.
- Nava Labs. 2025. Experimenting with AI-Powered Tools in Public Benefits. Case study. <https://www.navapbc.com/case-studies/ai-tools-public-benefits>.
- Nava Labs. 2026. Evaluating a GenAI-Powered Assistive Chatbot for Caseworkers. Case study. <https://www.navapbc.com/case-studies/evaluating-ai-assistive-chatbot-caseworkers>.
- OpenRouter. 2026. “Ox Alpha: API Pricing & Providers.” August. <https://openrouter.ai/stealth/ox-alpha>.
- PolicyEngine. 2024. PolicyEngine Launches State Income Tax Modeling Nationwide. PolicyEngine. <https://www.policyengine.org/us/research/state-tax-model-beta>.
- Shanahan, Catherine, Emma McCarthy, Yan Zhao, et al. 2025. “Performance of LLMs on VITA Test: Potential for AI-Assisted Tax Returns for Low Income Taxpayers.” Artificial Intelligence and Law, ahead of print. <https://doi.org/10.1007/s10506-025-09465-7>.
- Shorten, Connor, Charles Pierse, Thomas Benjamin Smith, et al. 2024. “StructuredRAG: JSON Response Formatting with Large Language Models.” arXiv Preprint arXiv:2408.11061, ahead of print. <https://doi.org/10.48550/arXiv.2408.11061>.
- Sutherland, Holly, and Francesco Figari. 2013. “EUROMOD: The European Union Tax-Benefit Microsimulation Model.” International Journal of Microsimulation 6 (1): 4–26. <https://doi.org/10.34196/ijm.00075>.
- Tech Times. 2026. “Kimi K3 Open Weights Arrive Sunday: Self-Hosting Cuts China Data Risk the API Never Can.” July. <https://www.techtimes.com/articles/321551/20260725/kimi-k3-open-weights-arrive-sunday-self-hosting-cuts-china-data-risk-api-never-can.htm>.
- TechCrunch. 2025. Google Launches Gemini 3 Flash, Makes It the Default Model in the Gemini App. TechCrunch, December 17, 2025. <https://techcrunch.com/2025/12/17/google-launches-gemini-3-flash-makes-it-the-default-model-in-the-gemini-app/>.
- TechCrunch. 2026. OpenAI Launches Its New Family of Models with GPT-5.6. TechCrunch, July 9, 2026. <https://techcrunch.com/2026/07/09/openai-launches-its-new-family-of-models-with-gpt-5-6/>.
- UK AI Security Institute. 2026. How Far Behind the Frontier Are Leading Open-Weight

Models on Cyber? Blog post. <https://www.aisi.gov.uk/blog/how-far-behind-the-frontier-are-leading-open-weight-models-on-cyber>.

Vals AI. 2026. Public Benefits Bench. Benchmark report, in partnership with the Center for Civic Futures and Code for America. <https://www.vals.ai/benchmarks/public-benefits-bench>.

Willison, Simon. 2026. Kimi K3, and What We Can Still Learn from the Pelican Benchmark. Simonwillison.net, July 16, 2026. <https://simonwillison.net/2026/Jul/16/kimi-k3/>.

Woodruff, Nikhil. 2026. Informing Policy Using Micro-Simulations. No10 Innovation Fellowship. <https://fellows.ai.gov.uk/articles/nikhil-woodruff-micro-simulation>.

ZenML LLMOps Database. 2025. Building and Automating Comprehensive LLM Evaluation Framework for SNAP Benefits. Industry writeup. <https://www.zenml.io/llmops-database/building-and-automating-comprehensive-llm-evaluation-framework-for-snap-benefits>.

Zhou, Ruiwen, Wenyue Hua, Liangming Pan, et al. 2025. “RuleArena: A Benchmark for Rule-Guided Reasoning with LLMs in Real-World Scenarios.” Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (Vienna, Austria), 550–72. <https://doi.org/10.18653/v1/2025.acl-long.27>.

Zhu, Fengbin, Wenqiang Lei, Youcheng Huang, et al. 2021. “TAT-QA: A Question Answering Benchmark on a Hybrid of Tabular and Textual Content in Finance.” arXiv Preprint arXiv:2105.07624. <https://arxiv.org/abs/2105.07624>.