

# PolicyBench

Benchmarking no-tool tax-and-benefit estimation in frontier language models

Max Ghenis\*

## Table of contents

|                                                          |    |
|----------------------------------------------------------|----|
| Abstract                                                 | 1  |
| 1 Introduction                                           | 2  |
| 2 Related work                                           | 2  |
| 3 Benchmark design                                       | 3  |
| 3.1 Headline metric: exact match . . . . .               | 4  |
| 3.2 Near-miss companion: within 1% . . . . .             | 4  |
| 3.3 Secondary metric: bounded continuous score . . . . . | 5  |
| 3.4 Frozen snapshot and open-set status . . . . .        | 7  |
| 4 Data and scenario construction                         | 12 |
| 4.1 United States . . . . .                              | 12 |
| 4.2 Reference-output credibility . . . . .               | 14 |
| 5 Results                                                | 17 |
| 5.1 United States leaderboard . . . . .                  | 17 |
| 5.2 Uncertainty . . . . .                                | 18 |
| 5.3 Cost and latency . . . . .                           | 18 |
| 5.4 Simple baselines . . . . .                           | 20 |
| 5.5 Sensitivity to benchmark view . . . . .              | 20 |
| 5.6 Hardest benchmark targets . . . . .                  | 22 |
| 5.7 Zero and positive cases . . . . .                    | 22 |
| 5.8 Open-weight time lag . . . . .                       | 23 |
| 6 Failure modes                                          | 24 |

---

\*PolicyEngine and Axiom Foundation.

|   |                                     |    |
|---|-------------------------------------|----|
| 7 | Limitations                         | 26 |
| 8 | Conclusion                          | 26 |
|   | Appendix A: Structured-output audit | 27 |
|   | Competing interests                 | 28 |

Source: [Article Notebook](#)

## Abstract

PolicyBench evaluates whether frontier language models can estimate household tax and benefit outputs from household facts without tools. This release covers the United States, drawing sampled households from the certified PolicyEngine populace microdataset, scored through policyengine.py 4.16.1. It evaluates 33 frontier models on 100 households across 18 output groups, and reports three complementary scores: an exact-match rate as the headline deployability bar (a prediction counts only if it matches the PolicyEngine reference to the dollar for amounts or to the eligibility flag for booleans), a within-1% hit rate as a near-miss-tolerant companion, and a continuous bounded score that awards partial credit for close answers. The headline is exact match on a household-impact-weighted leaderboard: 84% of US reference outputs are exact zeros, and the weighting concentrates scores in the high-dollar output groups a hedge-to-zero strategy cannot answer, so the top model leads an always-zero baseline of 66.9% weighted exact by 21.8 points (Section 3.1 details the construction). The manuscript snapshot is a 100-household public preview, not a protected held-out leaderboard: the current public scenario explorer exposes prompts and reference outputs, so open-set leakage is a central limitation of any public ranking. In the frozen snapshot used for this manuscript (2026-09-01), the top-scoring US model on exact match is GPT-5.6 Sol at 88.7%, while GPT-5.4 nano is lowest at 62.3%. Within model families, newer is not always stronger: Claude Opus 4.8 scores 72.6% exact, below the 77.4% of the earlier Claude Opus 4.7. Multi-step tax quantities and positive-dollar benefit cases are materially harder than zero cases. Every wrong cell is audited, with reference-suspect flags adjudicated against primary sources and engine traces; the frozen release has no known reference-computation defects, and every one of the 7,840 scored misses traces to a model-side failure (including the 1.0% of cells with a missing or unparsable answer, scored as misses rather than dropped). The live benchmark is available at <https://policybench.org>.

## 1 Introduction

Household tax-and-benefit estimation sits between arithmetic and policy discussion. Each case has numeric labels, but generating those labels requires filing status, household composition,

income concepts, program thresholds, and jurisdiction-specific rules. PolicyBench measures whether models can map household records to those outputs.

Most current language-model evaluations do not test this mapping directly. General math benchmarks emphasize symbolic manipulation and exact final answers (Cobbe et al. 2021; Hendrycks et al. 2021), while generic question-answering benchmarks emphasize recall or instruction following. PolicyBench instead asks whether models can transform household facts into policy outputs without access to tools or simulators.

This matters for public-facing calculators, analyst workflows, and tax-preparation or screening systems. The benchmark is intended to separate verbal policy fluency from household-level quantitative prediction.

## 2 Related work

The most relevant prior work is tax and statutory reasoning. SARA frames statutory interpretation, including tax-relevant rule application, as a language-understanding problem (Holzenberger et al. 2021). LegalBench broadens this view and includes tax-oriented numeric tasks such as `sara_numeric` (Guha et al. 2023). RuleArena evaluates rule-guided reasoning in regulation-style settings (Zhou et al. 2025), and TaxCalcBench evaluates frontier models on tax calculation from structured return-like inputs (Bock et al. 2025). Shanahan et al. (2025) report a similar split on the Volunteer Income Tax Assistance (VITA) test: models do better on tax knowledge questions than on open-ended calculation tasks.

PolicyBench also sits within a broader literature on quantitative reasoning benchmarks. Canonical math evaluations such as GSM8K and MATH normalized exact final-answer scoring for numeric tasks (Cobbe et al. 2021; Hendrycks et al. 2021). Recent work has questioned pure exact-match scoring in numeric QA and temporal reasoning, arguing for error-aware metrics instead (Abbood et al. 2025). PolicyBench treats this as a measurement design choice rather than a forced pick, leading with a household-impact-weighted exact-match rate and reporting near-miss and bounded companions; the metric definitions, and the weighting that keeps the zero-inflated panel from rewarding hedge-to-zero answers, are specified in the benchmark design section.

Finance-domain benchmarks provide another relevant precedent. FinQA and TAT-QA both show that realistic quantitative reasoning often requires combining structured numbers with domain knowledge rather than solving stylized school-math problems (Chen et al. 2021; Zhu et al. 2021). PolicyBench differs in that its reference outputs are generated by executable tax-benefit microsimulation rather than annotated document questions, but the underlying motivation is similar: domain-specific quantitative reasoning deserves dedicated evaluation.

The nearest public-benefits benchmark is the concurrent Public Benefits Bench, built by Vals AI with the Center for Civic Futures and Code for America (Vals AI 2026). It evaluates models on 459 expert-validated SNAP guidance scenarios (230 in the published test set) spanning

all 50 US states plus Guam and the Virgin Islands, and grades free-text answers against expert-written rubric criteria using a judge model that agreed with SNAP policy experts on 80.6% of grading decisions. Where PolicyBench isolates parametric knowledge by withholding tools, Public Benefits Bench varies tool access directly: across models, multi-turn conversation raises pass rates by roughly 7.6 percentage points and web search by about 6.9. The two benchmarks probe complementary failure surfaces — conversational guidance quality for a single program with tools available, versus no-tool numeric estimation across tax and transfer outputs scored deterministically against microsimulation references — and reach consistent conclusions. The best model passes 71.7% of rubric criteria under Public Benefits Bench’s strongest condition, and its authors conclude that “no general-purpose AI model performs well enough to be trusted with SNAP benefits guidance,” matching the deployability gaps PolicyBench reports for unaided calculation.

Finally, PolicyBench depends on two infrastructure literatures that are not themselves large language model (LLM) benchmarks. One is structured-output reliability, since the benchmark relies on multi-output JSON responses and parse coverage rather than free-form prose (Shorten et al. 2024). The other is tax-benefit microsimulation, where systems such as EUROMOD provide the methodological precedent for evaluating policy rules over household microdata (Sutherland and Figari 2013). There is also operational work on applying and evaluating AI systems in public-benefits settings, including caseworker-assist and Supplemental Nutrition Assistance Program (SNAP)-focused evaluations (Nava Labs 2026, 2025; ZenML LLMOps Database 2025). PolicyBench combines these strands into a public cross-model benchmark over household-level tax and benefit outputs.

### 3 Benchmark design

The canonical PolicyBench request asks each model to predict every benchmark output for a household in a whole-scenario response under its structured answer contract. 10 of the 33 models received the documented chunked accommodation and instead answered the same prompt template over one- or three-output subsets, reducing each response’s completion demand while preserving the household-level task.

#### 3.1 Headline metric: exact match

The headline metric is the exact-match rate: the share of requested outputs the model gets right to the dollar (or, for eligibility flags, to the boolean). This is the deployability bar. A tax filer, a benefit estimator, or a caseworker cannot ship “close” — a prediction off by \$50 is no more usable than one off by \$500, because neither matches the bottom line a downstream system or a household needs. For currency amounts, “exact” means within 1 currency unit of the reference value after numeric parsing. For binary outputs, “exact” means the parsed value is exactly the same 0/1 flag as the reference.

Exact match can be a deceptive headline in a zero-inflated panel — but only when the rate is computed unweighted. In the frozen manuscript snapshot, 84% of reference outputs are exact zeros: most sampled households are not eligible for any given program, owe no capital gains tax, do not phase into a refundable credit, and so on. A model that hedges to zero on every requested output earns full exact credit on every zero case, so a naive unweighted exact rate is dominated by those zeros — a hedge-to-zero model scores roughly 84% unweighted, leaving little room above the floor to separate models that actually understand the rules.

The public leaderboard corrects for this by household-impact weighting. Each output group’s weight is its population-average share of the household’s dollar stake, computed as  $|ref| / \max(|household\_net\_income|, \Sigma |ref|)$  per source household and averaged with calibrated household weights over the full populace population (Section 5.5 gives the full construction). Zero-reference rows are not removed — every row carries its output group’s weight — but the weighting concentrates mass in high-dollar output groups, which are rarely zero, so the share of weight on zero-reference rows falls from their 84% row share to 66.9%. An always-zero baseline therefore scores 66.9% weighted exact rather than the 84% it earns unweighted, and the top model clears that baseline by 21.8 points — a margin comparable to within-1%. Across the 33-model roster, the weighted exact rate spreads about as widely as within-1% and ranks models in nearly the same order, so exact match discriminates between rule comprehension and zero-hedging while preserving the to-the-dollar deployability meaning.

The exact rate is aggregated to a household score and then to a country score, with the same population-derived output weights that the within-1% companion and the bounded score use. Equal household weight preserves comparability across scenarios and slices. The leaderboard reports US results on this fixed test set.

### 3.2 Near-miss companion: within 1%

The within-1% hit rate is the near-miss-tolerant companion to the headline: the share of requested outputs the model gets within one percent of the PolicyEngine reference value. For currency amounts, “within 1%” means  $|pred - ref| \leq 0.01 \times |ref|$  when the reference is nonzero and  $|pred| \leq 1$  (the same one-currency-unit absolute tolerance used by exact match) when the reference is zero. Binary outputs are requested as integer 0/1 eligibility flags and scored identically to the exact-match metric.

Within-1% relaxes the exact bar by tolerating rounding and small rate or parameter drift on positive-reference cases, so it credits a model that is consistently close but rarely to-the-dollar. It uses the same per-output household-impact weighting and the same household-equal aggregation as the headline exact rate, and it is reported alongside exact in every leaderboard table. Because both metrics share the weighting, the two rank models almost identically; within-1% sits a few points above exact and is the right view when near-miss accuracy, rather than strict deployability, is the question.

### 3.3 Secondary metric: bounded continuous score

PolicyBench also reports a bounded 0-100 score that awards smooth partial credit for close-but-not-exact amount answers. This is informative for tracking conceptual progress year over year and captures the qualitative gap between a model that is consistently close and one that is consistently far from the reference value. For amount variables, the row-level bounded score is  $\max(0, 1 - |\text{pred} - \text{ref}| / |\text{ref}|)$  when the reference value is nonzero and  $1\{\text{pred} = 0\}$  when the reference value is zero. For binary outputs, it uses exact 0/1 matching. Note that the zero-reference branch is stricter than the headline’s one-currency-unit tolerance — a \$0.50 prediction on a zero-reference amount counts as exact but earns no bounded credit — so the bounded score is not a uniform relaxation of exact match. The exact, within-1%, and within-10% hit rates are reported separately as threshold diagnostics; they are not averaged to create the bounded score.

PolicyBench also tracks mean absolute error and related error metrics. These remain secondary diagnostics under the headline exact-match rate, the within-1% near-miss companion, and the bounded score. The three-metric design follows recent numeric-evaluation literature that questions pure exact-match scoring (Abbood et al. 2025) while keeping exactness as the production bar.

All requested US outputs are annual amounts or annual eligibility indicators for tax year 2026, scored under the tolerances defined above. Missing, unparseable, or non-0/1 answers receive zero score for that requested output.

Aggregation proceeds in three steps. First, each household-output prediction receives a 0-100 score. Second, each household receives one model score: requested output rows are weighted by output-group weights constructed from the full weighting population, then renormalized within the household over outputs that are actually requested for that household. Person-level coverage flags are scored at the person row; their output-group weight is split across the relevant people in the household and is based on PolicyEngine value proxies where available. Third, the country score averages those household scores with equal household weight. The leaderboard reports US results on this fixed test set. Equal-output-group scores and other alternative views are reported as sensitivity checks.

The benchmark requires each model response to include numeric answers and one explanation per requested output. Both the exact-match rate and the bounded score use only the numeric answers. Explanations are retained for scenario exploration and qualitative error analysis; they should not be interpreted as faithful traces of model reasoning.

Because explanations are required, the canonical task measures policy estimation under a public-facing structured-response contract, not isolated arithmetic accuracy. Prompt fairness is part of the benchmark contract. The current release uses one prompt template per country, with no model-specific tuning. Models receive the same household facts and requested outputs. They receive no web or external tool access. The prompt sets unlisted numeric inputs to 0, unlisted boolean or status facts to false, and household characteristics as constant over

the tax-benefit year. Provider-specific differences are limited to structured-output transport, request shape, and the reasoning setup in Table 2. The harness sends no reasoning-control or sampling parameters. It leaves reasoning effort, temperature, and related decoding controls at provider defaults. The frozen GPT-5.5 wave uses provider-default reasoning effort.

The forced `tool_choice` suppresses Claude’s extended thinking in this snapshot ([sensitivity/claude-thinking-2026-08.md](#)). A labeled sensitivity with `tool_choice`: “auto” is published beside the board. Claude Fable 5.1, released September 1, 2026, closes the interaction from the API side: it rejects forced tool use with an error, so its row answers through the JSON transport and reasons at the provider default; a sensitivity run with the tool declared under `tool_choice`: “auto” scores 1.2 points above that row. The next board version moves every model to `tool_choice`: “auto” so each provider’s default reasoning posture can engage.

Two dimensions on which requests were not identical across models are request shape and answer transport (Table 2). 10 of the 33 models could not reliably complete the whole-scenario request because their serving stacks rejected the structured-output call, exhausted completion budgets mid-response, or timed out. They instead answered the same prompt template over subsets of the requested outputs, one to three outputs per request. Subsetting gives each response fewer requested outputs and a fresh completion budget over the same household facts, so the accommodation, if it moves scores at all, should favor the chunked models; their scores are not strictly comparable to whole-scenario scores. This accommodation is closed going forward: a model added after this snapshot either answers the canonical whole-scenario request or is listed as not scorable rather than accommodated.

One row is a cloaked preview model listed as such. It is publicly callable under the canonical whole-scenario request and forced-tool transport, its maker is unnamed, and its date is its OpenRouter listing date.

### 3.4 Frozen snapshot and open-set status

Leaderboard positions are version-sensitive, so manuscript claims refer to the frozen source-run exports in Table 1 rather than to the live site. The committed source-run exports are the manuscript artifacts. The public site exposes the current prompts, predictions, explanations, and reference outputs for transparency. This makes the public leaderboard open-set: models, model providers, or benchmark users could learn from the released cases before later runs. Protected leaderboard claims would require a separate held-out or rotating set.

Source: [Article Notebook](#)

Table 1: Frozen manuscript snapshot.

| Item                | Value                    |
|---------------------|--------------------------|
| Snapshot date       | 2026-09-01               |
| Model response date | 2026-06-12 to 2026-09-01 |

Table 1: Frozen manuscript snapshot.

| Item                                 | Value                                                                                             |
|--------------------------------------|---------------------------------------------------------------------------------------------------|
| Policy period                        | US tax year 2026                                                                                  |
| Frozen export                        | paper/snapshot/20260501/runs/<br>us_full_run_20260612_poli-<br>cyengine_4_16_1_populace/data.json |
| Frozen export SHA-256 prefix         | US 80bb1988adcb                                                                                   |
| Snapshot manifest                    | paper/snapshot/20260501/manifest.json                                                             |
| Audit annotations                    | annotations/us_full_run_20260612_poli-<br>cyengine_4_16_1_populace                                |
| US scenarios SHA-256 prefix          | 71b16212f0c0                                                                                      |
| US reference outputs SHA-256 prefix  | b9136a15e285                                                                                      |
| Benchmark spec SHA-256 prefix        | ce233f8cbb05                                                                                      |
| Frozen prompt payload SHA-256 prefix | US 1c3f94362d0b                                                                                   |
| Scoring code SHA-256 prefix          | bde9e8de2bd2                                                                                      |
| US run label                         | us_full_run_20260612_poli-<br>cyengine_4_16_1_populace                                            |
| Household sample seed                | 42                                                                                                |
| PolicyEngine.py                      | 4.16.1                                                                                            |
| PolicyEngine-US                      | policyengine-us 1.755.4                                                                           |
| US dataset                           | populace_us_2024<br>(populace-us-2024-5da5a95-20260611)                                           |
| Households                           | 100 US                                                                                            |
| Models                               | 33                                                                                                |
| Output groups                        | 18 US                                                                                             |
| Condition                            | No tools, no web access; structured responses<br>covering every requested output                  |
| Response contract                    | Numeric answer and non-empty explanation<br>for every requested output                            |

Source: [Article Notebook](#)



Table 2: Model run configuration in the frozen manuscript snapshot. Every model answers through the transport its model card records, with no external tools, provider-default temperature and sampling, no reasoning-control parameters, and the same household facts and requested outputs. Transport is the answer contract: a forced answer-schema tool call (tool\_choice set to the answer tool) or a JSON object, the latter where the provider rejects a forced tool or where the model card selects JSON. Four rows carry supervised-run fingerprints for answer contract, request shape, tool choice, and completion ceiling; reasoning setup and timeouts for every row, and all fields for the other 29 rows, are the harness registry as frozen in the snapshot’s serving-configuration file. Rows with an output count answered the template over subsets of the requested outputs, an accommodation that predates the canonical whole-scenario rule described in the text. Released is the first public availability (paid tiers count; trusted-tester previews do not), compiled from vendor announcements and contemporaneous press; grok-build-0.1’s date rests on secondary trackers. Leaderboard artifact keys match the final segment of the provider id (qwen-3.7-max adds a hyphen). A dagger marks open-weight models; Kimi K3’s weights, announced at its API launch (Willison 2026), shipped on Hugging Face on July 27, 2026 under a custom license (Tech Times 2026). Ox Alpha is an OpenRouter stealth listing with no named maker; its Released date is its August 21, 2026 listing date.

| Model            | Provider ID               | Released   | Parsed          | Transport   | Request shape        | Reasoning setup                                                |
|------------------|---------------------------|------------|-----------------|-------------|----------------------|----------------------------------------------------------------|
| Claude Fable 5   | claude-fable-5            | 2026-06-09 | 1,984/<br>1,984 | forced tool | 1 output/<br>request | thinking<br>does not<br>engage<br>under<br>forced tool<br>call |
| Claude Fable 5.1 | claude-fable-5-1          | 2026-09-01 | 1,984/<br>1,984 | JSON        | whole<br>scenario    | provider<br>default;<br>16,384-<br>token<br>shared<br>budget   |
| Claude Haiku 4.5 | claude-haiku-4-5-20251001 | 2025-10-15 | 1,984/<br>1,984 | forced tool | 1 output/<br>request | provider<br>default                                            |
| Claude Opus 4.7  | claude-opus-4-7           | 2026-04-16 | 1,984/<br>1,984 | forced tool | 1 output/<br>request | provider<br>default                                            |
| Claude Opus 4.8  | claude-opus-4-8           | 2026-05-28 | 1,984/<br>1,984 | forced tool | 1 output/<br>request | provider<br>default                                            |

| Model                         | Provider ID                          | Released    | Parsed          | Transport   | Request shape    | Reasoning setup                                 |
|-------------------------------|--------------------------------------|-------------|-----------------|-------------|------------------|-------------------------------------------------|
| Claude Opus 5                 | claude-opus-5                        | 2026-07-24  | 1,984/<br>1,984 | forced tool | whole scenario   | thinking does not engage under forced tool call |
| Claude Sonnet 4.6             | claude-sonnet-4-6                    | 2026-02-17  | 1,984/<br>1,984 | forced tool | 1 output/request | provider default                                |
| Claude Sonnet 5               | claude-sonnet-5                      | 2026-06-30  | 1,984/<br>1,984 | forced tool | 1 output/request | thinking does not engage under forced tool call |
| DeepSeek V4 Pro               | deepseek/deepseek-v4-pro             | 2026-04-24† | 1,984/<br>1,984 | JSON        | whole scenario   | provider default; 16,384-token shared budget    |
| Gemini 3 Flash Preview        | gemini/gemini-3-flash-preview        | 2025-12-17  | 1,984/<br>1,984 | JSON        | whole scenario   | provider default                                |
| Gemini 3.1 Flash Lite Preview | gemini/gemini-3.1-flash-lite-preview | 2026-03-03  | 1,984/<br>1,984 | JSON        | whole scenario   | provider default                                |
| Gemini 3.1 Pro Preview        | gemini/gemini-3.1-pro-preview        | 2026-02-19  | 1,984/<br>1,984 | JSON        | whole scenario   | provider default                                |
| Gemini 3.5 Flash              | gemini/gemini-3.5-flash              | 2026-05-19  | 1,984/<br>1,984 | JSON        | whole scenario   | provider default                                |

| Model            | Provider ID                      | Released    | Parsed          | Transport   | Request shape     | Reasoning setup                                 |
|------------------|----------------------------------|-------------|-----------------|-------------|-------------------|-------------------------------------------------|
| Gemini 3.6 Flash | gemini/<br>gemini-3.6-flash      | 2026-07-21  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| Gemini 3.7 Flash | gemini/<br>gemini-3.7-flash      | 2026-08-13  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| GLM-5.2          | open-router/<br>z-ai/<br>glm-5.2 | 2026-06-13† | 1,845/<br>1,984 | JSON        | 3 outputs/request | provider default;<br>16,384-token shared budget |
| GPT-5.4 mini     | gpt-5.4-mini                     | 2026-03-17  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default                                |
| GPT-5.4 nano     | gpt-5.4-nano                     | 2026-03-17  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default                                |
| GPT-5.5          | gpt-5.5                          | 2026-04-23  | 1,984/<br>1,984 | forced tool | 3 outputs/request | provider default;<br>16,384-token shared budget |
| GPT-5.6 Luna     | gpt-5.6-luna                     | 2026-07-09  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| GPT-5.6 Sol      | gpt-5.6-sol                      | 2026-07-09  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |

| Model          | Provider ID                                   | Released    | Parsed          | Transport   | Request shape     | Reasoning setup                                 |
|----------------|-----------------------------------------------|-------------|-----------------|-------------|-------------------|-------------------------------------------------|
| GPT-5.6 Terra  | gpt-5.6-terra                                 | 2026-07-09  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| Grok 4.3       | xai/<br>grok-4.3                              | 2026-04-17  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default                                |
| Grok 4.5       | xai/<br>grok-4.5                              | 2026-07-08  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| Grok 4.6       | xai/<br>grok-4.6                              | 2026-08-12  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default;<br>16,384-token shared budget |
| Grok Build 0.1 | xai/grok-build-0.1                            | 2026-05-29  | 1,984/<br>1,984 | forced tool | whole scenario    | provider default                                |
| Inkling        | open-router/<br>thinking-machines/<br>inkling | 2026-07-15† | 1,984/<br>1,984 | forced tool | whole scenario    | provider default;<br>49,152-token shared budget |
| Kimi K2.6      | open-router/<br>moon-shotai/<br>kimi-k2.6     | 2026-04-20† | 1,558/<br>1,984 | JSON        | 3 outputs/request | provider default;<br>16,384-token shared budget |
| Kimi K3        | open-router/<br>moon-shotai/<br>kimi-k3       | 2026-07-16† | 1,920/<br>1,984 | JSON        | whole scenario    | provider default;<br>49,152-token shared budget |

| Model              | Provider ID                            | Released    | Parsed          | Transport   | Request shape         | Reasoning setup                                    |
|--------------------|----------------------------------------|-------------|-----------------|-------------|-----------------------|----------------------------------------------------|
| MiniMax M3         | open-router/<br>minimax/<br>minimax-m3 | 2026-06-01† | 1,984/<br>1,984 | forced tool | whole scenario        | provider default;<br>16,384-token<br>shared budget |
| Ox Alpha (preview) | open-router/<br>stealth/<br>ox-alpha   | 2026-08-21  | 1,984/<br>1,984 | forced tool | whole scenario        | provider default;<br>16,384-token<br>shared budget |
| Qwen3.7-max        | open-router/<br>qwen/<br>qwen3.7-max   | 2026-05-19  | 1,984/<br>1,984 | JSON        | 3 outputs/<br>request | provider default;<br>16,384-token<br>shared budget |
| Qwen3.8-Max        | open-router/<br>qwen/<br>qwen3.8-max   | 2026-08-03  | 1,984/<br>1,984 | JSON        | whole scenario        | provider default;<br>98,304-token<br>shared budget |

Source: [Article Notebook](#)

## 4 Data and scenario construction

### 4.1 United States

The US benchmark is built from the certified PolicyEngine populace dataset (populace\_us\_2024) using PolicyEngine US. The dataset contains 160,858 people in 75,112 households, with household, tax-unit, SPM-unit, family, marital-unit, and person records for PolicyEngine calculations. The sampled households are filtered to keep a single-tax-unit, single-family, single-Supplemental Poverty Measure (SPM)-unit structure with at least one adult and a supported filing status. Of the 75,112 households, 63,128 (84.0%) pass the filter and form the eligible draw (this filter-pass rate is coincidentally close to the 84% zero-output

share discussed under the benchmark design; the two figures are unrelated). The 16.0% excluded by the filter include multi-tax-unit households (e.g., adult roommates), multi-family households, multi-SPM-unit households, and households whose head reports a filing status outside the supported set. These excluded compositions are exactly the kind of cases where federal/state credit allocations and benefit-unit rules become hardest, so the eligible draw is a tractable subset rather than the full distribution of US households. Prompts include nonzero promptable raw inputs across relevant entities rather than a hand-curated summary, so the models see many of the same facts the simulator receives. Filing status is not stated in the prompt; the reference computation infers it from tax-unit role flags. Models therefore see the same household facts that drive the reference filing-status assignment, but they do not receive that assignment as a label.

The current US release requests 18 scored output groups spanning federal income tax, refundable credits, payroll and self-employment tax, state and local income tax, Supplemental Nutrition Assistance Program (SNAP), Supplemental Security Income (SSI), Temporary Assistance for Needy Families (TANF), school-meal eligibility, and person-level coverage eligibility for the Special Supplemental Nutrition Program for Women, Infants, and Children (WIC), Medicaid, the Children’s Health Insurance Program (CHIP), Medicare, Head Start, and Early Head Start; several of these categories span two scored groups each (federal and state refundable credits, payroll and self-employment tax, state and local income tax, free and reduced-price school meals), which is how the named categories reconcile to the 18 output groups. The source run also requested the Affordable Care Act (ACA) Premium Tax Credit (PTC), but explanation audits showed that the prompt could be misleading when households lacked plan-specific Marketplace information. We therefore preserve those raw responses but exclude PTC from the canonical scored leaderboard until the prompt contract is revised.

The output scope is intentionally narrower than the full PolicyEngine model. Table 3 summarizes the inclusion rule. The benchmark asks for WIC eligibility rather than a WIC dollar amount; WIC dollar values are used only as impact-weight proxies for coverage flags, not as requested model outputs.

Source: [Article Notebook](#)

Table 3: Output-selection rationale.

| Scope decision | Rationale                                                                                                                                                  |
|----------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Included       | Direct tax, credit, benefit, health-support, and coverage outputs that a household-facing model could plausibly be asked to estimate from household facts. |

Table 3: Output-selection rationale.

| Scope decision          | Rationale                                                                                                                                                                                                                     |
|-------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Excluded                | Intermediate tax bases, payroll subcomponents, and outputs that mainly require unavailable history, restricted local market data, restricted program-administration data, or take-up assignment rather than rule calculation. |
| Binary coverage outputs | Requested as 0/1 eligibility flags and scored as classification tasks; their dollar values are used only as impact-weight proxies, not as requested model outputs.                                                            |
| WIC                     | The benchmark asks for person-level WIC eligibility. It does not ask models to estimate a WIC dollar amount.                                                                                                                  |

Source: [Article Notebook](#)

## 4.2 Reference-output credibility

PolicyBench treats PolicyEngine outputs as benchmark reference outputs, not as administrative records. The reference source is nevertheless stronger than an ad hoc answer key: PolicyEngine is open source, used for household calculators and reform analysis, and externally checked in specific domains. No. 10 Downing Street’s data science team adapted PolicyEngine’s open-source microsimulation model for experimental policy simulation, with validation against external projections before use (Woodruff 2026; Ghenis 2026). In the US, PolicyEngine (2024) reports matching the National Bureau of Economic Research (NBER) TAXSIM-35 model (Feenberg and Coutts 1993) to the cent on the vast majority of cases for the 2021 tax year across hundreds of thousands of tax units per state, with state-specific differences documented in the integration tests. We do not restate that comparison as a single percentage because the published source uses qualitative phrasing rather than a headline accuracy number. PolicyEngine has also signed a memorandum of understanding (Ghenis and Makarchuk 2025) with the Federal Reserve Bank of Atlanta for future validation work against its Policy Rules Database (2026). The Atlanta Fed sources are a caveat rather than evidence of completed validation for this benchmark: they document planned collaboration and the comparison source, not finished checks of the frozen PolicyBench outputs. Taken together, these sources support using PolicyEngine as a transparent reference implementation with partial external validation, but they do not validate every benchmark output.

This does not make PolicyEngine infallible, so we treated the reference outputs as claims to be attacked rather than assumptions. Across development and three audit waves, every wrong cell received a model-assisted review, and every case in which a reviewing model suspected the reference itself was adjudicated by hand against primary sources — statutes, regulations, and published program parameters — and against the engine’s computation traces. Reviewers raised 35 reference-suspect flags in total. Thirty-one dissolved under adjudication, with the reviewer’s hypothesis failing against a primary source (stale thresholds, missed program mergers, missed federal-conformity elections, misread eligibility-income definitions) or against the engine trace. Four were genuine: two Head Start eligibility flags and one New Jersey state income tax cell reflected policyengine-us defects (an age-window bound and a filing-threshold floor the benchmarked path bypassed), and one SNAP cell rested on a scenario whose encoded inputs did not match the disability status its prompt implied. The engine defects were fixed upstream (policyengine-us #8845 and #8846, plus #8844, found in the same review but affecting no benchmark cell), and this snapshot’s reference outputs are generated with the fixed engine, policyengine-us 1.755.4. The mis-encoded scenario is disclosed under Limitations pending the next scenario refresh. The asymmetry of this record is itself evidence about the reference: even among the cells the reviewing models found most suspicious, 31 of 35 flags (89%) resolved to model-side errors, set against thousands of confirmed model errors overall and four reference defects. Table 4 summarizes the reviewed discrepancy classes and outcomes; no unresolved reference-computation defect is known in the frozen snapshot.

The audit tooling itself required hardening, which we report because it changes how much the annotations can be trusted. The derivation narratives displayed alongside reference values are generated from engine computation traces. After earlier model-written narratives were found to misattribute eligibility pathways, mechanical validators now enforce that each narrative states the engine’s actual mechanism — for example, that a child ineligible for CHIP because they qualify for Medicaid is never described as failing a CHIP income test. The failure-audit judge reads these grounded narratives together with engine facts and must return a per-model diagnosis naming the specific rule or computation the model missed. It is barred from re-adjudicating the reference in prose: hedged or reference-adjudicating verdicts are mechanically rejected and re-judged, and genuine reference doubt is routed to a structured flag that triggers the manual adjudication described above.

After freezing the snapshot and completing response-contract repairs, we annotated every wrong model-output row and every scenario-output case with at least one wrong row. Table 5 reports that all 7,840 rows receiving less than full score carry row-level and case-level annotations. Every row resolves to one of two final classes: `llm_error` (a substantive miscalculation, with a per-model diagnosis naming the rule or computation the model missed) or `parse_contract_failure` (a missing or unparsable answer, classified deterministically). No wrong row remains classified as a prompt ambiguity, unresolved reference issue, or needs-review item, and zero rows carry a standing reference-suspect flag. The audit pipeline is model-assisted with developer adjudication of every reference-suspect flag rather than independent external validation, but it is exhaustive over scored misses in the frozen snapshot.



Source: [Article Notebook](#)

Table 4: Development discrepancy review before the frozen snapshot.

| Discrepancy class                  | Review outcome                                                                                                                                                                                                                                                  |
|------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| US federal credits                 | Reviewed cases where explanations omitted EITC or refundable CTC; discrepancies reflected model credit-treatment errors, not PolicyEngine errors.                                                                                                               |
| US health coverage (Medicaid/CHIP) | Reviewed threshold-cluster suspects; all dissolved against 42 CFR 435 pathway rules (aged/disabled income tests, tax-dependent household composition, CHIP's exclusion of Medicaid-eligible children) confirmed in engine traces.                               |
| US SNAP, SSI, and school meals     | Reviewed false-zero suspects; references confirmed against categorical-eligibility, aged/blind/disabled, and universal-free-meal rules. One SNAP cell traced to a scenario encoding that omits the disability its prompt implies (disclosed under Limitations). |
| US state and local income tax      | Reviewed state-rule and bracket suspects; most applied wrong schedules or missed conformity elections and forgiveness programs. One genuine defect: New Jersey's filing-threshold floor was bypassed on the benchmarked path (fixed, policyengine-us #8845).    |
| US Head Start eligibility          | Two flags reflected a genuine age-window defect (fixed, policyengine-us #8846); references regenerated with the fix.                                                                                                                                            |
| US payroll and overtime            | Reviewed during development; upstream data fixes and prompt clarifications were applied before the frozen snapshot.                                                                                                                                             |

Source: [Article Notebook](#)

Table 5: Exhaustive annotation coverage for scored misses in the frozen snapshot.

| Country | Wrong rows audited | LLM response errors | Parse-contract failures | Unresolved prompt/reference rows | Missing row annotations | Missing case notes |
|---------|--------------------|---------------------|-------------------------|----------------------------------|-------------------------|--------------------|
| US      | 7840               | 7211                | 629                     | 0                                | 0                       | 0                  |
| Total   | 7840               | 7211                | 629                     | 0                                | 0                       | 0                  |

Source: [Article Notebook](#)

## 5 Results

### 5.1 United States leaderboard

The US leaderboard for the frozen manuscript snapshot is shown in Table 6. The top three models in that snapshot are GPT-5.6 Sol (88.7 exact, 95% CI 85.8–91.2; 90.9 within-1%), Claude Fable 5.1 (86.3 exact, 95% CI 83.2–89.1; 89.3 within-1%), and Kimi K3 (86.2 exact, 95% CI 82.1–89.6; 88.4 within-1%).

Source: [Article Notebook](#)

Table 6: Top US benchmark models in the frozen manuscript snapshot, ranked by the headline exact-match rate with the within-1% near-miss companion alongside. Exact 95% CIs and rank ranges are household-resampling bootstrap intervals (see Uncertainty).

|   | Model              | Exact % | Exact 95% CI | Exact rank | Within 1% % | Within 10% | Bounded score |
|---|--------------------|---------|--------------|------------|-------------|------------|---------------|
| 0 | GPT-5.6 Sol        | 88.7    | 85.8–91.2    | 1–2        | 90.9        | 93.7       | 95.3          |
| 1 | Claude Fable 5.1   | 86.3    | 83.2–89.1    | 1–4        | 89.3        | 93.4       | 95.3          |
| 2 | Kimi K3            | 86.2    | 82.1–89.6    | 1–6        | 88.4        | 91.9       | 93.4          |
| 3 | GPT-5.6 Luna       | 84.5    | 81.8–87.1    | 2–9        | 86.1        | 90.7       | 92.5          |
| 4 | Ox Alpha (preview) | 84.2    | 81.1–87.0    | 3–9        | 86.4        | 90.7       | 93.3          |

Source: [Article Notebook](#)

## 5.2 Uncertainty

Each leaderboard score is a mean over a sample of 100 households, so the rankings carry sampling uncertainty. We quantify it with a household-level bootstrap: the 100 households are resampled with replacement 10,000 times, and each model’s population-weighted exact-match rate is recomputed on every resample. The resampling unit is the household, carried with its full output vector, so within-household correlations across outputs are preserved. The same resampled set is used for every model in each replicate (a paired bootstrap), so the rank ranges in Table 6 reflect genuine ordering uncertainty rather than independent per-model noise.

Sampling uncertainty is substantial at this sample size. In the frozen snapshot, 7 of the other 32 models have an exact-match 95% interval that overlaps the leader’s. Small differences in the leaderboard ordering should therefore be read as ties rather than as a strict ranking, and the rank-range column makes the set of plausible positions explicit. These intervals are a manuscript artifact: the public leaderboard and app report point estimates only, because the benchmark sample, not run-to-run model variation, is the dominant source of uncertainty here.

## 5.3 Cost and latency

Accuracy is one axis; the price and speed of a no-tools answer are another. Table 7 reports each model’s estimated cost to compute one household’s full output vector and its median per-household latency, next to the headline exact rate. Per-household cost spans \\\\$0.002 (GPT-5.4 nano) to \\\\$0.541 (Claude Fable 5), and median per-household latency 3 s (Gemini 3.1 Flash Lite Preview) to 1023 s (Kimi K3). Neither dimension cleanly tracks accuracy: the top-scoring model, GPT-5.6 Sol at 88.7% exact, sits mid-pack at \\\\$0.090 per household and 36 s.

Source: [Article Notebook](#)

Table 7: Estimated cost and median latency per household for the frozen US snapshot, alongside the headline exact-match rate. Cost is reconstructed from logged token usage at provider list prices; latency is the median across households of each household’s summed request time. An em dash marks usage or wall-clock latency that was not measured per call.

|   | Model              | Exact (%) | Cost / household | Latency (median) |
|---|--------------------|-----------|------------------|------------------|
| 0 | GPT-5.6 Sol        | 88.7      | \\\$0.090        | 36 s             |
| 1 | Claude Fable 5.1   | 86.3      | \\\$0.257        | 49 s             |
| 2 | Kimi K3            | 86.2      | \\\$0.470        | 1023 s           |
| 3 | GPT-5.6 Luna       | 84.5      | \\\$0.019        | 14 s             |
| 4 | Ox Alpha (preview) | 84.2      | —                | 242 s            |
| 5 | Inkling            | 83.8      | \\\$0.091        | 250 s            |

Table 7: Estimated cost and median latency per household for the frozen US snapshot, alongside the headline exact-match rate. Cost is reconstructed from logged token usage at provider list prices; latency is the median across households of each household’s summed request time. An em dash marks usage or wall-clock latency that was not measured per call.

|    | Model                         | Exact (%) | Cost / household | Latency (median) |
|----|-------------------------------|-----------|------------------|------------------|
| 6  | GPT-5.5                       | 83.5      | \\$0.262         | —                |
| 7  | GPT-5.6 Terra                 | 83.4      | \\$0.039         | 19 s             |
| 8  | Grok 4.6                      | 82.8      | \\$0.087         | 219 s            |
| 9  | Grok 4.5                      | 80.9      | \\$0.073         | 112 s            |
| 10 | Claude Fable 5                | 79.9      | \\$0.541         | 98 s             |
| 11 | Claude Opus 5                 | 79.8      | \\$0.067         | 16 s             |
| 12 | Gemini 3.7 Flash              | 79.2      | \\$0.014         | 9 s              |
| 13 | Gemini 3.6 Flash              | 79.0      | \\$0.055         | 26 s             |
| 14 | Gemini 3.1 Pro Preview        | 77.9      | \\$0.085         | 45 s             |
| 15 | Claude Opus 4.7               | 77.4      | \\$0.293         | 53 s             |
| 16 | Grok 4.3                      | 77.2      | \\$0.015         | 22 s             |
| 17 | Claude Sonnet 4.6             | 77.1      | \\$0.186         | 114 s            |
| 18 | Gemini 3 Flash Preview        | 76.9      | \\$0.048         | 58 s             |
| 19 | Gemini 3.5 Flash              | 76.2      | \\$0.056         | 24 s             |
| 20 | Grok Build 0.1                | 76.1      | \\$0.045         | 135 s            |
| 21 | Gemini 3.1 Flash Lite Preview | 76.1      | \\$0.002         | 3 s              |
| 22 | DeepSeek V4 Pro               | 76.1      | \\$0.011         | 170 s            |
| 23 | Qwen3.7-max                   | 73.6      | \\$0.139         | —                |
| 24 | GLM-5.2                       | 73.1      | \\$0.082         | —                |
| 25 | Claude Opus 4.8               | 72.6      | \\$0.255         | 55 s             |
| 26 | MiniMax M3                    | 72.4      | \\$0.007         | —                |
| 27 | Claude Haiku 4.5              | 71.7      | \\$0.055         | 48 s             |
| 28 | Qwen3.8-Max                   | 71.5      | \\$0.224         | 861 s            |
| 29 | GPT-5.4 mini                  | 70.5      | \\$0.006         | 5 s              |
| 30 | Claude Sonnet 5               | 69.4      | \\$0.172         | 64 s             |
| 31 | Kimi K2.6                     | 64.6      | \\$0.391         | —                |
| 32 | GPT-5.4 nano                  | 62.3      | \\$0.002         | 6 s              |

Source: [Article Notebook](#)

Cost is reconstructed from each call’s logged prompt and completion tokens at provider list prices through litellm; for grok-build-0.1, which litellm’s price map did not yet cover, we apply the provider’s published rate of \$1 and \$2 per million input and output tokens. Latency is wall-clock request time as observed by the harness, summed within each household and taken as the median across the 100 households. The per-call timer spans any provider-side

retry or rate-limit backoff, so the median is reported in preference to the mean, which a few throttled calls inflate. Latency is not comparable across serving regimes: Claude-family models issue chunked serial calls that lengthen wall-clock time, several models ran through batch or asynchronous dispatch whose logged times reflect submission rather than generation (shown as —), and Claude Fable 5’s usage and latency are recorded at the run level rather than per call. Both cost and latency are properties of this run’s providers, prompts, and serving configurations at evaluation time rather than fixed model constants.

## 5.4 Simple baselines

Simple baselines help interpret score levels in a zero-heavy benchmark. Table 8 reports an always-zero response and a median-reference-by-output response on the same frozen household sample, scored with the household-impact weighting. These are not model competitors; they show how much of each metric can be earned without household-specific policy calculation. The always-zero baseline scores 66.9% on the headline exact-match rate — below the 84% of outputs that are exact zeros because the weighting concentrates mass in high-dollar output groups that are rarely zero (Section 3.1). The top model clears the baseline by 21.8 points, and the within-1% companion shows the same separation with near-miss tolerance.

Source: [Article Notebook](#)

Table 8: Simple baseline scores on the frozen manuscript snapshot, household-impact-weighted. The always-zero exact rate sits below the unweighted zero share because the weighting concentrates on high-dollar output groups that are rarely zero.

|   | Baseline                   | Exact % | Within 1% % | Bounded score |
|---|----------------------------|---------|-------------|---------------|
| 0 | Always zero                | 66.9    | 66.9        | 66.9          |
| 1 | Median reference by output | 61.0    | 61.0        | 64.3          |

Source: [Article Notebook](#)

## 5.5 Sensitivity to benchmark view

The primary leaderboard gives each household equal weight. Table 9 reports several alternative views, including the older equal-output-group score. The top model is stable across several views, but lower ranks move across sensitivity views. Scores are therefore best read with the weighting choice visible rather than collapsed into a single ordering.

Source: [Article Notebook](#)

Table 9: Leaderboard sensitivity to alternative scoring views.

|   | View                          | Rank 1                  | Rank 2                  | Rank 3                  |
|---|-------------------------------|-------------------------|-------------------------|-------------------------|
| 0 | Exact-match headline          | GPT-5.6 Sol (88.7)      | Claude Fable 5.1 (86.3) | Kimi K3 (86.2)          |
| 1 | Within-1% companion           | GPT-5.6 Sol (90.9)      | Claude Fable 5.1 (89.3) | Kimi K3 (88.4)          |
| 2 | Bounded score                 | GPT-5.6 Sol (95.3)      | Claude Fable 5.1 (95.3) | Inkling (94.9)          |
| 3 | Equal-output-group score      | Claude Fable 5.1 (96.9) | GPT-5.6 Sol (96.8)      | Gemini 3.6 Flash (96.1) |
| 4 | Amount outputs only           | Claude Fable 5.1 (96.1) | GPT-5.6 Sol (95.3)      | GPT-5.5 (94.5)          |
| 5 | Positive reference cases only | GPT-5.5 (85.8)          | Gemini 3.5 Flash (83.7) | GPT-5.6 Sol (81.5)      |
| 6 | Zero reference cases only     | Claude Fable 5.1 (98.8) | Gemini 3.6 Flash (98.7) | Grok 4.6 (98.7)         |

Source: [Article Notebook](#)

The bounded score is a global-output-weighted average of row scores. Amount outputs use  $\max(0, 1 - |\text{pred} - \text{ref}| / |\text{ref}|)$  when the reference is nonzero and  $1\{\text{pred} = 0\}$  when the reference is zero. Binary outputs use exact 0/1 matching. The exact-match rate uses the same weighting and aggregation but replaces the amount row score with the headline indicator:  $1\{|\text{pred} - \text{ref}| = 1\}$  (one currency unit). Each output group’s default household-impact weight is computed from a full weighting population, not from the 100 benchmark households: the certified PolicyEngine US populace dataset. For each source household, the share is  $|\text{ref\_ij}| / \max(|\text{household\_net\_income\_i}|, \sum_k |\text{ref\_ik}|)$ ; shares are averaged with calibrated household weights and renormalized so output weights sum to one. The  $\max(\dots)$  denominator anchors per-household shares to net income when net income is the dominant flow and falls back to the gross tax-benefit flow only when programs cancel each other out, so a \$1 benefit to a high-earner contributes essentially zero weight and per-household shares never exceed one. Booleans carry weight through PolicyEngine’s paired per-capita value (for example `medicaid_value`), so eligibility calls are graded as accuracy but weighted by the dollar stake. Table 10 compares the bounded-score ranking against two opt-in alternatives shown for transparency: equal weighting (every output the same) and budget-weighted (each output’s share of total absolute reference dollars in the same full weighting population).

Source: [Article Notebook](#)

Table 10: Top country ranks under the bounded secondary metric across three weightings (household, equal, budget).

|   | Weighting                    | Rank 1                  | Rank 2                  | Rank 3          |
|---|------------------------------|-------------------------|-------------------------|-----------------|
| 0 | Household-weighted (default) | GPT-5.6 Sol (95.3)      | Claude Fable 5.1 (95.3) | Inkling (94.9)  |
| 1 | Equal weights                | Claude Fable 5.1 (96.6) | GPT-5.6 Sol (96.4)      | Grok 4.6 (95.8) |
| 2 | Budget-weighted              | Claude Fable 5.1 (94.8) | GPT-5.6 Sol (94.5)      | Kimi K3 (92.7)  |

Source: [Article Notebook](#)

## 5.6 Hardest benchmark targets

The lowest-scoring US variables are multi-step tax quantities and sparse positive-dollar benefits. As shown in Table 11, federal income tax before refundable credits, state income tax before refundable credits, state refundable credits, SNAP, and federal refundable credits are the hardest US outputs by bounded score in the frozen manuscript snapshot.

Source: [Article Notebook](#)

Table 11: Hardest US variables by bounded score (the secondary tracking metric, used here because it orders the hardest variables differently from the hit rates and separates them more finely; on their positive-reference cases the hit rates collapse toward single digits — see the zero-versus-positive table).

|    | Variable                                     | Score | Exact | Within 1% | Within 10% |
|----|----------------------------------------------|-------|-------|-----------|------------|
| 0  | federal_income_tax_before_refundable_credits | 77.4  | 51.0  | 53.7      | 63.8       |
| 15 | state_income_tax_before_refundable_credits   | 79.6  | 54.7  | 57.0      | 68.5       |
| 16 | state_refundable_credits                     | 81.9  | 78.4  | 78.6      | 79.3       |
| 13 | snap                                         | 84.6  | 77.1  | 78.3      | 81.9       |
| 1  | federal_refundable_credits                   | 89.0  | 83.1  | 84.2      | 86.8       |

Source: [Article Notebook](#)

## 5.7 Zero and positive cases

Overall hit rates can overstate performance on sparse programs because correct zeros are common. Table 12 therefore reports within-10% performance for all cases, positive-reference cases, and zero-reference cases. The gap is large for several outputs: models often identify that a program or tax does not apply, but miss the amount when the reference value is positive. Figure 1 plots the same split by output group.

Source: [Article Notebook](#)

Table 12: US within-10% performance by zero versus positive reference cases.

|   | Variable                                     | All cases | Positive cases | Zero cases |
|---|----------------------------------------------|-----------|----------------|------------|
| 0 | federal_income_tax_before_refundable_credits | 65.0      | 36.4           | 90.8       |
| 1 | state_income_tax_before_refundable_credits   | 69.6      | 39.8           | 91.5       |
| 2 | state_refundable_credits                     | 80.6      | 11.4           | 98.9       |
| 3 | snap                                         | 82.4      | 24.9           | 96.6       |
| 4 | federal_refundable_credits                   | 88.5      | 46.5           | 94.7       |

Table 12: US within-10% performance by zero versus positive reference cases.

| Variable                   | All cases | Positive cases | Zero cases |
|----------------------------|-----------|----------------|------------|
| 5 payroll_tax              | 90.3      | 86.0           | 98.0       |
| 6 person_medicaid_eligible | 90.7      | 79.5           | 94.1       |
| 7 person_chip_eligible     | 95.9      | —              | 95.9       |

Source: [Article Notebook](#)




Figure 1: Output-level performance on zero-reference and positive-reference cases.

5.8 Open-weight time lag

Source: [Article Notebook](#)

Source: [Article Notebook](#)

The leaderboard also measures time. Following the lag construction the UK AI Security Institute applied to cyber capabilities (UK AI Security Institute 2026), the open-weight lag is



the time between the closed frontier first reaching a score and the leading open-weight model matching it. Two cuts of the frozen roster give the trajectory. Among models released before July 2026, the leading open-weight model is DeepSeek V4 Pro at 76.1% exact, whose weights shipped on April 24, 2026 (DeepSeek 2026); the earliest-released closed model above that score is Gemini 3 Flash Preview, released December 17, 2025 (TechCrunch 2025) — a lag of 128 days, about 4.2 months. On the full roster, Kimi K3 at 86.2% exact trails two closed models; the earliest released of them, GPT-5.6 Sol, generally available July 9, 2026 (TechCrunch 2026), puts the lag at 7 days by API availability, or 18 days dated to Kimi K3’s July 27 weight release on Hugging Face (Tech Times 2026).

Both figures are lower bounds in one direction: a closed model outside the evaluated roster that crossed either score earlier would lengthen the corresponding lag. They also carry the uncertainty of any single-benchmark, single-domain estimate — the AISI analysis reports substantially longer lags on cyber tasks — and a lag measured at a moving frontier compresses mechanically when a strong open-weight release lands days after a frontier release.

## 6 Failure modes

The benchmark surfaces a few recurring failure patterns.

First, models miss positive tax and benefit quantities more often than zero cases. The lowest-scoring outputs by bounded score (Table 11) require the model to choose the right income concepts, exclusions, program thresholds, and sequencing before applying any final subtraction. Positive benefit cases remain difficult, so the result should not be read as a general claim that benefits are easy.

Second, joint accuracy across interacting components is lower than marginal accuracy on either component. Table 13 shows within-10% accuracy for `federal_refundable_credits`, `state_refundable_credits`, and the conjunction of both within the same household. The joint hit rate can be no higher than either marginal and is strictly lower for all but the top model, so leaderboard scores that average across outputs understate how often a model gets a single household’s federal/state credit allocation jointly correct.

Source: [Article Notebook](#)

Table 13: US within-10% accuracy on federal vs state refundable credits and the household-level joint.

|    | Model            | Federal within 10% | State within 10% | Joint within 10% |
|----|------------------|--------------------|------------------|------------------|
| 1  | Claude Fable 5.1 | 98.0               | 90.0             | 90.0             |
| 20 | GPT-5.6 Sol      | 99.0               | 86.0             | 86.0             |
| 18 | GPT-5.5          | 95.0               | 85.0             | 82.0             |
| 24 | Grok 4.6         | 96.0               | 83.0             | 81.0             |

Table 13: US within-10% accuracy on federal vs state refundable credits and the household-level joint.

|    | Model                         | Federal within 10% | State within 10% | Joint within 10% |
|----|-------------------------------|--------------------|------------------|------------------|
| 13 | Gemini 3.6 Flash              | 95.0               | 82.0             | 81.0             |
| 26 | Inkling                       | 94.0               | 81.0             | 79.0             |
| 9  | Gemini 3 Flash Preview        | 92.0               | 82.0             | 78.0             |
| 14 | Gemini 3.7 Flash              | 90.0               | 83.0             | 78.0             |
| 21 | GPT-5.6 Terra                 | 95.0               | 79.0             | 77.0             |
| 28 | Kimi K3                       | 94.0               | 79.0             | 77.0             |
| 0  | Claude Fable 5                | 91.0               | 82.0             | 76.0             |
| 19 | GPT-5.6 Luna                  | 94.0               | 79.0             | 76.0             |
| 23 | Grok 4.5                      | 92.0               | 78.0             | 76.0             |
| 11 | Gemini 3.1 Pro Preview        | 90.0               | 82.0             | 76.0             |
| 8  | DeepSeek V4 Pro               | 90.0               | 80.0             | 76.0             |
| 30 | Ox Alpha (preview)            | 90.0               | 81.0             | 75.0             |
| 12 | Gemini 3.5 Flash              | 90.0               | 79.0             | 74.0             |
| 10 | Gemini 3.1 Flash Lite Preview | 86.0               | 79.0             | 73.0             |
| 6  | Claude Sonnet 4.6             | 93.0               | 79.0             | 73.0             |
| 5  | Claude Opus 5                 | 87.0               | 81.0             | 73.0             |
| 25 | Grok Build 0.1                | 88.0               | 78.0             | 72.0             |
| 17 | GPT-5.4 nano                  | 86.0               | 79.0             | 72.0             |
| 4  | Claude Opus 4.8               | 84.0               | 79.0             | 70.0             |
| 22 | Grok 4.3                      | 84.0               | 79.0             | 70.0             |
| 16 | GPT-5.4 mini                  | 81.0               | 79.0             | 69.0             |
| 7  | Claude Sonnet 5               | 80.0               | 79.0             | 68.0             |
| 3  | Claude Opus 4.7               | 85.0               | 78.0             | 67.0             |
| 29 | MiniMax M3                    | 80.0               | 79.0             | 67.0             |
| 15 | GLM-5.2                       | 80.0               | 72.0             | 65.0             |
| 2  | Claude Haiku 4.5              | 75.0               | 79.0             | 64.0             |
| 31 | Qwen3.7-max                   | 78.0               | 79.0             | 63.0             |
| 32 | Qwen3.8-Max                   | 67.0               | 79.0             | 56.0             |
| 27 | Kimi K2.6                     | 45.0               | 48.0             | 29.0             |

Source: [Article Notebook](#)

Third, structured-output reliability is part of the benchmark contract. Missing or unparseable numeric values are not dropped. Appendix A documents parser recovery and the supervised runner’s bounded retry rounds. The 629 cells (1.0%) that still returned no parsable value after retries remain in the canonical files and score zero as parse-contract failures. Appendix A gives the per-model breakdown.

## 7 Limitations

PolicyBench is not a substitute for a production tax-and-benefit calculator. Several caveats matter:

- output-contract reliability required a repair workflow; raw provider responses are retained within the canonical prediction files wherever the transport exposed them (no raw payload is retained for Claude Fable 5 (1,984 rows) and Kimi K3 (64 rows)), and cells that never parsed score zero
- request shape is not identical across models: 10 of 33 models answered the prompt template over subsets of the requested outputs rather than the whole scenario (Table 2), an accommodation that, if it moves scores at all, should favor those models; later additions must answer the whole-scenario request or are listed as not scorable
- zero-heavy outputs require separate positive-case interpretation
- the choice of headline metric (exact match), near-miss companion (within-1%), and secondary tracking metric (bounded score) is one benchmark view; all three are reported, but downstream evaluations should also consider error-magnitude metrics (mean absolute error, mean absolute percentage error) where appropriate
- the public scenario explorer exposes the current test set and reference outputs, so open-set leakage is a prominent limitation rather than a minor implementation detail
- the 100-household manuscript snapshot should be treated as a preview until larger frozen runs are published
- the scored-miss audit is exhaustive over this frozen snapshot, but it is developer-adjudicated and not an independent validation of every reference value
- one frozen scenario encodes household inputs that do not match its prompt narrative (a South Carolina household whose prompt implies a qualifying disability the inputs omit), and two input-semantics conventions (SSI disability status and adult tax-dependent relationships) are under-specified in ways a scheduled scenario refresh will address; affected cells score models against the encoded inputs
- benchmark success should not be interpreted as policy-advice readiness

The current paper is therefore an evaluation of model performance under a specific structured-output benchmark, not a general certification of tax or benefit competence.

## 8 Conclusion

PolicyBench shows a consistent pattern across the US benchmark. Models often identify non-applicability, but positive tax and benefit amounts remain difficult, especially for multi-step income tax and positive benefit cases. In the frozen manuscript snapshot, GPT-5.6 Sol is the top-scoring US model.

These results support a narrow conclusion: unaided frontier models still struggle to reproduce selected household-level microsimulation outputs under a structured public benchmark. They do not show that models cannot assist policy analysis, and they do not validate PolicyEngine outputs as administrative truth. They suggest that future evaluations should separate no-tool estimation, tool-using system design, and reference-output validation more explicitly.

Next steps are to expand country coverage, increase frozen sample sizes, and add protected or rotating evaluation sets so that public rankings are less exposed to open-set leakage. The benchmark should also continue reporting sensitivity views, because scores are useful summaries only when their weighting choices are visible.

## Appendix A: Structured-output audit

The benchmark requires one numeric value and one explanation for every requested output. Parse failures are benchmark failures, not missing data to drop from the denominator. Before publishing the manuscript snapshot, we audited missing numeric values and explanations, extended the parser only where explicit variable-keyed value and explanation fields were recoverable, and retried failing scenarios under the supervised runner’s bounded rounds. Table 14 summarizes that sequence.

Source: [Article Notebook](#)

Table 14: Structured-output parser audit and repair sequence.

| Step                         | Finding                                                                                                                                                                                                                                       |
|------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Initial parse-contract audit | The source run was audited for rows missing a parsed numeric value or non-empty explanation; those rows were not dropped from the denominator.                                                                                                |
| Parser repair                | The parser was extended to recover explicit value and non-empty explanation blocks from nested, escaped, or partially truncated provider JSON without scraping prose numbers.                                                                 |
| Retry policy                 | Later waves ran under a resumable supervised runner that retries failed or timed-out scenarios in bounded rounds; every model’s canonical file covers all 100 households. No separate whole-response retry artifacts exist for this snapshot. |

Table 14: Structured-output parser audit and repair sequence.

| Step                 | Finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
|----------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Final parse coverage | 629 of the 65,472 canonical model-output rows (1.0%) have no parsable numeric value after retries. The failures comprise Kimi K2.6 (426), GLM-5.2 (139), and Kimi K3 (64). They score zero as parse-contract failures. A further 61 rows (Grok 4.3 (56), Kimi K2.6 (4), and Claude Haiku 4.5 (1)) carry a parsed numeric value but no explanation; they are scored on the numeric value and counted here as explanation-contract violations, so 690 rows in all fall short of the numeric-plus-explanation contract. Every other row has both. |
| Preservation rule    | The snapshot retains raw provider responses in the compressed source-run prediction files wherever the transport exposed them; no raw payload is retained for Claude Fable 5 (1,984 rows) and Kimi K3 (64 rows) (batch-adaptor rows carry no per-row payload, and those parse failures captured none).                                                                                                                                                                                                                                         |

Source: [Article Notebook](#)

Model waves after June ran under a resumable supervised runner that retries failed or timed-out scenarios in bounded rounds. No separate retry artifacts exist; repairs are folded into each model’s canonical prediction file. After retries, 629 cells (1.0% of canonical rows) still returned no parsable numeric value. They comprise Kimi K2.6 (426), GLM-5.2 (139), and Kimi K3 (64). These cells remain in the canonical files, with their raw responses preserved wherever the transport exposed one (no raw payload is retained for Claude Fable 5 (1,984 rows) and Kimi K3 (64 rows)). They score zero as parse-contract failures, and the denominator is never reduced.

## Competing interests

The author leads PolicyEngine, which develops the microsimulation software used to produce the benchmark reference outputs, and is a founder of the Axiom Foundation, a nonprofit developing deterministic rules-as-code infrastructure. Both organizations’ missions are advanced by the finding that unaided language models do not reliably reproduce rule calculations.

Source: [Article Notebook](#)

- Abbood, Auss, Zaiqiao Meng, and Nigel Collier. 2025. “Time to Revisit Exact Match.” Findings of the Association for Computational Linguistics: EMNLP 2025 (Suzhou, China), 11903–26. <https://doi.org/10.18653/v1/2025.findings-emnlp.637>.
- Bock, Michael R., Kara Molisee, Zachary Ozer, and Sumit Shah. 2025. “TaxCalcBench: Evaluating Frontier Models on the Tax Calculation Task.” arXiv Preprint arXiv:2507.16126, ahead of print. <https://doi.org/10.48550/arXiv.2507.16126>.
- Chen, Zhiyu, Wenhui Chen, Charese Smiley, et al. 2021. “FinQA: A Dataset of Numerical Reasoning over Financial Data.” arXiv Preprint arXiv:2109.00122. <https://arxiv.org/abs/2109.00122>.
- Cobbe, Karl, Vineet Kosaraju, Mohammad Bavarian, et al. 2021. “Training Verifiers to Solve Math Word Problems.” arXiv Preprint arXiv:2110.14168. <https://arxiv.org/abs/2110.14168>.
- DeepSeek. 2026. DeepSeek-V4 Preview Release. DeepSeek API news, April 24, 2026. <https://api-docs.deepseek.com/news/news260424/>.
- Federal Reserve Bank of Atlanta. 2026. Policy Rules Database. Federal Reserve Bank of Atlanta. <https://www.atlantafed.org/what-we-study/workforce-development/advancing-careers-for-low-income-families/policy-rules-database>.
- Feenberg, Daniel R., and Elisabeth Coutts. 1993. “An Introduction to the TAXSIM Model.” *Journal of Policy Analysis and Management* 12 (1): 189–94.
- Ghenis, Max. 2026. PolicyEngine Powers Rapid Policy Analysis at No 10 Downing Street. PolicyEngine. <https://www.policyengine.org/us/research/policyengine-10-downing-street>.
- Ghenis, Max, and Pavel Makarchuk. 2025. PolicyEngine and Atlanta Fed Sign Memorandum of Understanding for Policy Rules Database Validation. PolicyEngine. <https://www.policyengine.org/us/research/policyengine-atlanta-fed-mou-prd>.
- Guha, Neel, Julian Nyarko, Daniel E. Ho, et al. 2023. “LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models.” arXiv Preprint arXiv:2308.11462. <https://arxiv.org/abs/2308.11462>.
- Hendrycks, Dan, Collin Burns, Saurav Kadavath, et al. 2021. “Measuring Mathematical Problem Solving with the MATH Dataset.” *NeurIPS Datasets and Benchmarks Track*. <https://arxiv.org/abs/2103.03874>.
- Holzenberger, Nils, Benjamin Van Durme, Sarah Lawsky, and Kyle Richardson. 2021. “Fac-

- toring Statutory Reasoning as Language Understanding Challenges.” arXiv Preprint arXiv:2105.07903. <https://arxiv.org/abs/2105.07903>.
- Nava Labs. 2025. Experimenting with AI-Powered Tools in Public Benefits. Case study. <https://www.navapbc.com/case-studies/ai-tools-public-benefits>.
- Nava Labs. 2026. Evaluating a GenAI-Powered Assistive Chatbot for Caseworkers. Case study. <https://www.navapbc.com/case-studies/evaluating-ai-assistive-chatbot-caseworkers>.
- PolicyEngine. 2024. PolicyEngine Launches State Income Tax Modeling Nationwide. PolicyEngine. <https://www.policyengine.org/us/research/state-tax-model-beta>.
- Shanahan, Catherine, Emma McCarthy, Yan Zhao, et al. 2025. “Performance of LLMs on VITA Test: Potential for AI-Assisted Tax Returns for Low Income Taxpayers.” Artificial Intelligence and Law, ahead of print. <https://doi.org/10.1007/s10506-025-09465-7>.
- Shorten, Connor, Charles Pierse, Thomas Benjamin Smith, et al. 2024. “StructuredRAG: JSON Response Formatting with Large Language Models.” arXiv Preprint arXiv:2408.11061, ahead of print. <https://doi.org/10.48550/arXiv.2408.11061>.
- Sutherland, Holly, and Francesco Figari. 2013. “EUROMOD: The European Union Tax-Benefit Microsimulation Model.” International Journal of Microsimulation 6 (1): 4–26. <https://doi.org/10.34196/ijm.00075>.
- Tech Times. 2026. “Kimi K3 Open Weights Arrive Sunday: Self-Hosting Cuts China Data Risk the API Never Can.” July. <https://www.techtimes.com/articles/321551/20260725/kimi-k3-open-weights-arrive-sunday-self-hosting-cuts-china-data-risk-api-never-can.htm>.
- TechCrunch. 2025. Google Launches Gemini 3 Flash, Makes It the Default Model in the Gemini App. TechCrunch, December 17, 2025. <https://techcrunch.com/2025/12/17/google-launches-gemini-3-flash-makes-it-the-default-model-in-the-gemini-app/>.
- TechCrunch. 2026. OpenAI Launches Its New Family of Models with GPT-5.6. TechCrunch, July 9, 2026. <https://techcrunch.com/2026/07/09/openai-launches-its-new-family-of-models-with-gpt-5-6/>.
- UK AI Security Institute. 2026. How Far Behind the Frontier Are Leading Open-Weight Models on Cyber? Blog post. <https://www.aisi.gov.uk/blog/how-far-behind-the-frontier-are-leading-open-weight-models-on-cyber>.
- Vals AI. 2026. Public Benefits Bench. Benchmark report, in partnership with the Center for Civic Futures and Code for America. <https://www.vals.ai/benchmarks/public-benefits-bench>.

- Willison, Simon. 2026. Kimi K3, and What We Can Still Learn from the Pelican Benchmark. Simonwillison.net, July 16, 2026. <https://simonwillison.net/2026/Jul/16/kimi-k3/>.
- Woodruff, Nikhil. 2026. Informing Policy Using Micro-Simulations. No10 Innovation Fellowship. <https://fellows.ai.gov.uk/articles/nikhil-woodruff-micro-simulation>.
- ZenML LLMOps Database. 2025. Building and Automating Comprehensive LLM Evaluation Framework for SNAP Benefits. Industry writeup. <https://www.zenml.io/llmops-database/building-and-automating-comprehensive-llm-evaluation-framework-for-snap-benefits>.
- Zhou, Ruiwen, Wenyue Hua, Liangming Pan, et al. 2025. “RuleArena: A Benchmark for Rule-Guided Reasoning with LLMs in Real-World Scenarios.” Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (Vienna, Austria), 550–72. <https://doi.org/10.18653/v1/2025.acl-long.27>.
- Zhu, Fengbin, Wenqiang Lei, Youcheng Huang, et al. 2021. “TAT-QA: A Question Answering Benchmark on a Hybrid of Tabular and Textual Content in Finance.” arXiv Preprint arXiv:2105.07624. <https://arxiv.org/abs/2105.07624>.