

# PolicyBench

Benchmarking no-tool tax-and-benefit estimation in frontier language models

Max Ghenis

## Table of contents

|     |                                               |    |
|-----|-----------------------------------------------|----|
| 1   | Abstract                                      | 1  |
| 2   | Introduction                                  | 1  |
| 3   | Related work                                  | 2  |
| 4   | Benchmark design                              | 3  |
| 4.1 | Frozen snapshot and open-set status . . . . . | 4  |
| 5   | Data and scenario construction                | 6  |
| 5.1 | United States . . . . .                       | 6  |
| 5.2 | United Kingdom . . . . .                      | 7  |
| 5.3 | Reference-output credibility . . . . .        | 7  |
| 6   | Results                                       | 9  |
| 6.1 | United States leaderboard . . . . .           | 9  |
| 6.2 | United Kingdom leaderboard . . . . .          | 9  |
| 6.3 | Global shared-model leaderboard . . . . .     | 10 |
| 6.4 | Simple baselines . . . . .                    | 10 |
| 6.5 | Sensitivity to benchmark view . . . . .       | 12 |
| 6.6 | Sampling uncertainty . . . . .                | 13 |
| 6.7 | Hardest benchmark targets . . . . .           | 13 |
| 6.8 | Zero and positive cases . . . . .             | 14 |
| 7   | Failure modes                                 | 15 |
| 8   | Limitations                                   | 17 |
| 9   | Conclusion                                    | 18 |

|                                        |    |
|----------------------------------------|----|
| 10 Appendix A: Structured-output audit | 18 |
| 11 Competing interests                 | 21 |

Source: [Article Notebook](#)

## 1 Abstract

PolicyBench evaluates whether frontier language models can estimate household tax and benefit outputs from household facts without tools. The benchmark covers the United States and the United Kingdom, uses sampled US Enhanced Current Population Survey (CPS)-derived scenarios and a public UK calibrated transfer dataset, and scores outputs with a bounded 0-100 metric. The manuscript snapshot is a 100-household-per-country public preview, not a protected held-out leaderboard: the current public scenario explorer exposes prompts and reference outputs, so open-set leakage is a central limitation of any public ranking. In the frozen snapshot used for this manuscript (2026-05-15), gpt-5.5 is top-scoring in the US, UK, and shared global views at 70.3, 50.3, and 60.3, respectively. In both countries, multi-step tax quantities and positive-dollar benefit cases are materially harder than zero cases. The live benchmark is available at <https://policybench.org>.

## 2 Introduction

Household tax-and-benefit estimation sits between arithmetic and policy discussion. Each case has numeric labels, but generating those labels requires filing status, household composition, income concepts, program thresholds, and jurisdiction-specific rules. PolicyBench measures whether models can map household records to those outputs.

Most current language-model evaluations do not test this mapping directly. General math benchmarks emphasize symbolic manipulation and exact final answers (Cobbe et al. 2021; Hendrycks et al. 2021), while generic question-answering benchmarks emphasize recall or instruction following. PolicyBench instead asks whether models can transform household facts into policy outputs without access to tools or simulators.

This matters for public-facing calculators, analyst workflows, and tax-preparation or screening systems. The benchmark is intended to separate verbal policy fluency from household-level quantitative prediction.

### 3 Related work

The most relevant prior work is tax and statutory reasoning. SARA frames statutory interpretation, including tax-relevant rule application, as a language-understanding problem (Holzenberger et al. 2021). LegalBench broadens this view and includes tax-oriented numeric tasks such as `sara_numeric` (Guha et al. 2023). RuleArena evaluates rule-guided reasoning in regulation-style settings (Zhou et al. 2025), and TaxCalcBench evaluates frontier models on tax calculation from structured return-like inputs (Bock et al. 2025). Shanahan et al. report a similar split on the Volunteer Income Tax Assistance (VITA) test: models do better on tax knowledge questions than on open-ended calculation tasks (Shanahan et al. 2025).

PolicyBench also sits within a broader literature on quantitative reasoning benchmarks. Canonical math evaluations such as GSM8K and MATH normalized exact final-answer scoring for numeric tasks (Cobbe et al. 2021; Hendrycks et al. 2021). That norm is less natural for continuous-valued policy outputs, where being within 1 percent is closer than being off by an order of magnitude. Recent work has questioned pure exact-match scoring in numeric QA and temporal reasoning, arguing for error-aware metrics instead (Abbood et al. 2025). PolicyBench follows that direction by combining exactness with near-miss thresholds rather than collapsing everything into exact match alone.

Finance-domain benchmarks provide another relevant precedent. FinQA and TAT-QA both show that realistic quantitative reasoning often requires combining structured numbers with domain knowledge rather than solving stylized school-math problems (Chen et al. 2021; Zhu et al. 2021). PolicyBench differs in that its reference outputs are generated by executable tax-benefit microsimulation rather than annotated document questions, but the underlying motivation is similar: domain-specific quantitative reasoning deserves dedicated evaluation.

Finally, PolicyBench depends on two infrastructure literatures that are not themselves large language model (LLM) benchmarks. One is structured-output reliability, since the benchmark relies on multi-output JSON responses and parse coverage rather than free-form prose (Shorten et al. 2024). The other is tax-benefit microsimulation, where systems such as EUROMOD provide the methodological precedent for evaluating policy rules over household microdata (Sutherland and Figari 2013). There is also operational work on applying and evaluating AI systems in public-benefits settings, including caseworker-assist and Supplemental Nutrition Assistance Program (SNAP)-focused evaluations (Nava Labs 2026, 2025; ZenML LLMOps Database 2025). PolicyBench combines these strands into a public cross-model benchmark over household-level tax and benefit outputs.

### 4 Benchmark design

PolicyBench asks models to predict all benchmark outputs for a household in a single structured response. One response per household reduces repeated prompt cost and keeps the task at the

household-return level rather than turning it into a sequence of unrelated one-output calls.

The benchmark uses a bounded 0-100 score. For amount variables, the score averages four hit rates:

1. exact
2. within 1%
3. within 5%
4. within 10%

For binary outputs, the score is exact accuracy. This keeps 100 as the ceiling while still giving partial credit for near misses on amount outputs. PolicyBench also tracks mean absolute error and related error metrics, but those are secondary to the bounded score. This choice preserves exact-match comparability while avoiding the failure mode that recent numeric-evaluation papers have criticized (Abbood et al. 2025).

Because the four thresholds are nested ( $\text{exact} \subseteq \text{within-1\%} \subseteq \text{within-5\%} \subseteq \text{within-10\%}$ ), averaging their indicator functions is equivalent to assigning step credit by error band: predictions that match exactly receive 1.00, predictions inside 1% but not exact receive 0.75, predictions inside 5% but outside 1% receive 0.50, predictions inside 10% but outside 5% receive 0.25, and predictions outside 10% receive 0.00. The score is therefore a partial-credit rule that rewards tighter agreement more than looser agreement, not an unweighted aggregate of independent hit rates.

All requested US outputs are annual amounts or annual eligibility indicators for tax year 2026; UK outputs are annual amounts or annual eligibility indicators for fiscal year 2026-27. For currency amounts, the “exact” hit rate means within 1 currency unit of the reference value after numeric parsing. Percentage-threshold hit rates use relative error when the reference value is nonzero and the same 1 currency-unit absolute tolerance when the reference value is zero. Binary outputs are parsed as 0 or 1 eligibility flags and scored by exact classification accuracy. Missing or unparseable numeric answers receive zero score for that requested output.

Aggregation proceeds in four steps. First, each household-output prediction receives a 0-100 score. Second, each household receives one model score: requested output rows are weighted inside the household by a blend of equal weighting and absolute PolicyEngine reference-impact share. Person-level coverage flags are scored at the person row, with value-proxy impact weights where available. Third, country scores average those household scores with equal household weight. Fourth, global scores average the US and UK country scores for models run in both countries. The global score is therefore an equal-country, household-equal summary for this benchmark sample, not a universally authoritative measure of tax-benefit competence. Equal-output-group scores and other alternative views are reported as sensitivity checks.

The benchmark requires each model response to include numeric answers and one explanation per requested output. The headline score uses only the numeric answers. Explanations are retained for scenario exploration and qualitative error analysis; they should not be interpreted as faithful traces of model reasoning.

Because explanations are required, the canonical task measures policy estimation under a public-facing structured-response contract, not isolated arithmetic accuracy. Prompt fairness is part of the benchmark contract. The current release uses one prompt template per country, with no model-specific tuning. Models receive the same household facts, the same requested outputs, and no web or tool access. The prompt states that unlisted numeric inputs are 0, unlisted boolean or status facts are false, and household characteristics are constant over the tax-benefit year. Provider-specific differences are limited to the response transport needed to obtain structured output.

#### 4.1 Frozen snapshot and open-set status

Leaderboard positions are version-sensitive, so manuscript claims refer to the frozen source-run exports in Table 1 rather than to the live site. The committed source-run exports are the manuscript artifacts. The public site exposes the current prompts, predictions, explanations, and reference outputs for transparency. This makes the public leaderboard open-set: models, model providers, or benchmark users could learn from the released cases before later runs. Protected leaderboard claims would require a separate held-out or rotating set.

Source: [Article Notebook](#)

|    | Item                                   | Value                                                                              |
|----|----------------------------------------|------------------------------------------------------------------------------------|
| 0  | Snapshot date                          | 2026-05-15                                                                         |
| 1  | Model response date                    | 2026-05-13                                                                         |
| 2  | Policy period                          | US tax year 2026; UK fiscal year 2026-27                                           |
| 3  | Frozen export                          | paper/snapshot/20260501/runs/us_full_run_20260513_policyengine_4_4_4_nested_output |
| 4  | Frozen export SHA-256 prefixes         | US 30ad2c0bb5cf; UK 7c221d7c1b47                                                   |
| 5  | Snapshot manifest                      | paper/snapshot/20260501/manifest.json                                              |
| 6  | US scenarios SHA-256 prefix            | b05091225c06                                                                       |
| 7  | US reference outputs SHA-256 prefix    | 04febaadd091                                                                       |
| 8  | UK scenarios SHA-256 prefix            | eee79fb1db36                                                                       |
| 9  | UK reference outputs SHA-256 prefix    | 6286a238e54a                                                                       |
| 10 | Benchmark spec SHA-256 prefix          | 532e15ba19bf                                                                       |
| 11 | Frozen prompt payload SHA-256 prefixes | US f3f208702ac6; UK 2f64b11a99aa                                                   |
| 12 | Scoring code SHA-256 prefix            | 8f465e64a84b                                                                       |
| 13 | US run label                           | us_full_run_20260513_policyengine_4_4_4_nested_output                              |
| 14 | UK run label                           | uk_full_run_20260513_policyengine_4_4_4_nested_output                              |
| 15 | Household sample seed                  | 42                                                                                 |
| 16 | PolicyEngine.py                        | 4.4.4                                                                              |
| 17 | PolicyEngine-US                        | policyengine-us 1.691.10                                                           |
| 18 | PolicyEngine-UK                        | policyengine-uk 2.88.16                                                            |

|    | Item                             | Value                                                          |
|----|----------------------------------|----------------------------------------------------------------|
| 19 | US data bundle                   | policyengine-us-data 1.113.1                                   |
| 20 | UK runtime bundle                | policyengine-uk-data 1.55.5                                    |
| 21 | UK scenario source               | enhanced_cps_2025.h5                                           |
| 22 | UK scenario-source repository    | PolicyEngine/policyengine-uk-data                              |
| 23 | UK scenario-source pinned commit | 9514dfb7ec607897c9f7122a2e073b922c9fd8b6                       |
| 24 | UK scenario-source SHA-256       | 199ebc61d29231b4799ad337a95393765b5fb5aede1834b93ff2ace        |
| 25 | Households                       | 100 US and 100 UK                                              |
| 26 | Models                           | 12 shared models                                               |
| 27 | Output groups                    | 19 US and 7 UK                                                 |
| 28 | Condition                        | No tools, no web access, one structured response per household |
| 29 | Response contract                | Numeric answer and non-empty explanation for every request     |

Source: [Article Notebook](#)

Table 2: Model run configuration in the fro

|    | Model                         | Provider ID                          | US parsed   | UK parsed | Structured |
|----|-------------------------------|--------------------------------------|-------------|-----------|------------|
| 0  | gpt-5.5                       | gpt-5.5                              | 2,188/2,188 | 700/700   | provider s |
| 1  | claude-opus-4.7               | claude-opus-4-7                      | 2,188/2,188 | 699/700   | provider s |
| 2  | grok-4.20                     | xai/grok-4.20-reasoning              | 2,188/2,188 | 700/700   | provider s |
| 3  | claude-sonnet-4.6             | claude-sonnet-4-6                    | 2,104/2,188 | 665/700   | provider s |
| 4  | gemini-3.1-pro-preview        | gemini/gemini-3.1-pro-preview        | 2,068/2,188 | 665/700   | provider s |
| 5  | gemini-3-flash-preview        | gemini/gemini-3-flash-preview        | 2,167/2,188 | 694/700   | provider s |
| 6  | grok-4.3                      | xai/grok-4.3                         | 2,188/2,188 | 700/700   | provider s |
| 7  | gpt-5.4-mini                  | gpt-5.4-mini                         | 2,188/2,188 | 700/700   | provider s |
| 8  | claude-haiku-4.5              | claude-haiku-4-5-20251001            | 2,183/2,188 | 699/700   | provider s |
| 9  | gemini-3.1-flash-lite-preview | gemini/gemini-3.1-flash-lite-preview | 2,161/2,188 | 427/700   | provider s |
| 10 | gpt-5.4-nano                  | gpt-5.4-nano                         | 2,188/2,188 | 700/700   | provider s |
| 11 | grok-4.1-fast                 | xai/grok-4-1-fast-non-reasoning      | 1,401/2,188 | 112/700   | provider s |

Source: [Article Notebook](#)

## 5 Data and scenario construction

### 5.1 United States

The US benchmark is built from Enhanced Current Population Survey (CPS)-derived households using PolicyEngine US. The sampled households are filtered to keep a single-tax-unit, single-family, single-Supplemental Poverty Measure (SPM)-unit structure with at least one adult and a supported filing status. The 2024 Enhanced CPS source contains 41,314 households; 30,173 (73.0%) pass the filter and form the eligible draw. The 27.0% excluded by the filter include multi-tax-unit households (e.g., adult roommates), multi-family households, multi-SPM-unit households, and households whose head reports a filing status outside the supported set. These excluded compositions are exactly the kind of cases where federal/state credit allocations and benefit-unit rules become hardest, so the eligible draw is a tractable subset rather than the full distribution of US households. Prompts include nonzero promptable raw inputs across relevant entities rather than a hand-curated summary, so the models see many of the same facts the simulator receives. Filing status is not stated in the prompt; the reference computation infers it from tax-unit role flags. Models therefore see the same household facts that drive the reference filing-status assignment, but they do not receive that assignment as a label.

The current US release evaluates 19 output groups spanning federal income tax, refundable credits, payroll and self-employment tax, state and local income tax, Supplemental Nutrition Assistance Program (SNAP), Supplemental Security Income (SSI), Temporary Assistance for Needy Families (TANF), Affordable Care Act (ACA) premium tax credits, school-meal eligibility, and person-level coverage eligibility for the Special Supplemental Nutrition Program for Women, Infants, and Children (WIC), Medicaid, the Children’s Health Insurance Program (CHIP), Medicare, Head Start, and Early Head Start.

The output scope is intentionally narrower than the full PolicyEngine model. Table 3 summarizes the inclusion rule. The benchmark asks for WIC eligibility rather than a WIC dollar amount; WIC dollar values are used only as impact-weight proxies for coverage flags, not as requested model outputs.

Source: [Article Notebook](#)

|   | Scope decision          | Rationale                                                                                                     |
|---|-------------------------|---------------------------------------------------------------------------------------------------------------|
| 0 | Included                | Direct tax, credit, benefit, health-support, and coverage outputs that a household receives                   |
| 1 | Excluded                | Intermediate tax bases, payroll subcomponents, and outputs that mainly require household-level data           |
| 2 | ACA Premium Tax Credit  | Retained as a deliberate health-support output; when local benchmark premium is zero                          |
| 3 | Binary coverage outputs | Requested as 0/1 eligibility flags and scored as classification tasks; their dollar amounts are not requested |
| 4 | WIC                     | The benchmark asks for person-level WIC eligibility. It does not ask models to output dollar amounts          |

## 5.2 United Kingdom

The UK benchmark is built from a calibrated public transfer dataset scored through PolicyEngine UK. The current public build starts from a public export of benchmark-compatible households from PolicyEngine US Enhanced CPS, maps those records into UK-facing inputs, and recalibrates them to selected UK targets. The resulting `enhanced_cps_2025.h5` artifact is checked in to the public `PolicyEngine/policyengine-uk-data` GitHub repository; the manuscript pins commit `9514dfb7ec607897c9f7122a2e073b922c9fd8b6` so that a third party can retrieve the exact file used here. The artifact contains 28,532 households; 28,502 (99.9%) pass the eligibility filter that retains households with one benefit unit and one or two adults. This creates a public UK-policy transfer benchmark without publishing restricted household microdata; it is not a representative evaluation over native UK household records and should not be treated as a substitute for Family Resources Survey (FRS)-based UK microdata. The current UK release evaluates seven outputs: Income Tax, National Insurance, Capital Gains Tax, Child Benefit, Universal Credit, Pension Credit, and Personal Independence Payment (PIP). Outputs that depend on status or award facts use prompt-visible facts rather than hidden take-up labels; for example, PIP-positive cases list the daily living and mobility award components used by PolicyEngine.

The UK data path is more synthetic than the enhanced FRS pipeline and inherits limitations from cross-country transfer, calibration choices, and the subset of variables that can be made prompt-visible. It supports the current public cross-country benchmark, but it is not equivalent to an enhanced-FRS-based benchmark and should not be used to make population-representative claims about UK households (Sutherland and Figari 2013).

## 5.3 Reference-output credibility

PolicyBench treats PolicyEngine outputs as benchmark reference outputs, not as administrative records. The reference source is nevertheless stronger than an ad hoc answer key: PolicyEngine is open source, used for household calculators and reform analysis, and externally checked in specific domains. In the UK, No. 10's data science team adapted PolicyEngine's open-source microsimulation model for experimental policy simulation, with validation against external projections before use (Woodruff 2026; Ghenis 2026). In the US, PolicyEngine reports matching the National Bureau of Economic Research (NBER) TAXSIM-35 model to the cent on the vast majority of cases for the 2021 tax year across hundreds of thousands of tax units per state, with state-specific differences documented in PolicyEngine's integration tests (PolicyEngine 2024; Feenberg and Coutts 1993). We do not restate that comparison as a single percentage because the published source uses qualitative phrasing rather than a headline accuracy number. PolicyEngine has also signed a memorandum of understanding with the Federal Reserve Bank of Atlanta for future validation work against the Atlanta Fed's Policy Rules Database (Ghenis



and Makarchuk 2025; Federal Reserve Bank of Atlanta 2026). The Atlanta Fed sources are a caveat rather than evidence of completed validation for this benchmark: they document planned collaboration and the comparison source, not finished checks of the frozen PolicyBench outputs. Taken together, these sources support using PolicyEngine as a transparent reference implementation with partial external validation, but they do not validate every benchmark output.

This does not make PolicyEngine infallible. During benchmark development, we performed a manual, developer-led discrepancy review, with LLM assistance used to triage and summarize candidate cases surfaced by model explanations. Table 4 summarizes the main discrepancy classes reviewed before the frozen snapshot. In those reviewed classes, discrepancies reflected model mistakes, prompt ambiguity later corrected, or upstream data/model issues fixed before the frozen snapshot; the reviewed discrepancy classes did not identify unresolved PolicyEngine reference-output defects. This is a qualitative development audit, not an independent sampled validation study or exhaustive validation of every reference value.

Source: [Article Notebook](#)

Table 4: Development discrepancy review

|   | Discrepancy class              | Review outcome                                                          |
|---|--------------------------------|-------------------------------------------------------------------------|
| 0 | US federal credits             | Reviewed cases where explanations omitted EITC or refundable CTC;       |
| 1 | US SNAP, SSI, and Medicaid     | Reviewed false-zero and false-positive explanations; common model fail  |
| 2 | US payroll and overtime        | Reviewed overtime-related cases during development; upstream data/m     |
| 3 | UK Child Benefit               | Reviewed HICBC-related misses; the benchmark now asks for gross Ch      |
| 4 | UK National Insurance          | Reviewed cases where models subtracted pension contributions or used    |
| 5 | UK UC, Pension Credit, and PIP | Reviewed positive-case misses and false positives; explanations usually |

Source: [Article Notebook](#)

## 6 Results

### 6.1 United States leaderboard

The US leaderboard for the frozen manuscript snapshot is shown in Table 5. The top three models in that snapshot are gpt-5.5 (70.3), claude-opus-4.7 (67.5), and grok-4.20 (65.3).

Source: [Article Notebook](#)

Table 5: Top US benchmark models in the frozen manuscript snapshot.

|   | Model                  | Score | Exact | Within 10% | Parsed | Total |
|---|------------------------|-------|-------|------------|--------|-------|
| 0 | gpt-5.5                | 70.3  | 86.8  | 92.5       | 2188   | 2188  |
| 1 | claude-opus-4.7        | 67.5  | 85.1  | 90.3       | 2188   | 2188  |
| 2 | grok-4.20              | 65.3  | 86.8  | 91.9       | 2188   | 2188  |
| 3 | claude-sonnet-4.6      | 63.8  | 85.7  | 88.9       | 2104   | 2188  |
| 4 | gemini-3-flash-preview | 62.4  | 85.1  | 88.7       | 2167   | 2188  |

Source: [Article Notebook](#)

## 6.2 United Kingdom leaderboard

The UK leaderboard for the frozen manuscript snapshot is shown in Table 6. The top three models in that snapshot are gpt-5.5 (50.3), gemini-3.1-pro-preview (45.2), and grok-4.20 (45.1).

Source: [Article Notebook](#)

Table 6: Top UK benchmark models in the frozen manuscript snapshot.

|   | Model                  | Score | Exact | Within 10% | Parsed | Total |
|---|------------------------|-------|-------|------------|--------|-------|
| 0 | gpt-5.5                | 50.3  | 70.7  | 91.4       | 700    | 700   |
| 1 | gemini-3.1-pro-preview | 45.2  | 67.3  | 86.1       | 665    | 700   |
| 2 | grok-4.20              | 45.1  | 70.3  | 89.0       | 700    | 700   |
| 3 | claude-opus-4.7        | 45.1  | 69.7  | 88.6       | 699    | 700   |
| 4 | claude-sonnet-4.6      | 42.0  | 69.4  | 86.1       | 665    | 700   |

Source: [Article Notebook](#)

## 6.3 Global shared-model leaderboard

The global leaderboard averages US and UK scores for the models that have been run in both countries. In the frozen manuscript snapshot, the top three are gpt-5.5 (60.3), claude-opus-4.7 (56.3), and grok-4.20 (55.2). The global table is an equal-country summary for this two-country benchmark sample; it should be read alongside the country tables and sensitivity views rather than as a universally authoritative ordering.

Source: [Article Notebook](#)

Table 7: Global shared-model leaderboard across US and UK runs.

|   | Model                         | Score | Exact | Within 10% | Parsed | Total |
|---|-------------------------------|-------|-------|------------|--------|-------|
| 0 | gpt-5.5                       | 60.3  | 78.8  | 92.0       | 2888   | 2888  |
| 1 | claude-opus-4.7               | 56.3  | 77.4  | 89.4       | 2887   | 2888  |
| 2 | grok-4.20                     | 55.2  | 78.5  | 90.5       | 2888   | 2888  |
| 3 | claude-sonnet-4.6             | 52.9  | 77.6  | 87.5       | 2769   | 2888  |
| 4 | gemini-3.1-pro-preview        | 52.3  | 74.8  | 86.2       | 2733   | 2888  |
| 5 | gemini-3-flash-preview        | 51.3  | 77.5  | 86.7       | 2861   | 2888  |
| 6 | grok-4.3                      | 49.1  | 77.2  | 85.6       | 2888   | 2888  |
| 7 | gpt-5.4-mini                  | 43.5  | 76.7  | 80.9       | 2888   | 2888  |
| 8 | claude-haiku-4.5              | 42.9  | 76.1  | 79.8       | 2882   | 2888  |
| 9 | gemini-3.1-flash-lite-preview | 41.7  | 65.0  | 68.4       | 2588   | 2888  |

Source: [Article Notebook](#)




Figure 1: Global shared-model leaderboard in the frozen manuscript snapshot.

## 6.4 Simple baselines

Simple baselines help interpret score levels in a zero-heavy benchmark. Table 8 reports an always-zero response and a median-reference-by-output response on the same frozen household sample. These are not model competitors; they show how much of the bounded score can be earned without household-specific policy calculation.

Source: [Article Notebook](#)

Table 8: Simple baseline scores on the frozen manuscript snapshot.

|   | Baseline                   | Global score | US score | UK score |
|---|----------------------------|--------------|----------|----------|
| 0 | Always zero                | 26.5         | 25.3     | 27.6     |
| 1 | Median reference by output | 24.2         | 24.6     | 23.7     |

Source: [Article Notebook](#)

## 6.5 Sensitivity to benchmark view

The primary leaderboard gives each country equal weight globally and each household equal weight within a country. Table 9 reports several alternative views, including the older equal-output-group score. The top model is stable across several views, but lower ranks move across sensitivity views. The global score is therefore best read as a summary statistic whose interpretation depends on these checks.

Source: [Article Notebook](#)

Table 9: Leaderboard sensitivity to alternative scoring views.

|   | View                              | Rank 1           | Rank 2                        | Rank 3                        |
|---|-----------------------------------|------------------|-------------------------------|-------------------------------|
| 0 | Main household-equal impact score | gpt-5.5 (60.3)   | claude-opus-4.7 (56.3)        | grok-4.20 (55.2)              |
| 1 | Equal-output-group score          | gpt-5.5 (85.1)   | grok-4.20 (83.5)              | claude-opus-4.7 (82.5)        |
| 2 | Amount outputs only               | gpt-5.5 (82.5)   | grok-4.20 (80.3)              | claude-opus-4.7 (80.3)        |
| 3 | Positive reference cases only     | gpt-5.5 (43.4)   | claude-opus-4.7 (40.1)        | grok-4.20 (35.9)              |
| 4 | Zero reference cases only         | gpt-5.5 (98.6)   | grok-4.20 (98.3)              | gemini-3-flash-preview (98.3) |
| 5 | US only                           | gpt-5.5 (70.3)   | claude-opus-4.7 (67.5)        | grok-4.20 (65.3)              |
| 6 | UK only                           | gpt-5.5 (50.3)   | gemini-3.1-pro-preview (45.2) | grok-4.20 (45.1)              |
| 7 | US binary coverage only           | grok-4.20 (97.7) | gemini-3-flash-preview (96.9) | gpt-5.5 (96.8)                |

Source: [Article Notebook](#)

The household-equal impact score weights each requested output row inside a household by a blend of an equal-weight floor and the absolute reference-impact share. The default uses a 0.3 floor. Table 10 reports the top three global ranks under floors of 0.0, 0.1, 0.3, 0.5, and 1.0. Floor 1.0 is equal-row weighting within each household; floor 0.0 is pure impact-share weighting. Rank stability across this range means the main view is not driven by the specific floor choice.

Source: [Article Notebook](#)

Table 10: Top global ranks under varying household-equal impact-score floors.

|   | Floor | Rank 1         | Rank 2                 | Rank 3                 |
|---|-------|----------------|------------------------|------------------------|
| 0 | 0.0   | gpt-5.5 (49.5) | claude-opus-4.7 (44.7) | grok-4.20 (42.8)       |
| 1 | 0.1   | gpt-5.5 (53.1) | claude-opus-4.7 (48.6) | grok-4.20 (46.9)       |
| 2 | 0.3   | gpt-5.5 (60.3) | claude-opus-4.7 (56.3) | grok-4.20 (55.2)       |
| 3 | 0.5   | gpt-5.5 (67.5) | claude-opus-4.7 (63.9) | grok-4.20 (63.4)       |
| 4 | 1.0   | gpt-5.5 (85.6) | grok-4.20 (84.1)       | claude-opus-4.7 (83.1) |

Source: [Article Notebook](#)

## 6.6 Sampling uncertainty

The manuscript snapshot uses 100 households per country, so small score gaps should not be overinterpreted. Table 11 reports a household-resampling bootstrap over the frozen sample using 1,000 percentile bootstrap draws. It resamples households within each country, recomputes household-equal impact country scores, then recomputes the equal-country global score. The intervals therefore quantify uncertainty from the specific 100-household draw used here. They do not capture (i) prompt-template variance, since the manuscript uses a single template per country with no paraphrase resamples; (ii) decoding stochasticity, since each model’s snapshot prediction is a single sample at the provider’s default decoding settings; (iii) provider-side drift after 2026-05-13 in alias-resolved model weights; or (iv) reference-output uncertainty inside PolicyEngine. Total uncertainty over those sources is wider than the household-bootstrap intervals.

Source: [Article Notebook](#)

Table 11: Household-resampling uncertainty for the global leaderboard.

|   | Model           | Score | 95% bootstrap score interval | Bootstrap rank range |
|---|-----------------|-------|------------------------------|----------------------|
| 0 | gpt-5.5         | 60.3  | 57.2-63.3                    | 1-1                  |
| 1 | claude-opus-4.7 | 56.3  | 53.5-59.1                    | 2-4                  |

Table 11: Household-resampling uncertainty for the global leaderboard.

|   | Model                         | Score | 95% bootstrap score interval | Bootstrap rank range |
|---|-------------------------------|-------|------------------------------|----------------------|
| 2 | grok-4.20                     | 55.2  | 52.2-58.2                    | 2-5                  |
| 3 | claude-sonnet-4.6             | 52.9  | 49.9-56.1                    | 3-6                  |
| 4 | gemini-3.1-pro-preview        | 52.3  | 48.8-55.7                    | 3-7                  |
| 5 | gemini-3-flash-preview        | 51.3  | 48.2-54.5                    | 4-7                  |
| 6 | grok-4.3                      | 49.1  | 46.2-52.2                    | 5-7                  |
| 7 | gpt-5.4-mini                  | 43.5  | 40.5-46.7                    | 8-10                 |
| 8 | claude-haiku-4.5              | 42.9  | 39.9-46.0                    | 8-10                 |
| 9 | gemini-3.1-flash-lite-preview | 41.7  | 38.4-45.2                    | 8-10                 |

Source: [Article Notebook](#)

## 6.7 Hardest benchmark targets

The lowest-scoring US variables are multi-step tax quantities and sparse positive-dollar benefits. As shown in Table 12, federal income tax before refundable credits, state income tax before refundable credits, payroll tax, and SNAP are the hardest US outputs in the frozen manuscript snapshot.

Source: [Article Notebook](#)

Table 12: Hardest US variables by bounded score.

|    | Variable                                     | Score | Exact | Within 10% |
|----|----------------------------------------------|-------|-------|------------|
| 0  | federal_income_tax_before_refundable_credits | 42.7  | 35.7  | 52.1       |
| 16 | state_income_tax_before_refundable_credits   | 43.9  | 39.2  | 51.8       |
| 4  | payroll_tax                                  | 68.7  | 56.5  | 82.2       |
| 14 | snap                                         | 71.9  | 70.5  | 73.5       |
| 1  | federal_refundable_credits                   | 74.3  | 72.7  | 76.7       |

Source: [Article Notebook](#)

The same pattern appears in the UK run. Income tax and National Insurance are the two hardest UK outputs, while zero-heavy benefits and capital gains tax score higher overall.

Source: [Article Notebook](#)

Table 13: Hardest UK variables by bounded score.

|   | Variable           | Score | Exact | Within 10% |
|---|--------------------|-------|-------|------------|
| 2 | income_tax         | 31.4  | 19.8  | 46.8       |
| 3 | national_insurance | 49.2  | 40.6  | 63.9       |
| 1 | child_benefit      | 69.4  | 56.9  | 82.7       |
| 6 | universal_credit   | 70.6  | 68.8  | 73.2       |
| 4 | pension_credit     | 82.8  | 82.8  | 82.8       |

Source: [Article Notebook](#)

### 6.8 Zero and positive cases

Overall hit rates can overstate performance on sparse programs because correct zeros are common. Table 14 and Table 15 therefore report within-10% performance for all cases, positive-reference cases, and zero-reference cases. The gap is large for several outputs: models often identify that a program or tax does not apply, but miss the amount when the reference value is positive.

Source: [Article Notebook](#)

Table 14: US within-10% performance by zero versus positive reference cases.

|   | Variable                                     | All cases | Positive cases | Zero cases |
|---|----------------------------------------------|-----------|----------------|------------|
| 0 | state_income_tax_before_refundable_credits   | 54.6      | 26.7           | 93.7       |
| 1 | federal_income_tax_before_refundable_credits | 56.2      | 30.6           | 94.2       |
| 2 | snap                                         | 76.8      | 12.9           | 97.4       |
| 3 | federal_refundable_credits                   | 80.5      | 24.6           | 94.7       |
| 4 | state_refundable_credits                     | 81.4      | 4.7            | 98.9       |
| 5 | payroll_tax                                  | 85.9      | 83.0           | 94.9       |
| 6 | person_medicaid_eligible                     | 89.6      | 74.3           | 95.7       |
| 7 | person_head_start_eligible                   | 92.4      | 90.9           | 92.5       |

Source: [Article Notebook](#)

Table 15: UK within-10% performance by zero versus positive reference cases.

|   | Variable           | All cases | Positive cases | Zero cases |
|---|--------------------|-----------|----------------|------------|
| 0 | income_tax         | 53.4      | 40.8           | 92.9       |
| 1 | national_insurance | 71.7      | 50.7           | 95.2       |

Table 15: UK within-10% performance by zero versus positive reference cases.

|   | Variable          | All cases | Positive cases | Zero cases |
|---|-------------------|-----------|----------------|------------|
| 2 | universal_credit  | 82.8      | 24.2           | 98.0       |
| 3 | child_benefit     | 92.6      | 79.7           | 100.0      |
| 4 | pension_credit    | 92.8      | 1.5            | 98.9       |
| 5 | capital_gains_tax | 94.5      | 23.0           | 99.8       |
| 6 | pip               | 99.0      | NaN            | 99.0       |

Source: [Article Notebook](#)



Figure 2: Output-level performance on zero-reference and positive-reference cases.

## 7 Failure modes

The benchmark surfaces a few recurring failure patterns.



First, models miss positive tax and benefit quantities more often than zero cases. In the US, federal income tax before refundable credits, state income tax before refundable credits, payroll tax, and SNAP are the four lowest-scoring outputs. These require the model to choose the right income concepts, exclusions, program thresholds, and sequencing before applying any final subtraction.

Second, the UK benchmark shows the same split between tax calculations and many benefit outputs. Income tax and National Insurance score below the benefit outputs in the frozen manuscript snapshot. Positive Universal Credit and Pension Credit cases remain difficult, so the result should not be read as a general claim that benefits are easy.

Third, joint accuracy across interacting components is lower than marginal accuracy on either component. Table 16 shows within-10% accuracy for federal\_refundable\_credits, state\_refundable\_credits, and the conjunction of both within the same household. The joint hit rate is consistently lower than either marginal hit rate, so leaderboard scores that average across outputs understate how often a model gets a single household’s federal/state credit allocation jointly correct.

Source: [Article Notebook](#)

Table 16: US within-10% accuracy on federal vs state refundable credits and the household-level joint.

|    | Model                         | Federal within 10% | State within 10% | Joint within 10% |
|----|-------------------------------|--------------------|------------------|------------------|
| 8  | gpt-5.5                       | 94.0               | 82.0             | 77.0             |
| 10 | grok-4.20                     | 88.0               | 83.0             | 76.0             |
| 5  | gemini-3.1-pro-preview        | 83.0               | 79.0             | 72.0             |
| 1  | claude-opus-4.7               | 87.0               | 78.0             | 71.0             |
| 7  | gpt-5.4-nano                  | 79.0               | 81.0             | 71.0             |
| 2  | claude-sonnet-4.6             | 82.0               | 80.0             | 68.0             |
| 3  | gemini-3-flash-preview        | 77.0               | 81.0             | 66.0             |
| 4  | gemini-3.1-flash-lite-preview | 75.0               | 80.0             | 66.0             |
| 11 | grok-4.3                      | 75.0               | 81.0             | 66.0             |
| 6  | gpt-5.4-mini                  | 73.0               | 81.0             | 65.0             |
| 0  | claude-haiku-4.5              | 70.0               | 81.0             | 62.0             |
| 9  | grok-4.1-fast                 | 37.0               | 51.0             | 31.0             |

Source: [Article Notebook](#)

Fourth, structured-output reliability affects rankings. Missing or unparseable numeric values count as benchmark failures, and Appendix A reports the post-run parser audit. After a conservative reparse of saved raw responses, 271 of 2,400 country-model-household responses still have at least one missing output.

## 8 Limitations

PolicyBench is not a substitute for a production tax-and-benefit calculator. Several caveats matter:

- model rankings can reflect output-contract reliability as well as policy reasoning
- zero-heavy outputs require separate positive-case interpretation
- the scoring rule is one benchmark view and should be checked against alternative metrics
- the global score is an equal-country, household-equal summary for models run in both countries, not a universal ranking of model quality
- the public UK calibrated transfer dataset is not equivalent to enhanced Family Resources Survey quality and is not population-representative UK microdata
- the public scenario explorer exposes the current test set and reference outputs, so open-set leakage is a prominent limitation rather than a minor implementation detail
- the 100-household manuscript snapshot should be treated as a preview until larger frozen runs are published
- the discrepancy review is a qualitative developer audit, not an exhaustive external validation of every reference value
- benchmark success should not be interpreted as policy-advice readiness

The current paper is therefore an evaluation of model performance under a specific structured-output benchmark, not a general certification of tax or benefit competence.

## 9 Conclusion

PolicyBench shows a consistent pattern across both countries. Models often identify non-applicability, but positive tax and benefit amounts remain difficult, especially for multi-step income tax, payroll tax, National Insurance, and positive benefit cases. In the frozen manuscript snapshot, gpt-5.5 is the top-scoring model in both the US and UK.

These results support a narrow conclusion: unaided frontier models still struggle to reproduce selected household-level microsimulation outputs under a structured public benchmark. They do not show that models cannot assist policy analysis, and they do not validate PolicyEngine outputs as administrative truth. They suggest that future evaluations should separate no-tool estimation, tool-using system design, and reference-output validation more explicitly.

Next steps are to expand country coverage, increase frozen sample sizes, improve UK data provenance, and add protected or rotating evaluation sets so that public rankings are less exposed to open-set leakage. The benchmark should also continue reporting sensitivity views, because equal-country household-equal scores are useful summaries only when their weighting choices are visible.

## 10 Appendix A: Structured-output audit

The benchmark requires one numeric value and one explanation for every requested output. Parse failures are therefore benchmark failures, not missing data to drop from the denominator. After the frozen model run, we audited all rows classified as parse-contract failures, repaired one parser defect, and reran the parser against saved raw responses without making new model calls. Table 17 summarizes that sequence.

Source: [Article Notebook](#)

Table 17: Structured-output

|   | Step                         | Finding                                                                                                       |
|---|------------------------------|---------------------------------------------------------------------------------------------------------------|
| 0 | Initial parse-contract audit | 2,025 output rows were classified as parse-contract failures after the frozen model run.                      |
| 1 | Parser repair                | The parser was extended to accept nested ‘outputs’ payloads serialized as JSON.                               |
| 2 | Recovered outputs            | A conservative reparse recovered 42 formerly missing US numeric predictions.                                  |
| 3 | Remaining parse failures     | 1,983 of 34,656 model-output rows remain missing or unparseable after conservative reparse.                   |
| 4 | Retry unit                   | Any future second-chance sensitivity run should retry and replace an entire country-model-household response. |

Source: [Article Notebook](#)

The main denominator for response reliability is the country-model-household response: one model answering one household in one country. Table 18 reports that 271 of 2,400 such responses had at least one remaining missing output after the conservative reparse. The output-row count is larger because each response contains multiple requested outputs.

Source: [Article Notebook](#)

Table 18: Remaining parse failures by country-model-household response.

|   | Scope          | Responses with any remaining parse failure | Total country-model-household responses |
|---|----------------|--------------------------------------------|-----------------------------------------|
| 0 | Overall        | 271                                        | 2400                                    |
| 1 | United States  | 110                                        | 1200                                    |
| 2 | United Kingdom | 161                                        | 1200                                    |

Source: [Article Notebook](#)

The remaining failures are not one homogeneous bug. Table 19 groups the 1,983 remaining missing output rows, or 5.7% of 34,656 model-output rows, by the raw-response shape observed during the audit. Some rows have no saved raw response and therefore cannot be recovered by parser changes. Others contain malformed JSON, malformed nested outputs strings, or empty

tool arguments after repair attempts. The final column counts affected saved-response or chunk groups for that raw-response shape; it is not a mutually exclusive full-response denominator.

Source: [Article Notebook](#)

Table 19: Remaining missing output rows by raw-response shape.

|   | Remaining raw-response shape                           | Output rows | US rows | UK rows | Affected raw |
|---|--------------------------------------------------------|-------------|---------|---------|--------------|
| 0 | No saved raw response                                  | 834         | 554     | 280     | 64           |
| 1 | Repair attempts returned malformed JSON text           | 476         | 168     | 308     | 52           |
| 2 | Tool call with malformed ‘outputs’ string              | 291         | 123     | 168     | 31           |
| 3 | Repair attempts returned malformed ‘outputs’ string    | 250         | 110     | 140     | 26           |
| 4 | Repair attempts returned empty tool arguments          | 125         | 89      | 36      | 125          |
| 5 | Invalid saved raw JSON plus top-level ‘outputs’ object | 6           | 0       | 6       | 1            |
| 6 | Double-nested ‘outputs’ object                         | 1           | 0       | 1       | 1            |

Source: [Article Notebook](#)

Parse failures are concentrated in a few models. Table 20 reports the model-level response denominator across both countries. Models not shown in the table had no country-model-household response with a remaining parse failure after the conservative reparse.

Source: [Article Notebook](#)

Table 20: Remaining response-level parse failures by model.

|   | Model                         | Responses with any remaining parse failure | Total responses | Share (%) |
|---|-------------------------------|--------------------------------------------|-----------------|-----------|
| 0 | grok-4.1-fast                 | 121                                        | 200             | 60.5      |
| 1 | claude-sonnet-4.6             | 90                                         | 200             | 45.0      |
| 2 | gemini-3.1-flash-lite-preview | 40                                         | 200             | 20.0      |
| 3 | gemini-3.1-pro-preview        | 11                                         | 200             | 5.5       |
| 4 | claude-haiku-4.5              | 6                                          | 200             | 3.0       |
| 5 | gemini-3-flash-preview        | 2                                          | 200             | 1.0       |
| 6 | claude-opus-4.7               | 1                                          | 200             | 0.5       |

Source: [Article Notebook](#)

The concentration also differs by country. Table 21 reports the nonzero country-model cells. This is the most direct table for deciding whether a second-chance sensitivity run is worth the cost: a retry should target whole country-model-household responses in these cells and should preserve both the original and retry raw responses.

Source: [Article Notebook](#)

Table 21: Remaining response-level parse failures by country and model.

|    | Country | Model                         | Responses with any remaining parse failure | Total responses | S |
|----|---------|-------------------------------|--------------------------------------------|-----------------|---|
| 0  | UK      | grok-4.1-fast                 | 84                                         | 100             | 8 |
| 1  | UK      | gemini-3.1-flash-lite-preview | 39                                         | 100             | 3 |
| 2  | UK      | claude-sonnet-4.6             | 30                                         | 100             | 3 |
| 3  | UK      | gemini-3.1-pro-preview        | 5                                          | 100             | 5 |
| 4  | UK      | claude-haiku-4.5              | 1                                          | 100             | 1 |
| 5  | UK      | claude-opus-4.7               | 1                                          | 100             | 1 |
| 6  | UK      | gemini-3-flash-preview        | 1                                          | 100             | 1 |
| 7  | US      | claude-sonnet-4.6             | 60                                         | 100             | 6 |
| 8  | US      | grok-4.1-fast                 | 37                                         | 100             | 3 |
| 9  | US      | gemini-3.1-pro-preview        | 6                                          | 100             | 6 |
| 10 | US      | claude-haiku-4.5              | 5                                          | 100             | 5 |
| 11 | US      | gemini-3-flash-preview        | 1                                          | 100             | 1 |
| 12 | US      | gemini-3.1-flash-lite-preview | 1                                          | 100             | 1 |

Source: [Article Notebook](#)

We do not mix replacement values within a single household-model response in the canonical snapshot. Several outputs in the same response are mechanically related, so replacing only missing variables could combine two inconsistent answer sets. A future second-chance analysis should therefore rerun each broken country-model-household response, accept the retry only if it satisfies the full response contract, and report the original-run and retry-run scores separately.

## 11 Competing interests

The author is affiliated with PolicyEngine, which develops the microsimulation software used to produce the benchmark reference outputs.

Source: [Article Notebook](#)

Abbood, Auss, Zaiqiao Meng, and Nigel Collier. 2025. “Time to Revisit Exact Match.” Findings of the Association for Computational Linguistics: EMNLP 2025 (Suzhou, China), 11903–26. <https://doi.org/10.18653/v1/2025.findings-emnlp.637>.

Bock, Michael R., Kara Molisee, Zachary Ozer, and Sumit Shah. 2025. “TaxCalcBench: Evaluating Frontier Models on the Tax Calculation Task.” arXiv Preprint arXiv:2507.16126, ahead of print. <https://doi.org/10.48550/arXiv.2507.16126>.

- Chen, Zhiyu, Wenhui Chen, Charese Smiley, et al. 2021. “FinQA: A Dataset of Numerical Reasoning over Financial Data.” arXiv Preprint arXiv:2109.00122. <https://arxiv.org/abs/2109.00122>.
- Cobbe, Karl, Vineet Kosaraju, Mohammad Bavarian, et al. 2021. “Training Verifiers to Solve Math Word Problems.” arXiv Preprint arXiv:2110.14168. <https://arxiv.org/abs/2110.14168>.
- Federal Reserve Bank of Atlanta. 2026. Policy Rules Database. Federal Reserve Bank of Atlanta. <https://www.atlantafed.org/what-we-study/workforce-development/advancing-careers-for-low-income-families/policy-rules-database>.
- Feenberg, Daniel R., and Elisabeth Coutts. 1993. “An Introduction to the TAXSIM Model.” *Journal of Policy Analysis and Management* 12 (1): 189–94.
- Ghenis, Max. 2026. PolicyEngine Powers Rapid Policy Analysis at No 10 Downing Street. PolicyEngine. <https://www.policyengine.org/us/research/policyengine-10-downing-street>.
- Ghenis, Max, and Pavel Makarchuk. 2025. PolicyEngine and Atlanta Fed Sign Memorandum of Understanding for Policy Rules Database Validation. PolicyEngine. <https://www.policyengine.org/us/research/policyengine-atlanta-fed-mou-prd>.
- Guha, Neel, Julian Nyarko, Daniel E. Ho, et al. 2023. “LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models.” arXiv Preprint arXiv:2308.11462. <https://arxiv.org/abs/2308.11462>.
- Hendrycks, Dan, Collin Burns, Saurav Kadavath, et al. 2021. “Measuring Mathematical Problem Solving with the MATH Dataset.” *NeurIPS Datasets and Benchmarks Track*. <https://arxiv.org/abs/2103.03874>.
- Holzenberger, Nils, Benjamin Van Durme, Sarah Lawsky, and Kyle Richardson. 2021. “Factoring Statutory Reasoning as Language Understanding Challenges.” arXiv Preprint arXiv:2105.07903. <https://arxiv.org/abs/2105.07903>.
- Nava Labs. 2025. Experimenting with AI-Powered Tools in Public Benefits. Case study. <https://www.navapbc.com/case-studies/ai-tools-public-benefits>.
- Nava Labs. 2026. Evaluating a GenAI-Powered Assistive Chatbot for Caseworkers. Case study. <https://www.navapbc.com/case-studies/evaluating-ai-assistive-chatbot-caseworkers>.
- PolicyEngine. 2024. PolicyEngine Launches State Income Tax Modeling Nationwide. PolicyEngine. <https://www.policyengine.org/us/research/state-tax-model-beta>.
- Shanahan, Catherine, Emma McCarthy, Yan Zhao, et al. 2025. “Performance of LLMs on

- VITA Test: Potential for AI-Assisted Tax Returns for Low Income Taxpayers.” Artificial Intelligence and Law, ahead of print. <https://doi.org/10.1007/s10506-025-09465-7>.
- Shorten, Connor, Charles Pierse, Thomas Benjamin Smith, et al. 2024. “StructuredRAG: JSON Response Formatting with Large Language Models.” arXiv Preprint arXiv:2408.11061, ahead of print. <https://doi.org/10.48550/arXiv.2408.11061>.
- Sutherland, Holly, and Francesco Figari. 2013. “EUROMOD: The European Union Tax-Benefit Microsimulation Model.” International Journal of Microsimulation 6 (1): 4–26. <https://doi.org/10.34196/ijm.00075>.
- Woodruff, Nikhil. 2026. Informing Policy Using Micro-Simulations. No10 Innovation Fellowship. <https://fellows.ai.gov.uk/articles/nikhil-woodruff-micro-simulation>.
- ZenML LLMOps Database. 2025. Building and Automating Comprehensive LLM Evaluation Framework for SNAP Benefits. Industry writeup. <https://www.zenml.io/llmops-database/building-and-automating-comprehensive-llm-evaluation-framework-for-snap-benefits>.
- Zhou, Ruiwen, Wenyue Hua, Liangming Pan, et al. 2025. “RuleArena: A Benchmark for Rule-Guided Reasoning with LLMs in Real-World Scenarios.” Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (Vienna, Austria), 550–72. <https://doi.org/10.18653/v1/2025.acl-long.27>.
- Zhu, Fengbin, Wenqiang Lei, Youcheng Huang, et al. 2021. “TAT-QA: A Question Answering Benchmark on a Hybrid of Tabular and Textual Content in Finance.” arXiv Preprint arXiv:2105.07624. <https://arxiv.org/abs/2105.07624>.